

Machine Learning with R and Python

Don Hale

2025-09-22

Table of contents

About	6
1 Introduction	7
1.1 Introduction to Machine Learning	7
1.2 What is Machine Learning?	7
1.3 Major Types of Machine Learning	7
1.3.1 1. Supervised Learning	8
1.3.2 2. Unsupervised Learning	8
1.3.3 3. Semi-Supervised Learning	8
1.3.4 4. Reinforcement Learning	8
1.4 Objectives of Machine Learning	8
1.5 Bias-Variance Tradeoff	9
1.6 The Machine Learning Process	9
1.7 Summary	10
2 Model Selection and Regularization - R Version	11
2.1 Overview	11
2.2 Step 0: Setup and Data Preparation	11
2.3 Step 1: Stepwise Regression with 10-Fold Cross Validation	12
2.4 Step 2: Ridge and Lasso Regression	14
2.4.1 Visualizing RMSE and Model Complexity	14
2.5 Step 3: Elastic Net (Balancing Ridge and Lasso)	16
2.5.1 Visualizing Effects	17
2.6 Step 4: Comparing All Models on the Test Set	18
2.7 Key Takeaways	23
3 Model Selection and Regularization - Python Version	24
3.1 Overview	24
3.2 Step 0: Setup and Data	24
3.3 Step 1: Ridge Regression	26
3.4 Step 2: Lasso Regression	26
3.5 Step 3: Visualize Coefficient Shrinkage and RMSE	27
3.6 Step 4: Elastic Net	28
3.7 Step 5: Compare Models	29
3.8 Summary	29

4 GAM – Piecewise, MARS, LOESS, and Splines - R Version	30
4.1 1. Setup and Data	30
4.2 2. Quick Visual: Nonlinearity is Common	31
4.3 3. Piecewise Regression	32
4.4 4. MARS (earth) – From Simple to Interactions	34
4.4.1 4A. Univariate MARS (Garage_Area)	34
4.4.2 4B. Additive MARS (degree = 1; no interactions)	35
4.4.3 4C. MARS with 2-way Interactions (degree = 2)	39
4.5 5. LOESS (Visual local regression)	44
4.6 6. GAM Splines with mgcv – Simple → Complex	45
4.6.1 6A. Univariate GAM: Sale_Price ~ s(Garage_Area)	45
4.6.2 6B. Two-smooth Additive GAM: s(Garage_Area) + s(Gr_Liv_Area) .	46
4.6.3 6C. Full Multivariate GAM (selected smooths + factors)	47
4.7 7. Model Comparison (R ² , RMSE, MSE)	49
4.8 8. Final Plot: Predicted vs Actual (Test Set)	50
4.9 9. How to Explain	51
5 GAM— Piecewise, LOESS, and GAM Splines - Python Version	52
5.1 1. Setup and Data	52
5.2 2. Visual Check for Nonlinearity	54
5.3 3. Piecewise Regression Example	55
5.4 4. MARS (Commented Out — For Reference Only)	56
5.5 5. LOESS Visualization	57
5.6 6. GAM and Spline Approaches (Modern Alternatives)	58
5.7 7. Regression Spline Example (Fallback)	61
5.8 8. Model Comparison (Piecewise vs GAM vs Spline)	61
5.9 9. Predicted vs Actual Comparison	62
5.10 10. Interpretation Summary	63
6 Tree Based Models - R Version	64
6.1 Visual Roadmap	64
6.2 Libraries, Data, Split	65
6.3 Decision Tree (Baseline)	66
6.4 Random Forest: Random Variable, PDPs, and Importance	67
6.4.1 Random Forest Results	70
6.5 XGBoost: Random Variable, CV, Tuning, PDPs	71
6.5.1 XGBoost Results (Aligned Columns)	75
6.6 Final Model Comparison	75
7 Tree Based Models - Python Version	77
7.1 Visual Roadmap	77
7.2 Setup and Data Preparation	78
7.3 Decision Tree (Baseline)	79

7.4	Random Forest: Random Variable, Importance, PDPs, and Tuning	80
7.5	XGBoost: Random Variable, Importance, PDPs, and Tuning	83
7.6	Final Model Comparison	85
8	Understanding Random Forest and XGBoost Hyperparameters	87
8.1	Random Forest Hyperparameters	87
8.1.1	Switching Between Regression and Classification	87
8.1.2	Thinking About Random Forest Tuning	89
8.2	XGBoost Hyperparameters	89
8.2.1	Thinking About XGBoost Tuning	91
8.3	Practical Tips and Heuristics	91
8.4	Rule-of-Thumb Summary	92
8.4.1	Key Takeaways	92
9	Neural Networks in R	93
9.1	Introduction	93
9.2	1. Feed-Forward Neural Networks (Ames Data)	93
9.2.1	1.2 Cross-Validation Tuning	95
9.2.2	1.3 Evaluate Performance	97
9.3	2. Feed-Forward Neural Networks (Keras 3)	97
9.4	3. Recurrent Neural Network (RNN)	99
9.5	4. Convolutional Neural Network (CNN)	101
9.5.1	Explanation of the MNIST Preprocessing	102
9.5.2	4.1 Train and Evaluate the CNN	102
9.5.3	4.2 Visualizing Correct and Incorrect Predictions	104
9.6	Summary	105
10	Feedforward Neural Network - Python	106
10.1	Feedforward Neural Network (MLP) Using Ames Housing (Python)	106
10.1.1	1. Load and Prepare Data	106
10.1.2	2. Train/Test Split & Scaling	107
10.1.3	3. Simple MLP Model	107
10.1.4	4. Grid Search Tuning	108
11	Naive Bayes Classification R version	109
11.1	1. Overview and Learning Objectives	109
11.2	2. Data Setup and Binary Target Creation	110
11.3	3. Naive Bayes Intuition	111
11.4	4. Fitting Naive Bayes (Baseline Model)	111
11.5	5. Predicting on Test Data	114
11.6	6. Confusion Matrix Interpretation	115
11.7	7. ROC Curve & AUC	116
11.8	8. Tuning Naive Bayes Using <code>caret::train()</code>	118

11.9 9. Evaluating the Tuned Model on the Test Set	120
11.10 10. Summary	121
12 Naive Bayes Classification - Python Version	123
12.1 1. Overview and Learning Objectives	123
12.2 2. Setup and Data Preparation	123
12.3 2.1 Select Variables of Interest	124
12.4 2.2 Create a Binary Target Variable	124
12.5 2.3 Train/Test Split	125
12.6 3. Gaussian Naive Bayes Model	125
12.7 4. Predictions	126
12.8 5. Confusion Matrix & Classification Report	126
12.9 6. ROC Curve (Corrected) and AUC	127
12.10 7. Interpretation Notes	129
12.10.1 Naive Bayes Assumptions	129
12.10.2 Strengths	129
12.10.3 Weaknesses	129
12.11 8. Summary	129
13 model-agnostic-interpretability	130
14 support-vector-machines	131

About

This book is a companion to the Machine Learning Class. It will be a repository of R and Python Code. The chapters will be updated as we progress through the class.

1 Introduction

1.1 Introduction to Machine Learning

Machine Learning (ML) is a branch of artificial intelligence (AI) that focuses on building systems that **learn patterns from data** and make predictions or decisions without being explicitly programmed to perform specific tasks. Instead of writing rules by hand, a machine learning algorithm “learns” from examples in the data.

1.2 What is Machine Learning?

Formally, machine learning is the process of using data to **train a model** to make accurate predictions or decisions. The model identifies patterns in the training data and generalizes these patterns to new, unseen data.

ML can be thought of as:

- **Predictive modeling** – learning a function that maps inputs (features) to outputs (targets).
 - **Automated decision-making** – systems that improve their performance over time without explicit reprogramming.
-

1.3 Major Types of Machine Learning

Machine learning is commonly categorized into several types based on the **nature of the task** and **availability of labeled data**.

1.3.1 1. Supervised Learning

- Uses **labeled data**, meaning that each training example has an input and a known output.
- The goal is to **learn a function** that maps inputs to outputs accurately.
- Typical tasks:
 - **Regression**: Predict a continuous variable (e.g., house prices).
 - **Classification**: Predict a categorical variable (e.g., spam vs. non-spam email).

1.3.2 2. Unsupervised Learning

- Works with **unlabeled data**, where the output is unknown.
- The goal is to **discover patterns or structure** in the data.
- Typical tasks:
 - **Clustering**: Group similar observations together (e.g., customer segmentation).
 - **Dimensionality reduction**: Reduce the number of features while retaining important information (e.g., PCA).

1.3.3 3. Semi-Supervised Learning

- Combines a small amount of labeled data with a large amount of unlabeled data.
- Useful when labeling data is expensive or time-consuming.

1.3.4 4. Reinforcement Learning

- Focuses on **learning through trial and error**.
 - An agent learns to take actions in an environment to **maximize cumulative reward**.
 - Examples: Robotics, game AI, recommendation systems.
-

1.4 Objectives of Machine Learning

The main objective of machine learning is to **build models that generalize well** from historical data to make accurate predictions or decisions on **new, unseen data**. Key considerations include:

1. **Accuracy** – How well does the model predict outcomes?

2. **Generalization** – Does the model perform well on unseen data, or is it overfitting the training data?
 3. **Interpretability** – Can humans understand how the model makes predictions?
-

1.5 Bias-Variance Tradeoff

A fundamental concept in machine learning is the **bias-variance tradeoff**, which explains the balance between:

- **Bias**: Error due to overly simplistic models that **underfit** the data.
- **Variance**: Error due to overly complex models that **overfit** the training data.

The goal is to find a model that **minimizes the total prediction error**:

$$[\text{Total Error} = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}]$$

- **Underfitting** → High bias, low variance
- **Overfitting** → Low bias, high variance

The optimal model balances bias and variance for the best predictive performance.

1.6 The Machine Learning Process

A typical workflow in machine learning consists of the following steps:

1. **Data Collection**
 - Gather data from experiments, databases, or external sources.
2. **Data Preprocessing**
 - Handle missing values, outliers, and feature engineering.
3. **Model Selection**
 - Choose an appropriate algorithm (e.g., linear regression, random forest, neural networks).
4. **Model Training**

- Fit the model to the training data.

5. Model Evaluation

- Assess performance using metrics such as RMSE, accuracy, precision, recall, or AUC.

6. Model Tuning

- Adjust hyperparameters to improve performance.

7. Model Deployment

- Use the trained model to make predictions on new data.

8. Monitoring and Maintenance

- Continuously track model performance and retrain if necessary.
-

1.7 Summary

Machine learning allows us to build predictive models that **learn from data**. Understanding the different types of ML, their objectives, and the tradeoffs between bias and variance is critical to applying ML effectively. The remainder of this book explores practical algorithms, their implementation in **R** and **Python**, and strategies for model selection and interpretation.

2 Model Selection and Regularization - R Version

2.1 Overview

In this section, we'll use the **Ames Housing** dataset to demonstrate model selection and regularization.

We'll cover:

1. Splitting data into training and testing sets
 2. Performing **stepwise regression** with cross-validation
 3. Running **Ridge** and **Lasso** regression, and visualizing how α affects RMSE and model complexity
 4. Exploring **Elastic Net**, varying α to balance Ridge and Lasso
 5. Comparing all models on the test set
-

2.2 Step 0: Setup and Data Preparation

We'll use a reduced set of variables for speed and clarity.

```
# Suppress messages and warnings globally
knitr::opts_chunk$set(echo = TRUE, message = FALSE, warning = FALSE)

library(tidyverse)
library(caret)
library(glmnet)
library(AmesHousing)
```

```

library(GGally)

set.seed(123)

# Load data and split into 70% training, 30% testing
ames <- make_ordinal_ames() %>% mutate(id = row_number())
train <- ames %>% sample_frac(0.7)
test <- anti_join(ames, train, by = "id")

# Select a manageable subset of predictors
keep <- c("Sale_Price", "Bedroom_AbvGr", "Year_Built", "Mo_Sold", "Lot_Area",
          "Street", "Central_Air", "First_Flr_SF", "Second_Flr_SF", "Full_Bath",
          "Half_Bath", "Fireplaces", "Garage_Area", "Gr_Liv_Area", "TotRms_AbvGrd")

train <- train %>% select(all_of(keep))
test <- test %>% select(all_of(keep))

# Convert categorical variables to factors
train <- train %>% mutate(across(c(Street, Central_Air), as.factor))
test <- test %>% mutate(across(c(Street, Central_Air), as.factor))

```

2.3 Step 1: Stepwise Regression with 10-Fold Cross Validation

We'll use backward stepwise regression as a traditional benchmark.

```

ctrl <- trainControl(method = "cv", number = 10)

step_model <- train(
  Sale_Price ~ ., data = train,
  method = "lmStepAIC",
  trControl = ctrl,
  trace = FALSE,
  direction = "backward"
)


# Evaluate on the test set
step_pred <- predict(step_model, newdata = test)

```

```
step_perf <- postResample(step_pred, test$Sale_Price)
step_perf
```

```
RMSE      Rsquared      MAE
3.820776e+04 7.670269e-01 2.613098e+04
```

```
# Visualize predicted vs observed
ggplot(data.frame(obs = test$Sale_Price, pred = step_pred), aes(obs, pred)) +
  geom_point(alpha = 0.6) +
  geom_abline(linetype = "dashed", color = "red") +
  labs(title = "Stepwise Regression: Observed vs Predicted",
       x = "Observed Price", y = "Predicted Price")
```



Note:

Stepwise regression is intuitive and fast, but it can be unstable.

Stepwise regression (forward, backward, or both) selects predictors one at a time based on how much they improve a criterion (like AIC or adjusted R²).

2.4 Step 2: Ridge and Lasso Regression

Ridge and Lasso both shrink coefficients, but in different ways:

- Ridge shrinks all coefficients toward zero.
- Lasso can set some coefficients *exactly* to zero, performing variable selection.

```
x_train <- model.matrix(Sale_Price ~ ., train)[, -1]
y_train <- train$Sale_Price
x_test  <- model.matrix(Sale_Price ~ ., test)[, -1]
y_test  <- test$Sale_Price

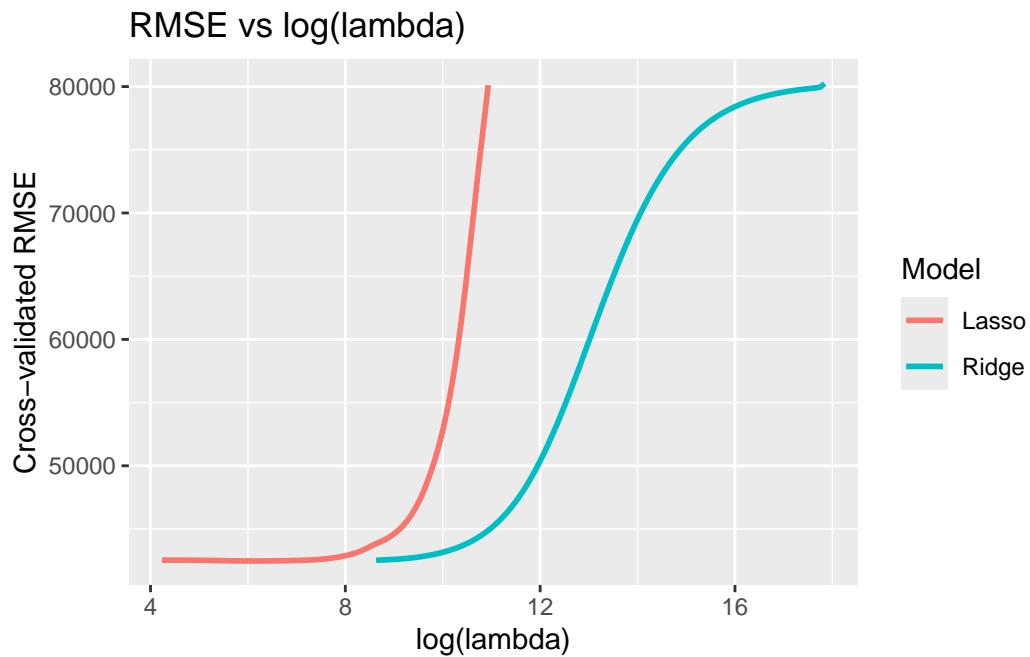
set.seed(123)
cv_ridge <- cv.glmnet(x_train, y_train, alpha = 0, nfolds = 10)
cv_lasso <- cv.glmnet(x_train, y_train, alpha = 1, nfolds = 10)
```

2.4.1 Visualizing RMSE and Model Complexity

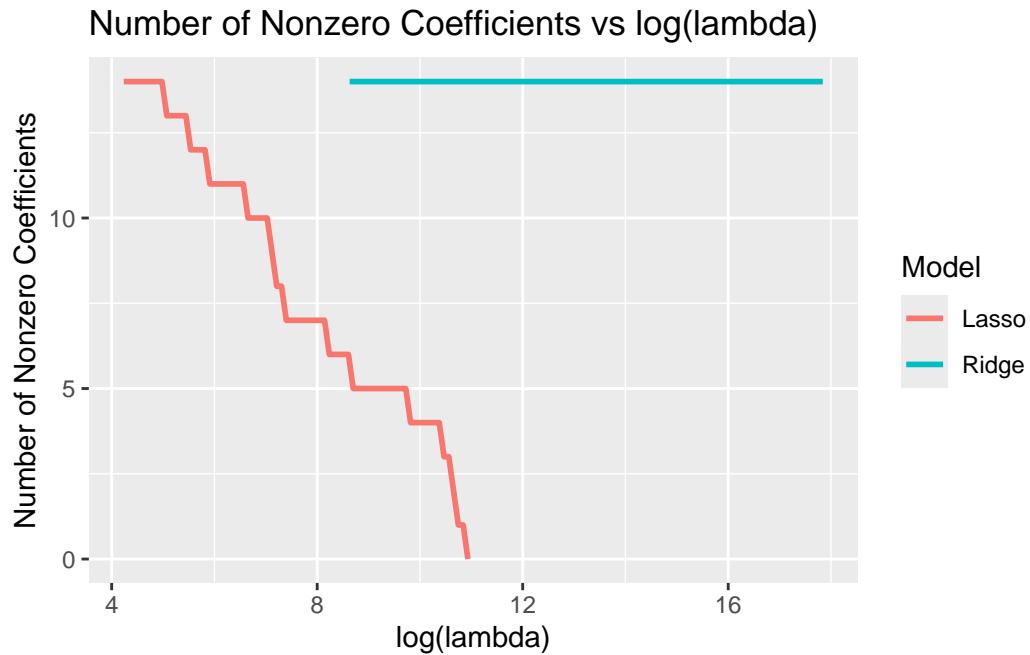
```
build_path_df <- function(cvfit, label) {
  fit <- cvfit$glmnet.fit
  tibble(
    lambda = fit$lambda,
    log_lambda = log(fit$lambda),
    RMSE = sqrt(cvfit$cvm),
    nonzero = colSums(abs(fit$beta) > 0),
    Model = label
  )
}

ridge_df <- build_path_df(cv_ridge, "Ridge")
lasso_df <- build_path_df(cv_lasso, "Lasso")
df <- bind_rows(ridge_df, lasso_df)

ggplot(df, aes(log_lambda, RMSE, color = Model)) +
  geom_line(size = 1) +
  labs(title = "RMSE vs log(lambda)", y = "Cross-validated RMSE", x = "log(lambda)")
```



```
ggplot(df, aes(log_lambda, nonzero, color = Model)) +  
  geom_line(size = 1) +  
  labs(title = "Number of Nonzero Coefficients vs log(lambda)",  
       y = "Number of Nonzero Coefficients", x = "log(lambda)")
```



Note:

- Increasing $\log(\lambda)$ increases regularization.
 - Ridge never eliminates variables, Lasso can.
 - There's a sweet spot where RMSE is minimized.
-

2.5 Step 3: Elastic Net (Balancing Ridge and Lasso)

Elastic Net introduces α to control the mix between Ridge ($\alpha = 0$) and Lasso ($\alpha = 1$).

```
alpha_grid <- seq(0, 1, by = 0.25)

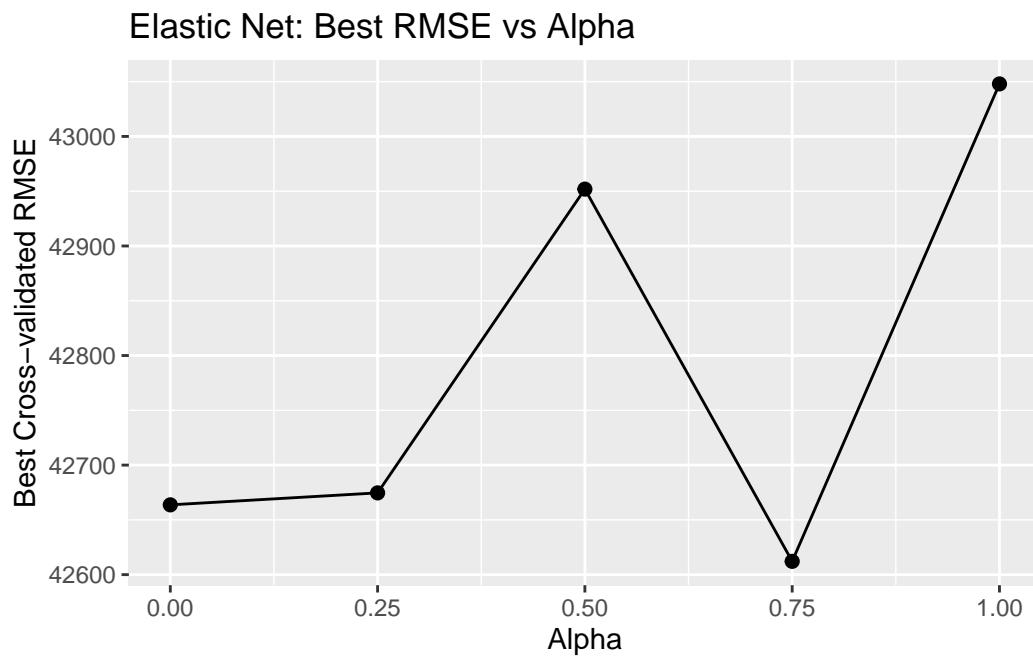
elastic_results <- map_df(alpha_grid, function(a) {
  cv_fit <- cv.glmnet(x_train, y_train, alpha = a, nfolds = 10)
  tibble(
    alpha = a,
    best_lambda = cv_fit$lambda.min,
    best_RMSE = sqrt(min(cv_fit$cvm)),
    nonzero = sum(abs(coef(cv_fit, s = "lambda.min")[-1]) > 0)
  )
})
```

```
})  
  
elastic_results
```

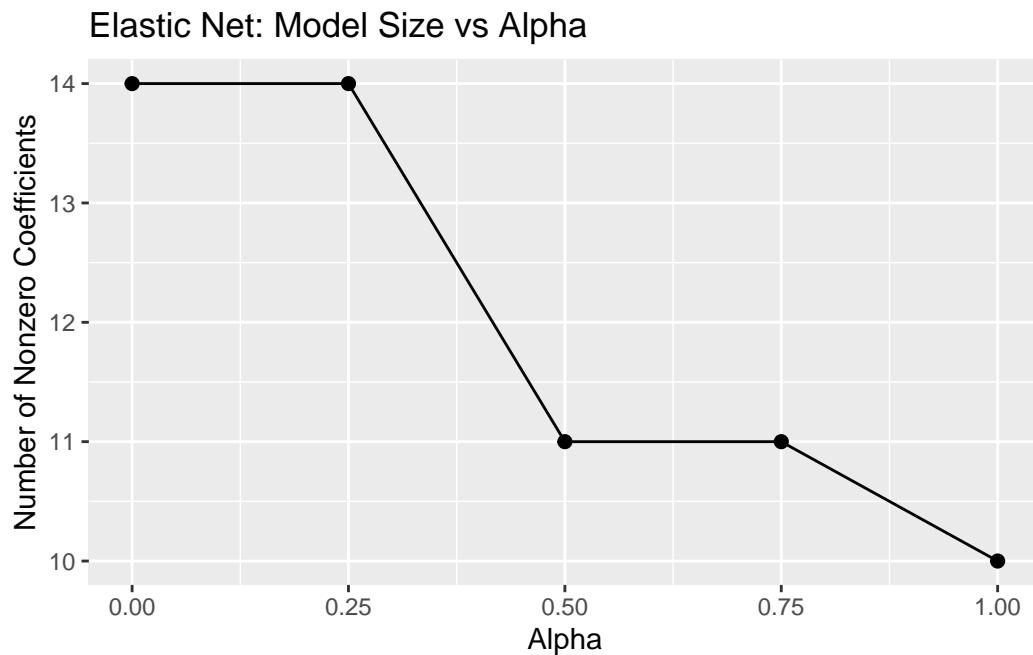
```
# A tibble: 5 x 4  
  alpha best_lambda best_RMSE nonzero  
  <dbl>      <dbl>     <dbl>    <int>  
1 0          5603.    42664.     14  
2 0.25       1224.    42675.     14  
3 0.5        1552.    42952.     11  
4 0.75       650.     42612.     11  
5 1          1126.    43048.     10
```

2.5.1 Visualizing Effects

```
ggplot(elastic_results, aes(alpha, best_RMSE)) +  
  geom_line() + geom_point(size = 2) +  
  labs(title = "Elastic Net: Best RMSE vs Alpha",  
       y = "Best Cross-validated RMSE", x = "Alpha")
```



```
ggplot(elastic_results, aes(alpha, nonzero)) +
  geom_line() + geom_point(size = 2) +
  labs(title = "Elastic Net: Model Size vs Alpha",
       y = "Number of Nonzero Coefficients", x = "Alpha")
```



Note:

Elastic Net combines the strengths of Ridge and Lasso: - Ridge for stability. - Lasso for sparsity. - Often performs best at intermediate values.

2.6 Step 4: Comparing All Models on the Test Set

We'll compare Stepwise, Ridge, Lasso, and the best Elastic Net.

```
ridge_pred <- predict(cv_ridge, newx = x_test, s = "lambda.min")
lasso_pred <- predict(cv_lasso, newx = x_test, s = "lambda.min")

best_alpha <- elastic_results %>% arrange(best_RMSE) %>% slice(1) %>% pull(alpha)
cv_en_best <- cv.glmnet(x_train, y_train, alpha = best_alpha, nfolds = 10)
elastic_pred <- predict(cv_en_best, newx = x_test, s = "lambda.min")
```

```

compare_tbl <- bind_rows(
  Stepwise = as_tibble_row(postResample(step_pred, test$Sale_Price)),
  Ridge = as_tibble_row(postResample(ridge_pred, y_test)),
  Lasso = as_tibble_row(postResample(lasso_pred, y_test)),
  Elastic = as_tibble_row(postResample(elastic_pred, y_test)),
  .id = "Model"
)

```

```

colnames(compare_tbl) <- c("Model", "RMSE", "Rsquared", "MAE")
compare_tbl <- compare_tbl %>% arrange(RMSE)
compare_tbl

```

```

# A tibble: 4 x 4
  Model      RMSE Rsquared     MAE
  <chr>    <dbl>   <dbl>   <dbl>
1 Stepwise 38208.    0.767 26131.
2 Lasso    38297.    0.767 26045.
3 Elastic  38301.    0.767 26039.
4 Ridge    38412.    0.766 26022.

```

```

# -----
# Display the best cross-validated Lasso, Ridge, and Elastic Net models

# -----
library(glmnet)
library(knitr)

# Helper function to summarize a fitted cv.glmnet model

summarize_glmnet <- function(cvfit, X_train, y_train, X_test, y_test, model_name) {
  cat("\n=====\\n")
  cat("Model:", model_name, "\\n")
  cat("=====\\n")

  # Extract best lambda
  best_lambda <- cvfit$lambda.min

```

```

cat("Best Lambda (lambda.min):", round(best_lambda, 6), "\n")

# Coefficients (non-zero)
coefs <- as.matrix(coef(cvfit, s = "lambda.min"))
nonzero <- coefs[coefs != 0, , drop = FALSE]

# Print nonzero coefficients
cat("\nNonzero Coefficients:\n")
print(knitr::kable(as.data.frame(nonzero), digits = 4))

# Predictions and performance
pred_train <- predict(cvfit, newx = X_train, s = "lambda.min")
pred_test  <- predict(cvfit, newx = X_test,   s = "lambda.min")

rmse <- function(y, yhat) sqrt(mean((y - yhat)^2))
rsq  <- function(y, yhat) 1 - sum((y - yhat)^2) / sum((y - mean(y))^2)

train_rmse <- rmse(y_train, pred_train)
test_rmse  <- rmse(y_test, pred_test)
test_r2     <- rsq(y_test, pred_test)

# Return one-row summary
data.frame(
  Model = model_name,
  Lambda = round(best_lambda, 6),
  Train_RMSE = round(train_rmse, 2),
  Test_RMSE = round(test_rmse, 2),
  Test_R2 = round(test_r2, 3),
  Nonzero_Coeff = nrow(nonzero)
)
}

# Summarize and display all three models
results <- rbind(
  summarize_glmnet(cv_lasso, x_train, y_train, x_test, y_test, "Lasso"),
  summarize_glmnet(cv_ridge, x_train, y_train, x_test, y_test, "Ridge"),
  summarize_glmnet(cv_en_best, x_train, y_train, x_test, y_test, "Elastic Net")
)
=====

Model: Lasso

```

```
=====
```

```
Best Lambda (lambda.min): 444.0297
```

```
Nonzero Coefficients:
```

	lambda.min
(Intercept)	-1536712.4502
Bedroom_AbvGr	-12871.7248
Year_Built	786.0097
Lot_Area	0.3320
StreetPave	14782.6929
Central_AirY	3463.4742
First_Flr_SF	22.9872
Half_Bath	-1509.7005
Fireplaces	11647.3233
Garage_Area	62.2549
Gr_Liv_Area	74.6167
TotRms_AbvGrd	1301.9029

```
=====
```

```
Model: Ridge
```

```
=====
```

```
Best Lambda (lambda.min): 5603.019
```

```
Nonzero Coefficients:
```

	lambda.min
(Intercept)	-1398127.9251
Bedroom_AbvGr	-13051.9562
Year_Built	706.9534
Mo_Sold	105.9352
Lot_Area	0.4001
StreetPave	22908.4375
Central_AirY	6951.0203
First_Flr_SF	45.3168
Second_Flr_SF	21.9030
Full_Bath	4203.4612
Half_Bath	-26.1950
Fireplaces	12666.1199

```

| Garage_Area |      64.5695 |
| Gr_Liv_Area |      42.3845 |
| TotRms_AbvGrd | 3067.4622 |

=====
Model: Elastic Net
=====
Best Lambda (lambda.min): 592.0396

Nonzero Coefficients:

|           | lambda.min |
|:-----|-----:|
|(Intercept) | -1532932.2050 |
|Bedroom_AbvGr | -12808.7530 |
|Year_Built | 783.9250 |
|Lot_Area | 0.3319 |
|StreetPave | 14821.6671 |
|Central_Airy | 3497.9772 |
|First_Flr_SF | 23.2544 |
|Half_Bath | -1319.3051 |
|Fireplaces | 11686.0135 |
|Garage_Area | 62.4230 |
|Gr_Liv_Area | 73.9778 |
|TotRms_AbvGrd | 1388.0576 |

```

```

# Combined performance table
cat("\n\n***Model Performance Comparison***\n")

```

Model Performance Comparison

```
kable(results, caption = "Performance of Cross-Validated Regularized Models")
```

Table 2.1: Performance of Cross-Validated Regularized Models

Model	Lambda	Train_RMSE	Test_RMSE	Test_R2	Nonzero_Coeff
Lasso	444.0297	41745.42	38296.94	0.764	12

Model	Lambda	Train_RMSE	Test_RMSE	Test_R2	Nonzero_Coeff
Ridge	5603.0194	41830.59	38411.78	0.763	15
Elastic Net	592.0396	41748.11	38300.92	0.764	12

2.7 Key Takeaways

- Stepwise selection is simple but can lead to unstable models.
- Ridge adds stability via coefficient shrinkage.
- Lasso enforces sparsity by removing weak predictors.
- Elastic Net balances both effects.
- Regularization often produces models that generalize better than purely stepwise ones.

3 Model Selection and Regularization - Python Version

3.1 Overview

This section parallels the R version of Model Selection and Regularization.
We will:

1. Load the Ames Housing data
 2. Split into training and test sets
 3. Perform Ridge, Lasso, and Elastic Net regressions with 10-fold CV
 4. Visualize RMSE and coefficient shrinkage patterns
 5. Compare test set performance
-

3.2 Step 0: Setup and Data

We'll use scikit-learn for modeling.

If the CSVs `ames_training.csv` and `ames_testing.csv` are not present, you can load Ames Housing via the `fetch_openml` function.

```
import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split, KFold, cross_val_score
```

```

from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import RidgeCV, LassoCV, ElasticNetCV
from sklearn.metrics import mean_squared_error

# Load Ames dataset directly
ames = fetch_openml(name="house_prices", as_frame=True)
df = ames.frame

# Subset variables to match the R version
keep = [
    "SalePrice", "BedroomAbvGr", "YearBuilt", "MoSold", "LotArea", "Street", "CentralAir",
    "1stFlrSF", "2ndFlrSF", "FullBath", "HalfBath", "Fireplaces", "GarageArea",
    "GrLivArea", "TotRmsAbvGrd"
]
df = df[keep].copy()

# Clean up names to match Python variable rules
df.columns = ["Sale_Price", "Bedroom_AbvGr", "Year_Built", "Mo_Sold", "Lot_Area", "Street",
              "Central_Air", "First_Flr_SF", "Second_Flr_SF", "Full_Bath", "Half_Bath",
              "Fireplaces", "Garage_Area", "Gr_Liv_Area", "TotRms_AbvGrd"]

# Drop missing rows
df = df.dropna()

# Train/test split (70/30)
train, test = train_test_split(df, test_size=0.3, random_state=123)

# Identify predictors and target
y_train = train["Sale_Price"]
y_test = test["Sale_Price"]
X_train = train.drop(columns="Sale_Price")
X_test = test.drop(columns="Sale_Price")

cat_vars = ["Street", "Central_Air"]
num_vars = [c for c in X_train.columns if c not in cat_vars]

# Preprocess: scale numeric, one-hot encode categorical
preprocessor = ColumnTransformer([
    ("num", StandardScaler(), num_vars),
    ("cat", OneHotEncoder(drop="first"), cat_vars)
])

```

```
])
```

3.3 Step 1: Ridge Regression

We'll use cross-validation to tune α (alpha).

Then we visualize RMSE vs log(lambda) and the number of nonzero coefficients.

```
alphas = np.logspace(4, -4, 80)
ridge = Pipeline([
    ("prep", preprocessor),
    ("model", RidgeCV(alphas=alphas, scoring="neg_mean_squared_error", cv=10))
])
ridge.fit(X_train, y_train)

ridge_best = ridge.named_steps["model"].alpha_
print("Best Ridge alpha:", ridge_best)

# Evaluate RMSE on test set
ridge_pred = ridge.predict(X_test)
ridge_rmse = np.sqrt(mean_squared_error(y_test, ridge_pred))
print("Test RMSE (Ridge):", ridge_rmse)
```

Best Ridge alpha: 94.33732216299774

Test RMSE (Ridge): 37055.84511170759

3.4 Step 2: Lasso Regression

Lasso adds variable selection — some coefficients go exactly to zero.

```
lasso = Pipeline([
    ("prep", preprocessor),
    ("model", LassoCV(alphas=alphas, cv=10, max_iter=20000))
])
lasso.fit(X_train, y_train)
```

```

lasso_best = lasso.named_steps["model"].alpha_
print("Best Lasso alpha:", lasso_best)

lasso_pred = lasso.predict(X_test)
lasso_rmse = np.sqrt(mean_squared_error(y_test, lasso_pred))
print("Test RMSE (Lasso):", lasso_rmse)

```

Best Lasso alpha: 11.568875283162821
 Test RMSE (Lasso): 37074.74978063246

3.5 Step 3: Visualize Coefficient Shrinkage and RMSE

Let's visualize how RMSE and model sparsity change with λ .

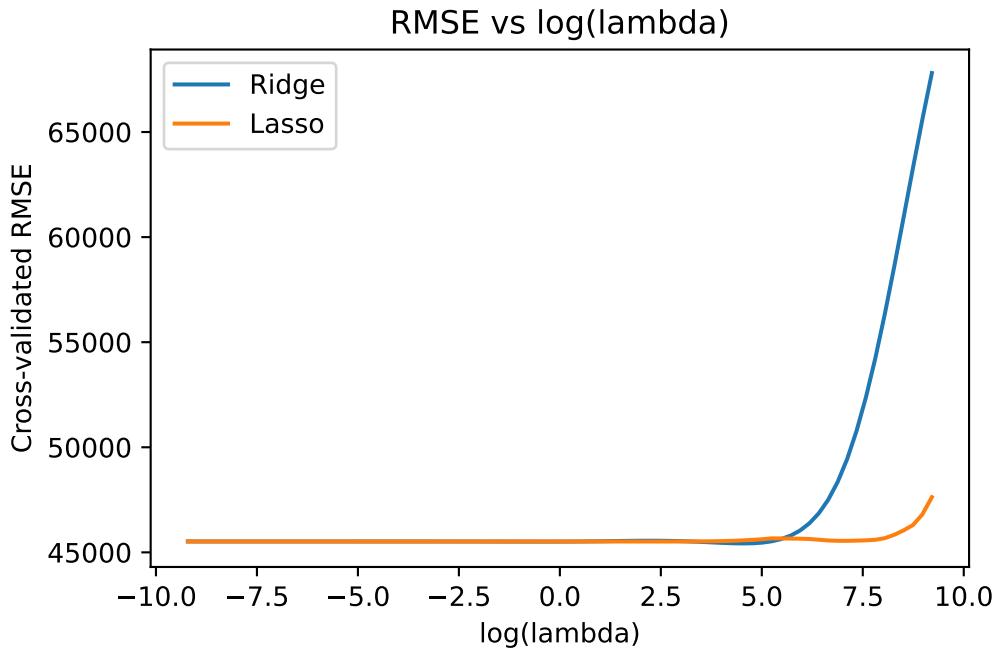
```

# Compute RMSE path manually
def cv_rmse(model_class, X, y, alphas):
    rmse = []
    for a in alphas:
        model = Pipeline([
            ("prep", preprocessor),
            ("model", model_class(alpha=a, max_iter=20000))
        ])
        scores = cross_val_score(model, X, y, scoring="neg_mean_squared_error", cv=KFold(10,
            rmse.append(np.sqrt(-scores.mean())))
    return np.array(rmse)

from sklearn.linear_model import Ridge, Lasso
ridge_rmse_path = cv_rmse(Ridge, X_train, y_train, alphas)
lasso_rmse_path = cv_rmse(Lasso, X_train, y_train, alphas)

plt.figure()
plt.plot(np.log(alphas), ridge_rmse_path, label="Ridge")
plt.plot(np.log(alphas), lasso_rmse_path, label="Lasso")
plt.xlabel("log(lambda)")
plt.ylabel("Cross-validated RMSE")
plt.title("RMSE vs log(lambda)")
plt.legend()
plt.show()

```



3.6 Step 4: Elastic Net

Elastic Net blends Ridge and Lasso.

We'll vary α (the mix) and visualize the tradeoff.

```
l1_ratios = np.linspace(0, 1, 6)
en_cv = Pipeline([
    ("prep", preprocessor),
    ("model", ElasticNetCV(l1_ratio=l1_ratios, alphas=alphas, cv=10, max_iter=50000))
])
en_cv.fit(X_train, y_train)

print("Best Elastic Net alpha:", en_cv.named_steps["model"].alpha_)
print("Best Elastic Net l1_ratio:", en_cv.named_steps["model"].l1_ratio_)

en_pred = en_cv.predict(X_test)
en_rmse = np.sqrt(mean_squared_error(y_test, en_pred))
print("Test RMSE (Elastic Net):", en_rmse)
```

```
Best Elastic Net alpha: 0.55825862688627
Best Elastic Net l1_ratio: 0.8
Test RMSE (Elastic Net): 37134.497858434275
```

3.7 Step 5: Compare Models

```
results = pd.DataFrame({
    "Model": ["Ridge", "Lasso", "Elastic Net"],
    "Test_RMSE": [ridge_rmse, lasso_rmse, en_rmse]
})
print(results.sort_values("Test_RMSE"))
```

	Model	Test_RMSE
0	Ridge	37055.845112
1	Lasso	37074.749781
2	Elastic Net	37134.497858

3.8 Summary

- Ridge stabilizes coefficients by shrinking them.
- Lasso enforces sparsity by zeroing weak predictors.
- Elastic Net blends both for balance.
- Regularization often beats pure stepwise regression on unseen data.

4 GAM – Piecewise, MARS, LOESS, and Splines - R Version

Goal: Predict `Sale_Price` from a focused set of Ames features using increasingly flexible methods in the GAM framework.

4.1 1. Setup and Data

```
set.seed(4321)

library(tidyverse)
library(dplyr)
library(tidyr)
library(AmesHousing)    # make_ordinal_ames()
library(earth)           # MARS
library(segmented)       # piecewise (segmented) regression
library(splines)          # regression splines (bs, ns) - kept for reference
library(mgcv)             # smoothing splines via GAM
library(caret)            # data splitting / utilities
library(Metrics)          # rmse, mse helpers (plus we'll compute R2 manually)
theme_set(theme_minimal())

# Limit to requested columns
cols <- c("Sale_Price", "Bedroom_AbvGr", "Year_Built", "Mo_Sold", "Lot_Area", "Street",
         "Central_Air", "First_Flr_SF", "Second_Flr_SF", "Full_Bath", "Half_Bath",
         "Fireplaces", "Garage_Area", "Gr_Liv_Area", "TotRms_AbvGrd")

ames <- make_ordinal_ames() |>
  dplyr::select(dplyr::all_of(cols)) |>
  tidyr::drop_na()

# Train/test split
set.seed(4321)
```

```

idx <- sample.int(nrow(ames), size = floor(0.7*nrow(ames)))
train <- ames[idx,]
test <- ames[-idx,]

glimpse(train)

```

Rows: 2,051
 Columns: 15

```

$ Sale_Price      <int> 274000, 75200, 329900, 145400, 108000, 184000, 176000, 1~  

$ Bedroom_AbvGr <int> 3, 2, 4, 3, 2, 2, 3, 3, 3, 3, 3, 3, 4, 3, 2, 3, 3, 2, ~  

$ Year_Built     <int> 2001, 1922, 2005, 1926, 1949, 1999, 1962, 1915, 1999, 19~  

$ Mo_Sold        <int> 1, 9, 8, 5, 5, 6, 5, 6, 9, 6, 8, 5, 7, 7, 4, 8, 5, 11, 5~  

$ Lot_Area       <int> 9720, 3672, 11643, 7000, 8777, 5858, 19296, 8094, 3768, ~  

$ Street          <fct> Pave, Pave, Pave, Pave, Pave, Pave, Pave, Pa~  

$ Central_Air    <fct> Y, ~  

$ First_Flr_SF   <int> 1366, 816, 1544, 861, 1126, 1337, 1382, 1048, 713, 792, ~  

$ Second_Flr_SF  <int> 581, 0, 814, 424, 0, 0, 0, 720, 739, 725, 0, 1151, 695, ~  

$ Full_Bath       <int> 2, 1, 2, 1, 2, 2, 1, 1, 2, 2, 2, 1, 2, 2, 1, 1, ~  

$ Half_Bath       <int> 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, ~  

$ Fireplaces      <int> 1, 0, 1, 0, 0, 1, 1, 0, 0, 2, 0, 1, 2, 1, 0, 1, 0, 0, 1, ~  

$ Garage_Area    <dbl> 725, 100, 784, 506, 520, 511, 884, 576, 506, 400, 0, 434~  

$ Gr_Liv_Area    <int> 1947, 816, 2358, 1285, 1126, 1337, 1382, 1768, 1452, 151~  

$ TotRms_AbvGrd  <int> 7, 5, 10, 6, 5, 5, 6, 8, 6, 7, 5, 8, 7, 5, 6, 7, 6, 5~
```

```
summary(train$Sale_Price)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	12789	130000	161000	180897	215000	755000

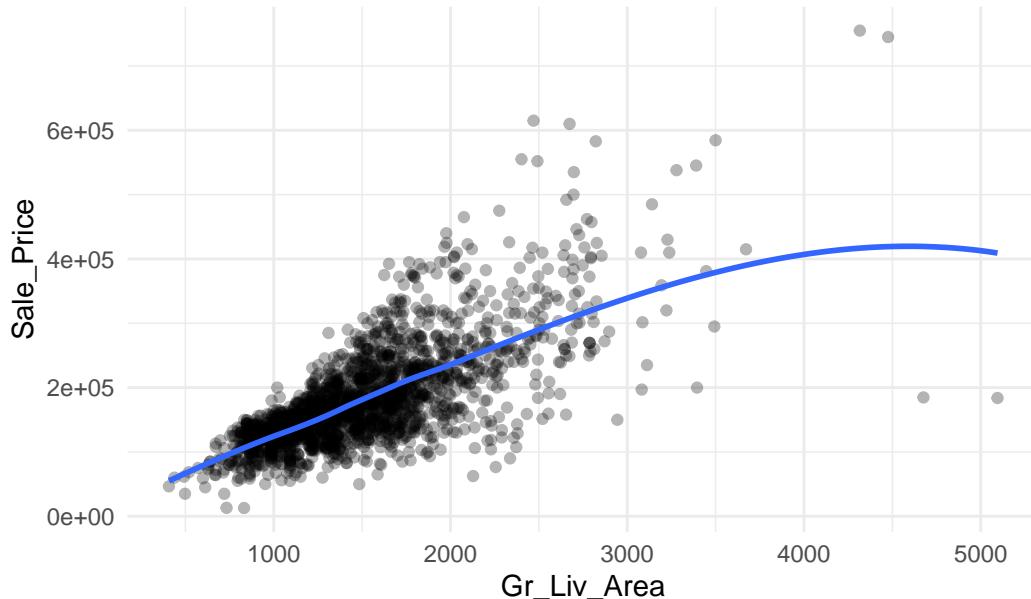
4.2 2. Quick Visual: Nonlinearity is Common

```

ggplot(train, aes(Gr_Liv_Area, Sale_Price)) +
  geom_point(alpha=.3) +
  geom_smooth(method="loess", se=FALSE) +
  labs(title="Sale_Price vs Gr_Liv_Area (LOESS smoother)")

```

Sale_Price vs Gr_Liv_Area (LOESS smoother)



Takeaway: The relationship bends—nonlinear methods can help.

4.3 3. Piecewise Regression

```
m_lin <- lm(Sale_Price ~ Gr_Liv_Area, data=train)
m_seg_init <- lm(Sale_Price ~ Gr_Liv_Area, data=train)
m_seg <- segmented(m_seg_init, seg.Z = ~Gr_Liv_Area, psi = list(Gr_Liv_Area = 2000))
summary(m_seg)
```

```
***Regression Model with Segmented Relationship(s)***

Call:
segmented.lm(obj = m_seg_init, seg.Z = ~Gr_Liv_Area, psi = list(Gr_Liv_Area = 2000))

Estimated Break-Point(s):
      Est.  St.Err
psi1.Gr_Liv_Area 2466 332.881
```

```

Coefficients of the linear terms:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 7901.186   4698.641   1.682   0.0928 .  
Gr_Liv_Area 115.705     3.171  36.488 <2e-16 *** 
U1.Gr_Liv_Area -23.479    12.436  -1.888      NA    
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```


Residual standard error: 55110 on 2047 degrees of freedom
Multiple R-Squared: 0.5124, Adjusted R-squared: 0.5116

Boot restarting based on 6 samples. Last fit:
Convergence attained in 1 iterations (rel. change 0)

```

plot(train$Gr_Liv_Area, train$Sale_Price, pch=16, cex=.5,
      xlab="Gr_Liv_Area", ylab="Sale_Price", main="Piecewise fit at estimated knot")
plot(m_seg, add=TRUE, col=2, lwd=2)
abline(v = m_seg$psi[, "Est."], lty=2, col=2)

```



Interpretation (piecewise): Two linear slopes before/after a knot. Useful when you expect a threshold effect.

4.4 4. MARS (earth) – From Simple to Interactions

MARS = *Multivariate Adaptive Regression Splines*. It builds **hinge** basis functions like $h(x - c) = \max(0, x - c)$ and can **interact** them.

- A term such as $h(\text{Gr_Liv_Area} - 1500)$ means *once* `Gr_Liv_Area` exceeds 1500, the fitted line's slope can change.
- An interaction like $h(\text{Gr_Liv_Area} - 1500) * \text{Central_AirY}$ means the slope change applies when `Central_Air == "Y"`.

4.4.1 4A. Univariate MARS (Garage_Area)

```
mars1 <- earth(Sale_Price ~ Garage_Area, data=train)
summary(mars1)
```

```
Call: earth(formula=Sale_Price~Garage_Area, data=train)
```

```
              coefficients
(Intercept)          124159.039
h(286-Garage_Area)      -60.257
h(Garage_Area-286)       297.277
h(Garage_Area-521)      -483.642
h(Garage_Area-576)       733.859
h(Garage_Area-758)      -356.460
h(Garage_Area-1043)     -490.873
```

```
Selected 7 of 7 terms, and 1 of 1 predictors
```

```
Termination condition: RSq changed by less than 0.001 at 7 terms
```

```
Importance: Garage_Area
```

```
Number of terms at each degree of interaction: 1 6 (additive model)
```

```
GCV 3427475346    RSS 6.94092e+12    GRSq 0.4492014    RSq 0.4556309
```

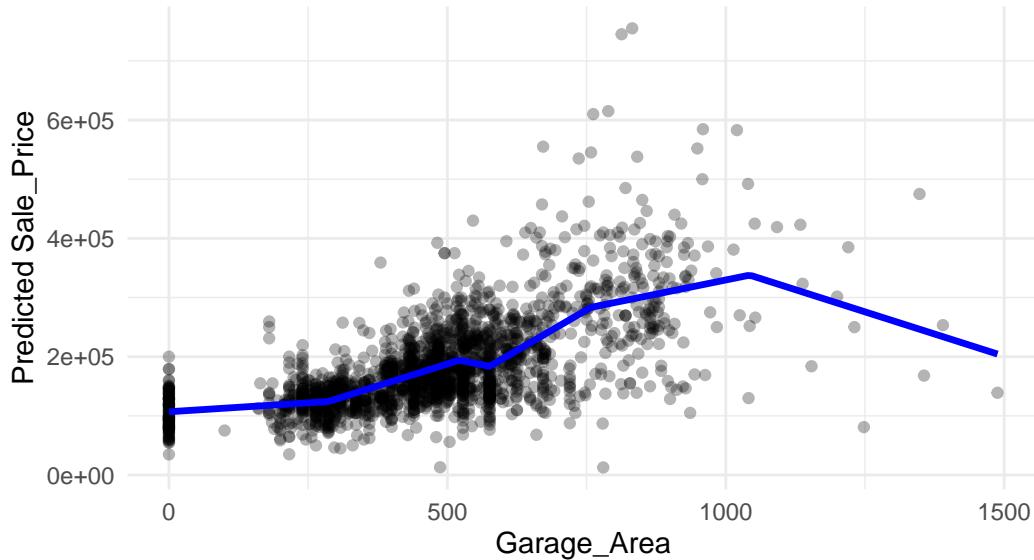
```
grid <- tibble(Garage_Area = seq(min(train$Garage_Area), max(train$Garage_Area), length.out=100))
grid$pred <- predict(mars1, newdata=grid)

ggplot(train, aes(Garage_Area, Sale_Price)) +
  geom_point(alpha=.3) +
  geom_line(data=grid, aes(Garage_Area, pred), color="blue", linewidth=1.2) +
```

```
labs(title="Step 1: MARS with One Predictor (Garage_Area)",
     subtitle="Piecewise linear fit with automatically chosen knots",
     y="Predicted Sale_Price")
```

Step 1: MARS with One Predictor (Garage_Area)

Piecewise linear fit with automatically chosen knots



4.4.2 4B. Additive MARS (degree = 1; no interactions)

```
mars2 <- earth(Sale_Price ~ Bedroom_AbvGr + Year_Built + Mo_Sold + Lot_Area +
                 Street + Central_Air + First_Flr_SF + Second_Flr_SF + Full_Bath +
                 Half_Bath + Fireplaces + Garage_Area + Gr_Liv_Area + TotRms_AbvGrd,
                 data=train, degree=1, nfold=5)
summary(mars2)
```

```
Call: earth(formula=Sale_Price~Bedroom_AbvGr+Year_Built+Mo_Sold+Lot_Ar...),
           data=train, degree=1, nfold=5)
```

	coefficients
(Intercept)	319493.46
Central_AirY	20289.49
h(4-Bedroom_AbvGr)	9214.66
h(Bedroom_AbvGr-4)	-23009.05

h(Year_Built-1977)	1275.57
h(2004-Year_Built)	-336.64
h(Year_Built-2004)	5315.57
h(13869-Lot_Area)	-2.09
h(Lot_Area-13869)	0.22
h(First_Flr_SF-1600)	104.91
h(2402-First_Flr_SF)	-71.56
h(First_Flr_SF-2402)	-176.61
h(1523-Second_Flr_SF)	-53.13
h(Second_Flr_SF-1523)	426.63
h(Half_Bath-1)	-45378.31
h(2-Fireplaces)	-14408.56
h(Fireplaces-2)	-26072.58
h(Garage_Area-539)	101.97
h(Garage_Area-1043)	-294.30
h(Gr_Liv_Area-2049)	65.21
h(Gr_Liv_Area-3194)	-159.79

Selected 21 of 24 terms, and 10 of 14 predictors

Termination condition: Reached nk 29

Importance: First_Flr_SF, Second_Flr_SF, Year_Built, Garage_Area, ...

Number of terms at each degree of interaction: 1 20 (additive model)

GCV 1033819964 RSS 2.036439e+12 GRSq 0.8338641 RSq 0.8402842 CVRSq 0.7978671

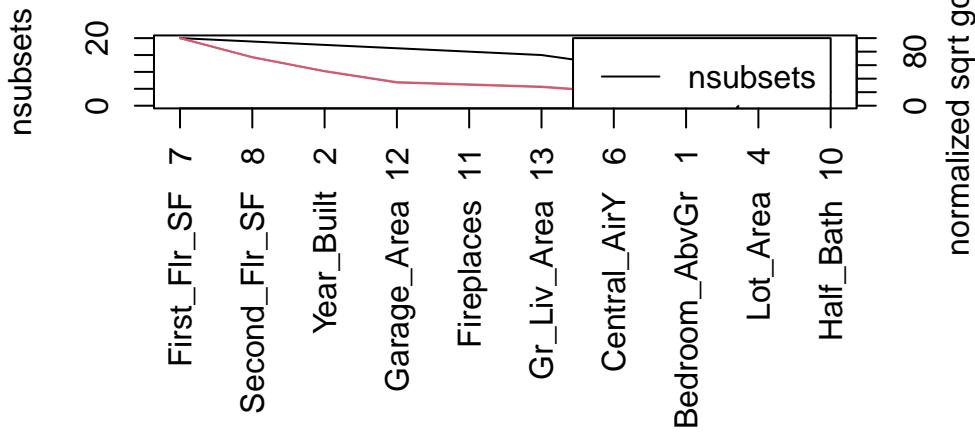
Note: the cross-validation sd's below are standard deviations across folds

Cross validation: nterms 21.00 sd 0.71 nvars 9.20 sd 0.84

CVRSq	sd	MaxErr	sd
0.798	0.043	-428941	282429

```
ev2 <- evimp(mars2); plot(ev2, main="Step 2: Variable Importance (degree=1, additive)")
```

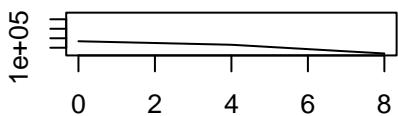
Step 2: Variable Importance (degree=1, additive)




```
par(mfrow=c(2,2)); plotmo(mars2, type="response", nresponse=1, do.par=FALSE); par(mfrow=c(1,1))
```

```
plotmo grid:      Bedroom_AbvGr Year_Built Mo_Sold Lot_Area Street Central_Air
                  3           1974     6       9480   Pave        Y
First_Flr_SF Second_Flr_SF Full_Bath Half_Bath Fireplaces Garage_Area
      1092          0         2         0         1        480
Gr_Liv_Area TotRms_AbvGrd
      1442          6
```

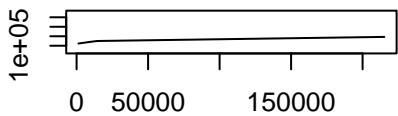
1 Bdrm_AbvG




2 Year_Bult



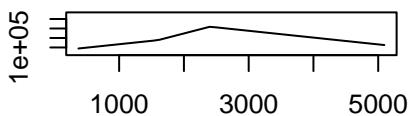
3 Lot_Area




4 Centrl_Ar



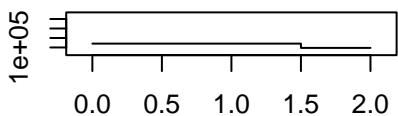
5 Frst_F_SF




6 Scnd_F_SF

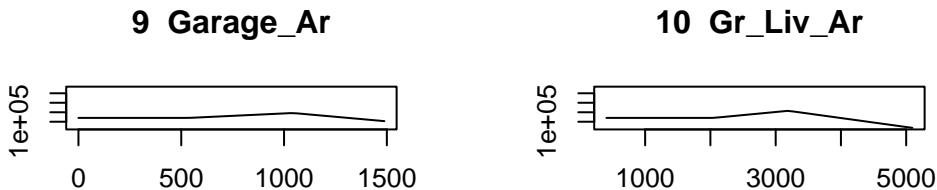


7 Half_Bath



8 Fireplacs





4.4.3 4C. MARS with 2-way Interactions (degree = 2)

```

mars3 <- earth(Sale_Price ~ Bedroom_AbvGr + Year_Built + Mo_Sold + Lot_Area +
                 Street + Central_Air + First_Flr_SF + Second_Flr_SF + Full_Bath +
                 Half_Bath + Fireplaces + Garage_Area + Gr_Liv_Area + TotRms_AbvGrd,
                 data=train, degree=2, nfold=5)
summary(mars3)

```

Call: earth(formula=Sale_Price~Bedroom_AbvGr+Year_Built+Mo_Sold+Lot_Ar...),
 data=train, degree=2, nfold=5)

	coefficients
(Intercept)	411949.27
h(Year_Built-2004)	169869.37
h(14803-Lot_Area)	-1.75
h(Lot_Area-14803)	0.48
h(1570-First_Flr_SF)	-150.79
h(First_Flr_SF-1570)	224.90
h(1523-Second_Flr_SF)	-86.91
h(Second_Flr_SF-1523)	156.02
h(1-Half_Bath)	-7743.19

h(Half_Bath-1)	-54819.24
h(2-Fireplaces)	-10516.92
h(1043-Garage_Area)	-70.97
h(Garage_Area-1043)	-141.86
h(Gr_Liv_Area-3194)	106.27
h(3-Bedroom_AbvGr) * h(1523-Second_Flr_SF)	7.01
h(Bedroom_AbvGr-3) * h(1523-Second_Flr_SF)	-9.61
h(2004-Year_Built) * h(First_Flr_SF-876)	-1.86
h(2004-Year_Built) * h(3194-Gr_Liv_Area)	-0.17
h(Year_Built-2004) * h(Gr_Liv_Area-2320)	-159.80
h(Year_Built-2004) * h(2320-Gr_Liv_Area)	156.64
h(Year_Built-2004) * h(3194-Gr_Liv_Area)	-173.05
h(First_Flr_SF-1778) * h(2-Fireplaces)	-78.19
h(1043-Garage_Area) * h(2122-Gr_Liv_Area)	0.05

Selected 23 of 29 terms, and 9 of 14 predictors

Termination condition: Reached nk 29

Importance: First_Flr_SF, Second_Flr_SF, Year_Built, Gr_Liv_Area, ...

Number of terms at each degree of interaction: 1 13 9

GCV 795873462 RSS 1.544416e+12 GRSq 0.8721024 RSq 0.8788731 CVRSq 0.7824362

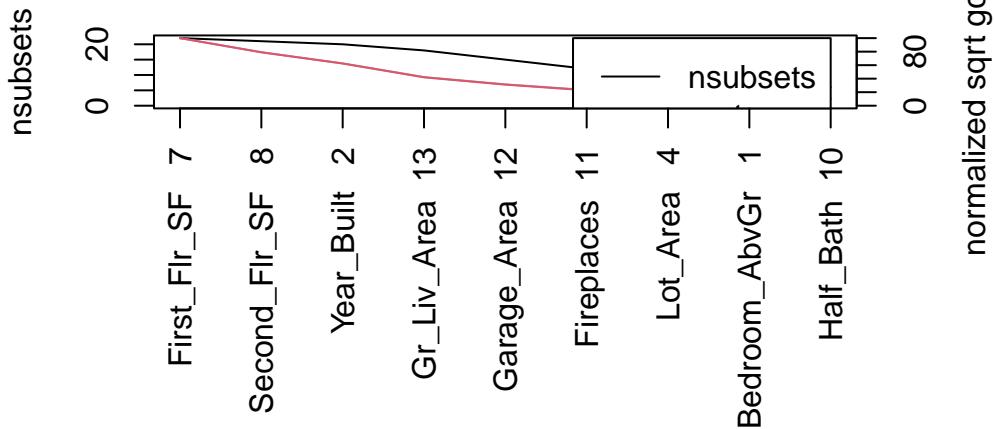
Note: the cross-validation sd's below are standard deviations across folds

Cross validation: nterms 22.60 sd 1.82 nvars 9.40 sd 1.14

CVRSq	sd	MaxErr	sd
0.782	0.114	-945502	422348

```
ev3 <- evimp(mars3); plot(ev3, main="Step 3: Variable Importance (degree=2, interactions)")
```

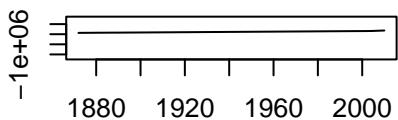
Step 3: Variable Importance (degree=2, interactions)



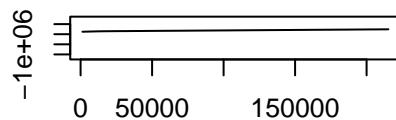
```
par(mfrow=c(2,2)); plotmo(mars3, type="response", nresponse=1, do.par=FALSE); par(mfrow=c(1,1))
```

```
plotmo grid: Bedroom_AbvGr Year_Built Mo_Sold Lot_Area Street Central_Air
            3           1974      6     9480   Pave       Y
First_Flr_SF Second_Flr_SF Full_Bath Half_Bath Fireplaces Garage_Area
          1092           0         2         0         1        480
Gr_Liv_Area TotRms_AbvGrd
          1442           6
```

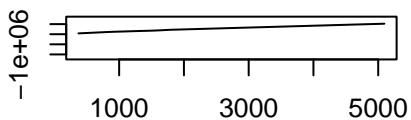
1 Year_Bult



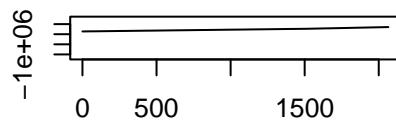
2 Lot_Area



3 Frst_F_SF



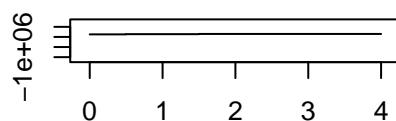
4 Scnd_F_SF



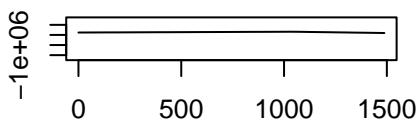
5 Half_Bath



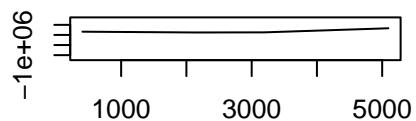
6 Fireplacs




7 Garage_Ar



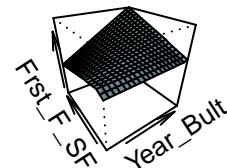
8 Gr_Liv_Ar



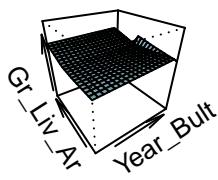
1 Bdrm_AbvG: Scnd_F_SF



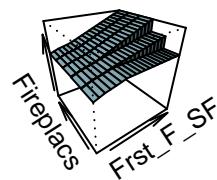
2 Year_Bult: Frst_F_SF



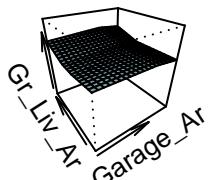
3 Year_Bult: Gr_Liv_Ar



4 Frst_F_SF: Fireplacs



5 Garage_Ar: Gr_Liv_Ar



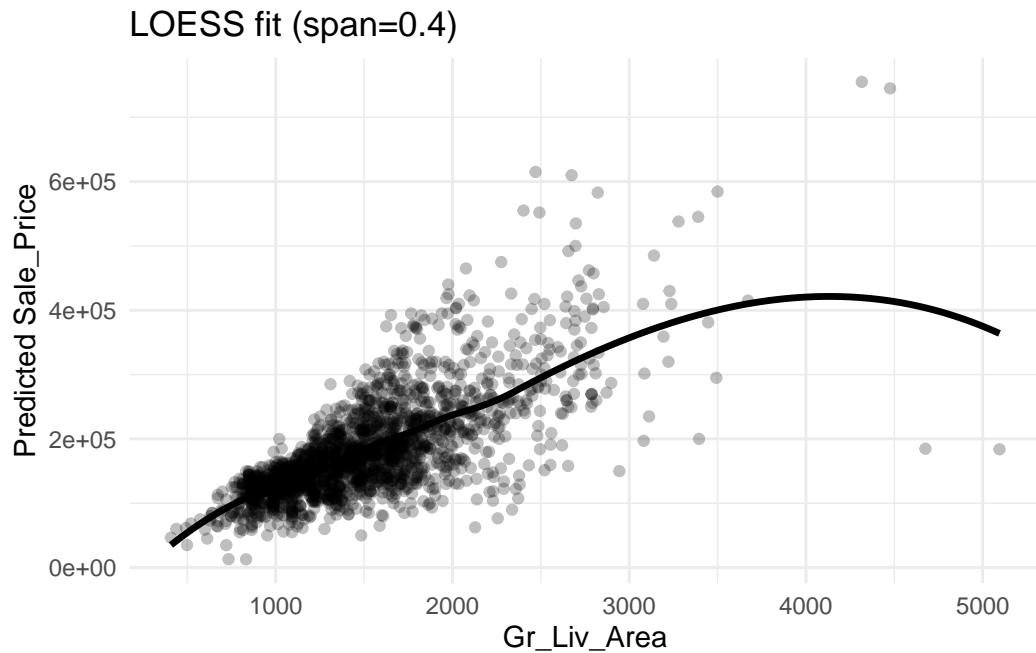
Interpreting MARS terms:

Look for $h()$ pieces and products. Coefficients on $h(x-c)$ indicate how slope changes after the knot c . Interaction products mean “the slope change depends on another variable”.

4.5 5. LOESS (Visual local regression)

```
loess_fit <- loess(Sale_Price ~ Gr_Liv_Area, data=train, span=0.4)
grid_lo <- tibble(Gr_Liv_Area = seq(min(train$Gr_Liv_Area), max(train$Gr_Liv_Area), length.out=100))
grid_lo$pred <- predict(loess_fit, newdata=grid_lo)

ggplot() +
  geom_point(data=train, aes(Gr_Liv_Area, Sale_Price), alpha=.25) +
  geom_line(data=grid_lo, aes(Gr_Liv_Area, pred), linewidth=1.2) +
  labs(title="LOESS fit (span=0.4)", y="Predicted Sale_Price")
```



We keep LOESS as a visualization-oriented smoother (not used in the multivariate comparison below).

4.6 6. GAM Splines with mgcv – Simple → Complex

GAM smooths are written `s(x)` and estimated with penalized splines. The `edf` (effective degrees of freedom) in the summary tells you the smooth's complexity (larger edf = wigglier function).

4.6.1 6A. Univariate GAM: Sale_Price ~ s(Garage_Area)

```
gam1 <- gam(Sale_Price ~ s(Garage_Area, k=9), data=train, method="REML")
summary(gam1) # inspect edf and significance
```

```
Family: gaussian
Link function: identity

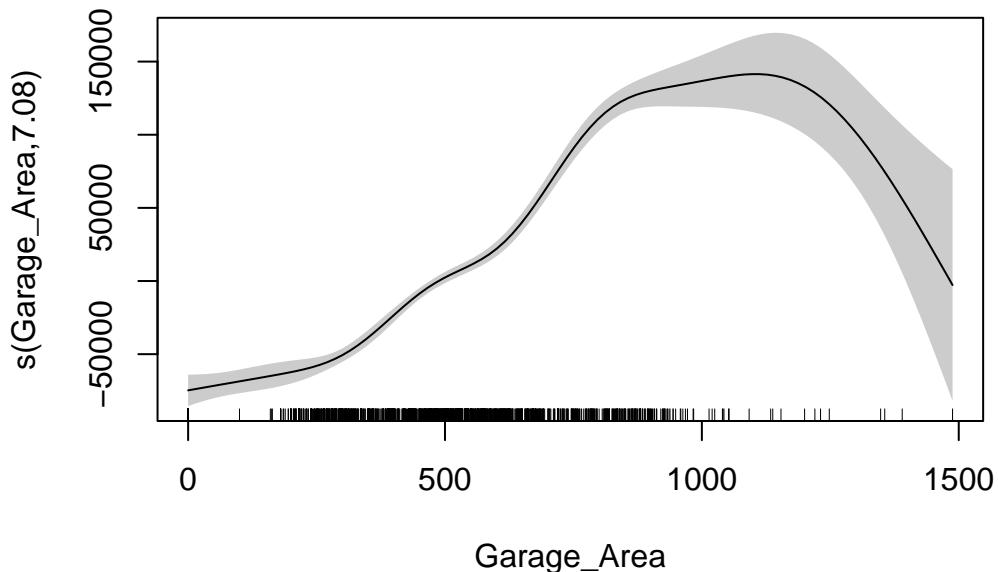
Formula:
Sale_Price ~ s(Garage_Area, k = 9)

Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) 180897        1292      140    <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:
            edf Ref.df     F p-value
s(Garage_Area) 7.078 7.719 217.1 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) =  0.45   Deviance explained = 45.2%
-REML = 25418  Scale est. = 3.4232e+09 n = 2051
```

```
plot(gam1, pages=1, shade=TRUE, scheme=1, scale=0)
```



4.6.2 6B. Two-smooth Additive GAM: $s(\text{Garage_Area}) + s(\text{Gr_Liv_Area})$

```
gam2 <- gam(Sale_Price ~ s(Garage_Area, k=9) + s(Gr_Liv_Area, k=9), data=train, method="REML")
summary(gam2)
```

```

Family: gaussian
Link function: identity

Formula:
Sale_Price ~ s(Garage_Area, k = 9) + s(Gr_Liv_Area, k = 9)

Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) 180897      1026    176.2   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:
            edf Ref.df     F p-value
s(Garage_Area) 6.446 7.293 103.3   <2e-16 ***
```

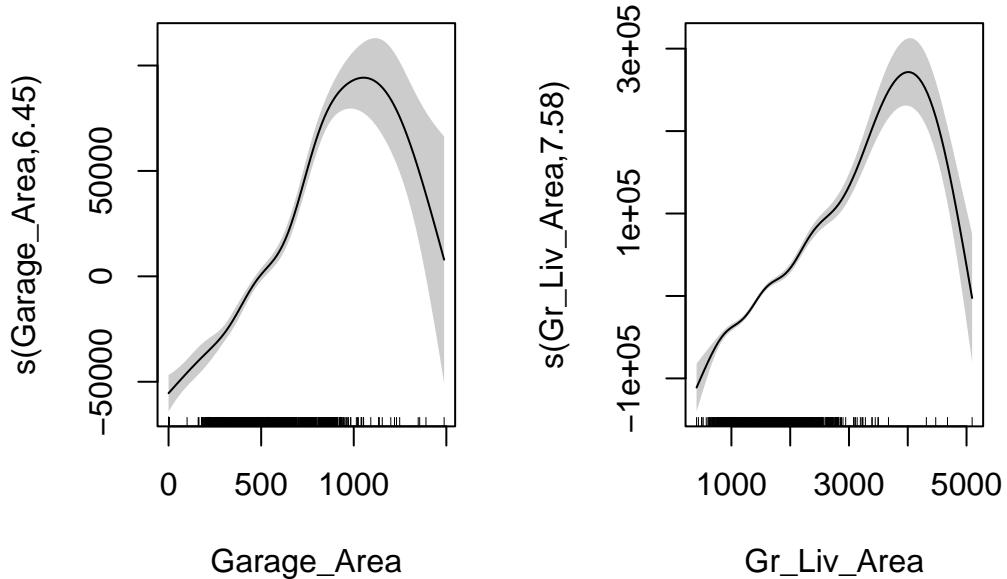
```

s(Gr_Liv_Area) 7.585 7.941 152.6 <2e-16 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.653 Deviance explained = 65.5%
-REML = 24949 Scale est. = 2.1608e+09 n = 2051

```

```
plot(gam2, pages=1, shade=TRUE, scheme=1, scale=0)
```



4.6.3 6C. Full Multivariate GAM (selected smooths + factors)

```

gam3 <- gam(Sale_Price ~
              s(Garage_Area, k=9) +
              s(Gr_Liv_Area, k=9) +
              s(Year_Built, k=7) +
              s(Lot_Area, k=7) +
              Bedroom_AbvGr + Mo_Sold +
              Street + Central_Air +
              First_Flr_SF + Second_Flr_SF +
              Full_Bath + Half_Bath +

```

```

    Fireplaces + TotRms_AbvGrd,
    data=train, method="REML")
summary(gam3)

Family: gaussian
Link function: identity

Formula:
Sale_Price ~ s(Garage_Area, k = 9) + s(Gr_Liv_Area, k = 9) +
  s(Year_Built, k = 7) + s(Lot_Area, k = 7) + Bedroom_AbvGr +
  Mo_Sold + Street + Central_Air + First_Flr_SF + Second_Flr_SF +
  Full_Bath + Half_Bath + Fireplaces + TotRms_AbvGrd

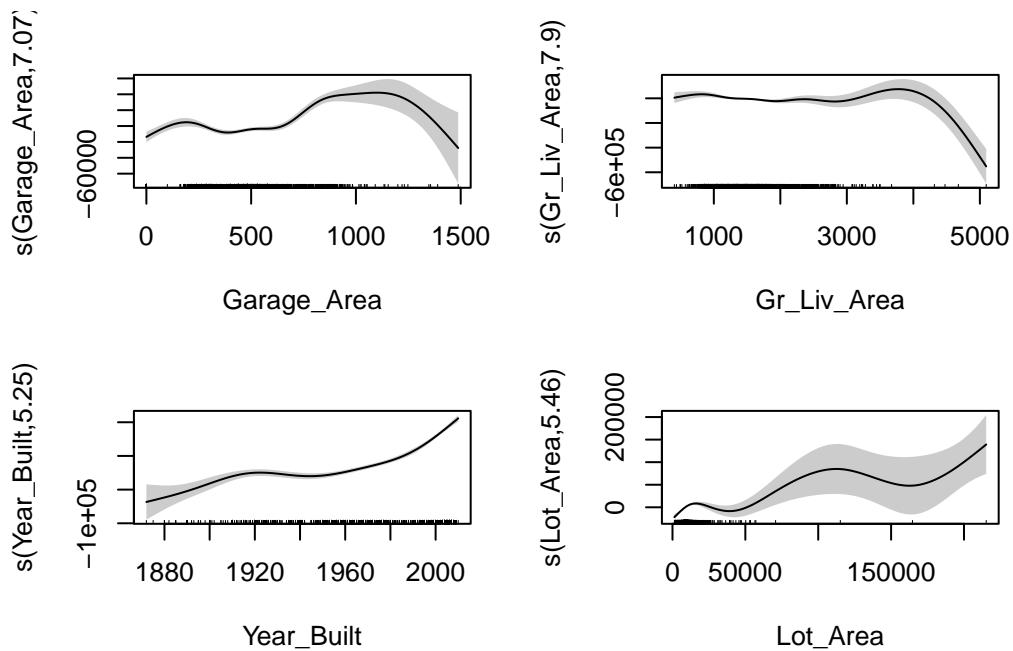
Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -21237.92   30106.84 -0.705  0.4806
Bedroom_AbvGr -10350.20    1382.38 -7.487 1.05e-13 ***
Mo_Sold      -409.12     270.24 -1.514  0.1302
StreetPave    30589.21   13486.61  2.268  0.0234 *
Central_AirY 17189.68    3366.03  5.107 3.59e-07 ***
First_Flr_SF  140.27      17.44  8.044 1.47e-15 ***
Second_Flr_SF 101.93      17.36  5.871 5.06e-09 ***
Full_Bath     -4548.94   2211.69 -2.057  0.0398 *
Half_Bath      2355.98   2175.40  1.083  0.2789
Fireplaces    13066.61   1371.59  9.527 < 2e-16 ***
TotRms_AbvGrd -1640.87    948.67 -1.730  0.0838 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:
        edf Ref.df      F p-value
s(Garage_Area) 7.067 7.712 24.96 <2e-16 ***
s(Gr_Liv_Area) 7.899 7.996 46.44 <2e-16 ***
s(Year_Built)  5.255 5.764 128.05 <2e-16 ***
s(Lot_Area)    5.465 5.883 16.44 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) =  0.827  Deviance explained =  83%
-REML =  24167  Scale est. = 1.0786e+09 n = 2051

```

```
plot(gam3, pages=1, shade=TRUE, scheme=1, scale=0)
```



Interpreting GAM splines:

A smooth $s(x)$ shows how the expected outcome varies with x *holding other terms constant*. The edf indicates complexity; wide confidence bands suggest greater uncertainty in those regions.

4.7 7. Model Comparison (R^2 , RMSE, MSE)

```
# Helper for R^2 on test
rsq <- function(y, yhat) {
  1 - sum((y - yhat)^2) / sum((y - mean(y))^2)
}

# Predictions
p_piece <- predict(m_seg, newdata=test)
p_mars  <- predict(mars3, newdata=test)      # final MARS with interactions
p_gam   <- predict(gam3, newdata=test)        # full multivariate GAM
```

```

# Metrics
tbl <- tibble(
  Model = c("Piecewise (segmented)", "MARS (degree=2)", "GAM splines (mgcv full)"),
  RMSE  = c(rmse(test$Sale_Price, p_piece),
            rmse(test$Sale_Price, p_mars),
            rmse(test$Sale_Price, p_gam)),
  MSE   = c(mse(test$Sale_Price, p_piece),
            mse(test$Sale_Price, p_mars),
            mse(test$Sale_Price, p_gam)),
  R2    = c(rsq(test$Sale_Price, p_piece),
            rsq(test$Sale_Price, p_mars),
            rsq(test$Sale_Price, p_gam))
) |> arrange(desc(R2))

tbl

```

Model	RMSE	MSE	R2
<chr>	<dbl>	<dbl>	<dbl>
1 MARS (degree=2)	33028.	1090840972.	0.839
2 GAM splines (mgcv full)	36796.	1353918715.	0.800
3 Piecewise (segmented)	59034.	3484963067.	0.484


4.8 8. Final Plot: Predicted vs Actual (Test Set)

```

plot_df <- bind_rows(
  tibble(Model="Piecewise (segmented)", Actual=test$Sale_Price, Pred=p_piece),
  tibble(Model="MARS (degree=2)", Actual=test$Sale_Price, Pred=p_mars),
  tibble(Model="GAM splines (mgcv full)", Actual=test$Sale_Price, Pred=p_gam)
)

ggplot(plot_df, aes(Actual, Pred, color=Model)) +
  geom_point(alpha=.35) +
  geom_abline(intercept=0, slope=1, linetype="dashed") +
  labs(title="Predicted vs Actual (Test Set)", y="Predicted Sale_Price") +
  coord_equal()

```



4.9 9. How to Explain

- **Piecewise:** “There’s a threshold where the effect changes.” Simple story, limited flexibility.
- **MARS:** “The model finds breakpoints and sometimes *when* they matter (interactions).” Explain $h(x-c)$ as *extra slope after c*.
- **GAM (splines):** “We model smooth curves for key variables; edf shows wigginess. More edf more flexibility.”
- **Tradeoff:** More flexibility usually improves error (RMSE/MSE) and R^2 , but reduce interpretability. Validate with cross-validation and keep the story aligned with business intuition.

5 GAM— Piecewise, LOESS, and GAM Splines - Python Version

Note: The original MARS section used the `sklearn-contrib-py-earth` package, which is **no longer actively maintained** and may fail to install on modern Python versions.

The MARS code is left **commented out** for reference — you can un-comment it if you successfully install `py-earth` (try `conda install -c conda-forge sklearn-contrib-py-earth`).

Equivalent flexibility can be achieved through **GAM** or **spline-based models**, which are shown below.

5.1 1. Setup and Data

```
import sys
print("Active Python interpreter:", sys.executable)

import numpy as np, pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression, Ridge
import inspect

# --- Fix for SciPy >= 1.13 removing .A from sparse matrices ---
import scipy.sparse as sp

if not hasattr(sp.spmatrix, "A"):
    def _toarray(self):
```

```

        return self.toarray()
    sp.spmatrix.A = property(_toarray)
# -----



np.random.seed(4321)

Active Python interpreter: /opt/anaconda3/bin/python

from sklearn.datasets import fetch_openml

ames = fetch_openml(name="house_prices", as_frame=True).frame

keep = {
    "SalePrice": "SalePrice",
    "BedroomAbvGr": "Bedroom_AbvGr",
    "YearBuilt": "Year_Built",
    "MoSold": "Mo_Sold",
    "LotArea": "Lot_Area",
    "Street": "Street",
    "CentralAir": "Central_Air",
    "1stFlrSF": "First_Flr_SF",
    "2ndFlrSF": "Second_Flr_SF",
    "FullBath": "Full_Bath",
    "HalfBath": "Half_Bath",
    "Fireplaces": "Fireplaces",
    "GarageArea": "Garage_Area",
    "GrLivArea": "Gr_Liv_Area",
    "TotRmsAbvGrd": "TotRms_AbvGrd"
}

df = ames[list(keep.keys())].rename(columns=keep).dropna().copy()

X = df.drop(columns=["SalePrice"])
y = df["SalePrice"].values

num_cols = X.select_dtypes(include=[np.number]).columns.tolist()
cat_cols = [c for c in X.columns if c not in num_cols]

# Handle OneHotEncoder version changes and ensure dense output
if "sparse_output" in inspect.signature(OneHotEncoder).parameters:
    encoder = OneHotEncoder(drop="first", handle_unknown="ignore", sparse_output=False)

```

```

else:
    encoder = OneHotEncoder(drop="first", handle_unknown="ignore", sparse=False)

pre = ColumnTransformer([
    ("num", StandardScaler(), num_cols),
    ("cat", encoder, cat_cols)
], sparse_threshold=0)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=4321)
df.head()

```

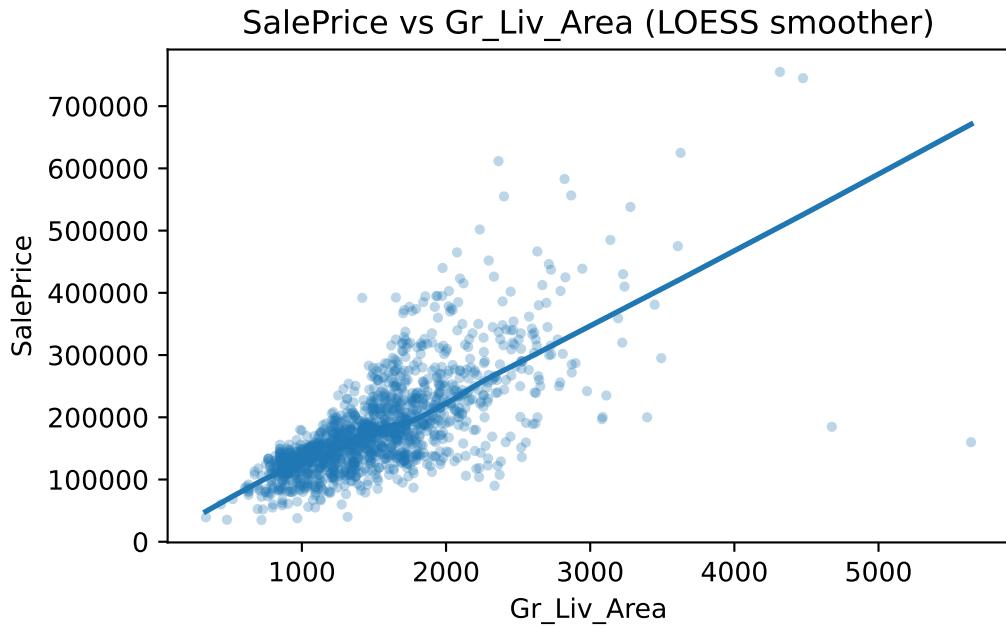
	SalePrice	Bedroom_AbvGr	Year_Built	Mo_Sold	Lot_Area	Street	Central_Air	First_Flr_SF
0	208500	3	2003	2	8450	Pave	Y	856
1	181500	3	1976	5	9600	Pave	Y	1262
2	223500	3	2001	9	11250	Pave	Y	920
3	140000	3	1915	2	9550	Pave	Y	961
4	250000	4	2000	12	14260	Pave	Y	1145

5.2 2. Visual Check for Nonlinearity

```

import statsmodels.api as sm
fig = plt.figure()
plt.scatter(df["Gr_Liv_Area"], df["SalePrice"], s=8, alpha=0.3)
low = sm.nonparametric.lowess(df["SalePrice"], df["Gr_Liv_Area"], frac=0.3, return_sorted=True)
plt.plot(low[:,0], low[:,1], linewidth=2)
plt.title("SalePrice vs Gr_Liv_Area (LOESS smoother)")
plt.xlabel("Gr_Liv_Area"); plt.ylabel("SalePrice")
plt.tight_layout()

```



5.3 3. Piecewise Regression Example

```

class PiecewiseLinear:
    def __init__(self, knot=2000.0):
        self.knot = knot
        self.lin = LinearRegression()

    def _transform(self, x):
        x1 = np.asarray(x).reshape(-1,1)
        hx = np.maximum(0, x1 - self.knot)
        return np.hstack([x1, hx])

    def fit(self, x, y):
        Z = self._transform(x)
        self.lin.fit(Z, y)
        return self

    def predict(self, x):
        Z = self._transform(x)

```

```

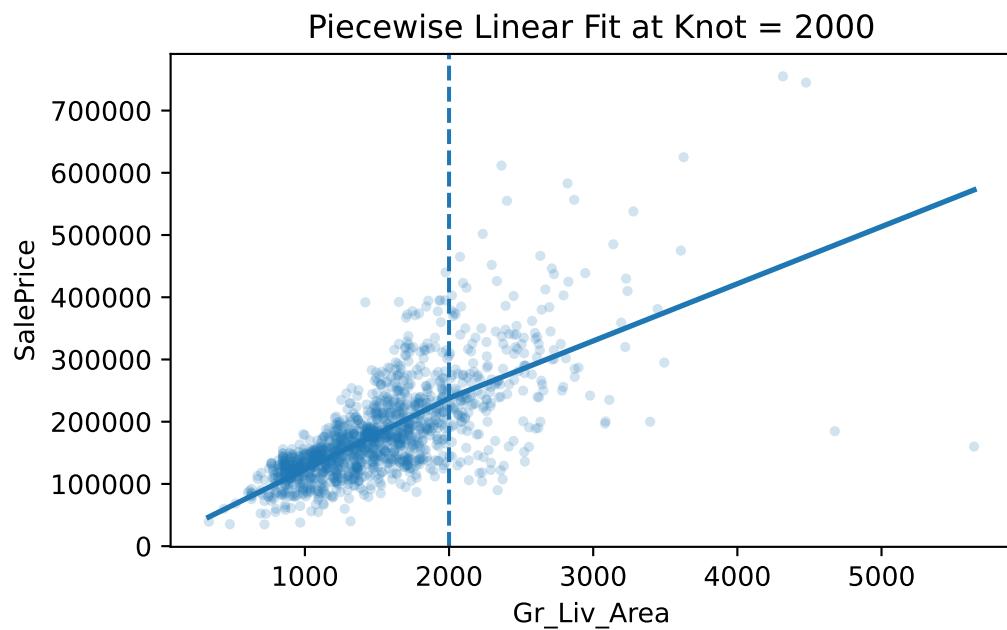
        return self.lin.predict(Z)

pw = PiecewiseLinear(knot=2000).fit(df["Gr_Liv_Area"].values, df["SalePrice"].values)

grid = np.linspace(df["Gr_Liv_Area"].min(), df["Gr_Liv_Area"].max(), 200)
pred = pw.predict(grid)

plt.figure()
plt.scatter(df["Gr_Liv_Area"], df["SalePrice"], s=8, alpha=0.2)
plt.plot(grid, pred, linewidth=2)
plt.axvline(2000, linestyle="--")
plt.title("Piecewise Linear Fit at Knot = 2000")
plt.xlabel("Gr_Liv_Area"); plt.ylabel("SalePrice")
plt.tight_layout()

```



5.4 4. MARS (Commented Out — For Reference Only)

```

# The MARS implementation (py-earth) is not actively maintained and may fail to install on m
# If you wish to try it, install from conda-forge:
#     conda install -c conda-forge sklearn-contrib-py-earth
#
# Example (commented out):
#
# from pyearth import Earth
#
# mars = Pipeline([("prep", pre), ("model", Earth(max_degree=2))]).fit(X_train, y_train)
# print(mars.named_steps["model"].summary())
#
# xs = np.linspace(X["Garage_Area"].min(), X["Garage_Area"].max(), 200)
# yhat = mars.predict(pd.DataFrame({"Garage_Area": xs}))
#
# plt.figure()
# plt.scatter(X_train["Garage_Area"], y_train, s=8, alpha=0.2)
# plt.plot(xs, yhat, linewidth=2)
# plt.title("MARS with Garage_Area (Commented Out Example)")
# plt.xlabel("Garage_Area"); plt.ylabel("SalePrice")
# plt.tight_layout()

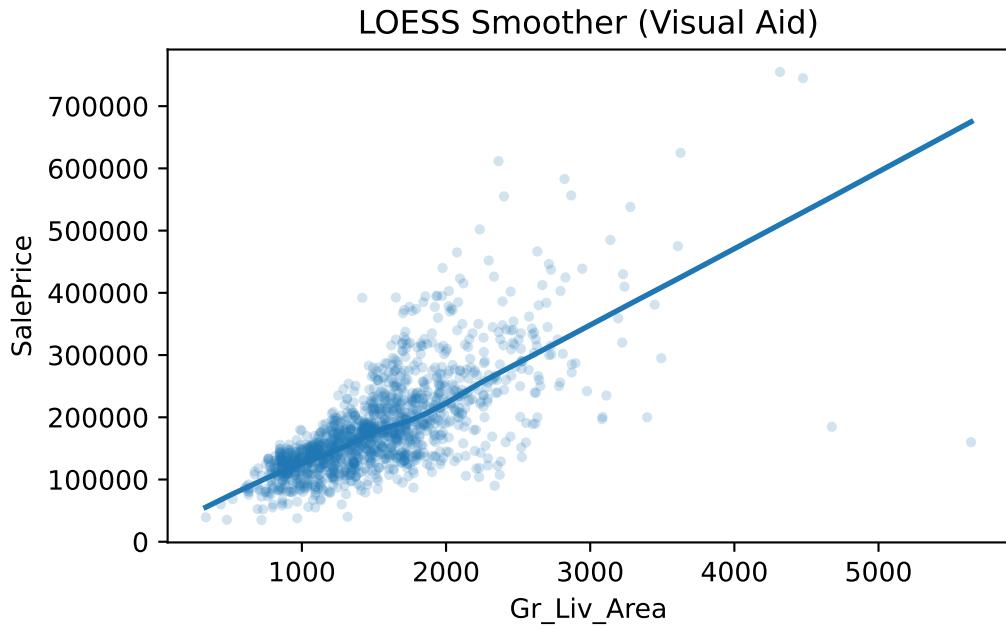
```

5.5 5. LOESS Visualization

```

fig = plt.figure()
plt.scatter(df["Gr_Liv_Area"], df["SalePrice"], s=8, alpha=0.2)
low = sm.nonparametric.lowess(df["SalePrice"], df["Gr_Liv_Area"], frac=0.4, return_sorted=True)
plt.plot(low[:,0], low[:,1], linewidth=2)
plt.title("LOESS Smoother (Visual Aid)")
plt.xlabel("Gr_Liv_Area"); plt.ylabel("SalePrice")
plt.tight_layout()

```



5.6 6. GAM and Spline Approaches (Modern Alternatives)

```
from pygam import LinearGAM, s
gam1 = LinearGAM(s(0)).fit(X_train[["Garage_Area"]].values, y_train)
print(gam1.summary())
gam2 = LinearGAM(s(0) + s(1)).fit(X_train[["Garage_Area", "Gr_Liv_Area"]].values, y_train)
print(gam2.summary())
num_train = X_train[num_cols].values
gam3 = LinearGAM().fit(num_train, y_train)
print(gam3.summary())
```

```
LinearGAM
=====
Distribution: NormalDist Effective DoF:
Link Function: IdentityLink Log Likelihood:
23263.1874
Number of Samples: 1022 AIC:
AICc:
```

Feature Function	Lambda	Rank	EDoF	P > x
s(0)	[0.6]	20	11.7	1.11e-
16 ***				
intercept		1	0.0	1.11e-
16 ***				
Significance codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1				
WARNING: Fitting splines and a linear function to a feature introduces a model identifiability problem which can cause p-values to appear significant when they are not.				
WARNING: p-values calculated in this manner behave correctly for un-penalized models or models with known smoothing parameters, but when smoothing parameters have been estimated, the p-values are typically lower than they should be, meaning that the tests reject the null too easily.				
None				
LinearGAM				
Distribution:	NormalDist	Effective DoF:		
Link Function:	IdentityLink	Log Likelihood:		
22770.8323				
Number of Samples:	1022	AIC:		
		AICc:		
		GCV:		
		Scale:		
		Pseudo R-Squared:		
Feature Function	Lambda	Rank	EDoF	P > x
s(0)	[0.6]	20	12.6	1.11e-
16 ***				
s(1)	[0.6]	20	8.3	1.11e-
16 ***				
intercept		1	0.0	1.11e-
16 ***				
Significance codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1				

WARNING: Fitting splines and a linear function to a feature introduces a model identifiability problem which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models with known smoothing parameters, but when smoothing parameters have been estimated, the p-values

are typically lower than they should be, meaning that the tests reject the null too often.

None

LinearGAM

```
=====
Distribution: NormalDist Effective DoF:
Link Function: IdentityLink Log Likelihood:
21995.395
Number of Samples: 1022 AIC:
AICc:
GCV:
Scale:
Pseudo R-Squared:
=====
```

Feature Function	Lambda	Rank	EDoF	P > x
s(0)	[0.6]	20	8.7	6.58e-
07 ***				
s(1)	[0.6]	20	12.8	1.11e-
16 ***				
s(2)	[0.6]	20	11.0	5.15e-
01				
s(3)	[0.6]	20	8.4	1.11e-
16 ***				
s(4)	[0.6]	20	8.6	1.11e-
16 ***				
s(5)	[0.6]	20	10.1	2.17e-
03 **				
s(6)	[0.6]	20	3.2	9.76e-
02 .				
s(7)	[0.6]	20	2.2	4.19e-
01				
s(8)	[0.6]	20	2.9	3.16e-
10 ***				
s(9)	[0.6]	20	9.0	1.11e-
16 ***				
s(10)	[0.6]	20	5.7	1.11e-
16 ***				

```

s(11)                      [0.6]          20      6.5      3.78e-
01
intercept                   1      0.0      1.11e-
16    ***
=====
Significance codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

WARNING: Fitting splines and a linear function to a feature introduces a model identifiability
         which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models
         with known smoothing parameters, but when smoothing parameters have been estimated, the p-
         values
         are typically lower than they should be, meaning that the tests reject the null too
         None

```

5.7 7. Regression Spline Example (Fallback)

```

from patsy import dmatrix

design_tr = dmatrix("bs(Garage_Area, df=6) + bs(Gr_Liv_Area, df=6)", data=X_train, return_type='dataframe')
design_te = dmatrix("bs(Garage_Area, df=6) + bs(Gr_Liv_Area, df=6)", data=X_test, return_type='dataframe')

ridge = Ridge(alpha=1.0).fit(design_tr, y_train)
pred_spline = ridge.predict(design_te)

```

5.8 8. Model Comparison (Piecewise vs GAM vs Spline)

```

def rmse(y, yhat):
    try:
        return mean_squared_error(y, yhat, squared=False)
    except TypeError:
        return np.sqrt(mean_squared_error(y, yhat))

```

```

def mse(y, yhat): return mean_squared_error(y, yhat)
def rsq(y, yhat): return r2_score(y, yhat)

p_piece = pw.predict(X_test["Gr_Liv_Area"].values)

try:
    p_gam = gam3.predict(X_test[num_cols].values)
except Exception:
    p_gam = np.full_like(y_test, np.nan, dtype=float)

p_spline = pred_spline

cmp = pd.DataFrame({
    "Model": ["Piecewise", "GAM (pyGAM)" if not np.isnan(p_gam).all() else "GAM (not available)"],
    "RMSE": [rmse(y_test, p_piece), rmse(y_test, p_gam) if not np.isnan(p_gam).all() else np.nan],
    "MSE": [mse(y_test, p_piece), mse(y_test, p_gam) if not np.isnan(p_gam).all() else np.nan],
    "R2": [rsq(y_test, p_piece), rsq(y_test, p_gam) if not np.isnan(p_gam).all() else np.nan]
}).sort_values("R2", ascending=False)

cmp

```

	Model	RMSE	MSE	R2
1	GAM (pyGAM)	40805.867162	1.665119e+09	0.766556
0	Piecewise	58108.537553	3.376602e+09	0.526612
2	Regression Spline	60289.000524	3.634764e+09	0.490419

5.9 9. Predicted vs Actual Comparison

```

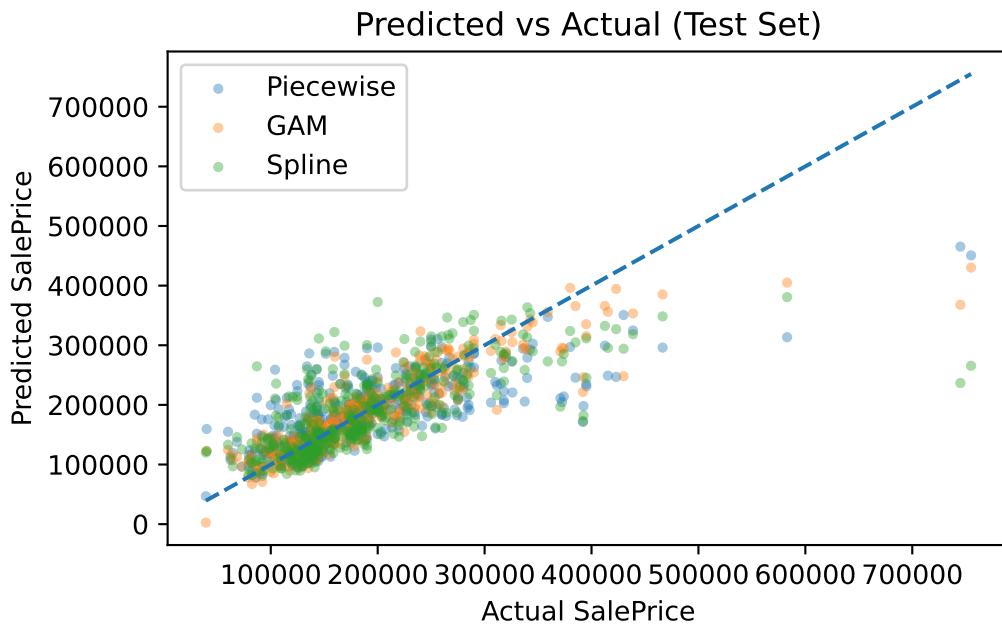
plt.figure()
plt.scatter(y_test, p_piece, s=8, alpha=0.4, label="Piecewise")
if not np.isnan(p_gam).all():
    plt.scatter(y_test, p_gam, s=8, alpha=0.4, label="GAM")
plt.scatter(y_test, p_spline, s=8, alpha=0.4, label="Spline")
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], linestyle="--")
plt.legend()

```

```

plt.title("Predicted vs Actual (Test Set)")
plt.xlabel("Actual SalePrice"); plt.ylabel("Predicted SalePrice")
plt.tight_layout()

```



5.10 10. Interpretation Summary

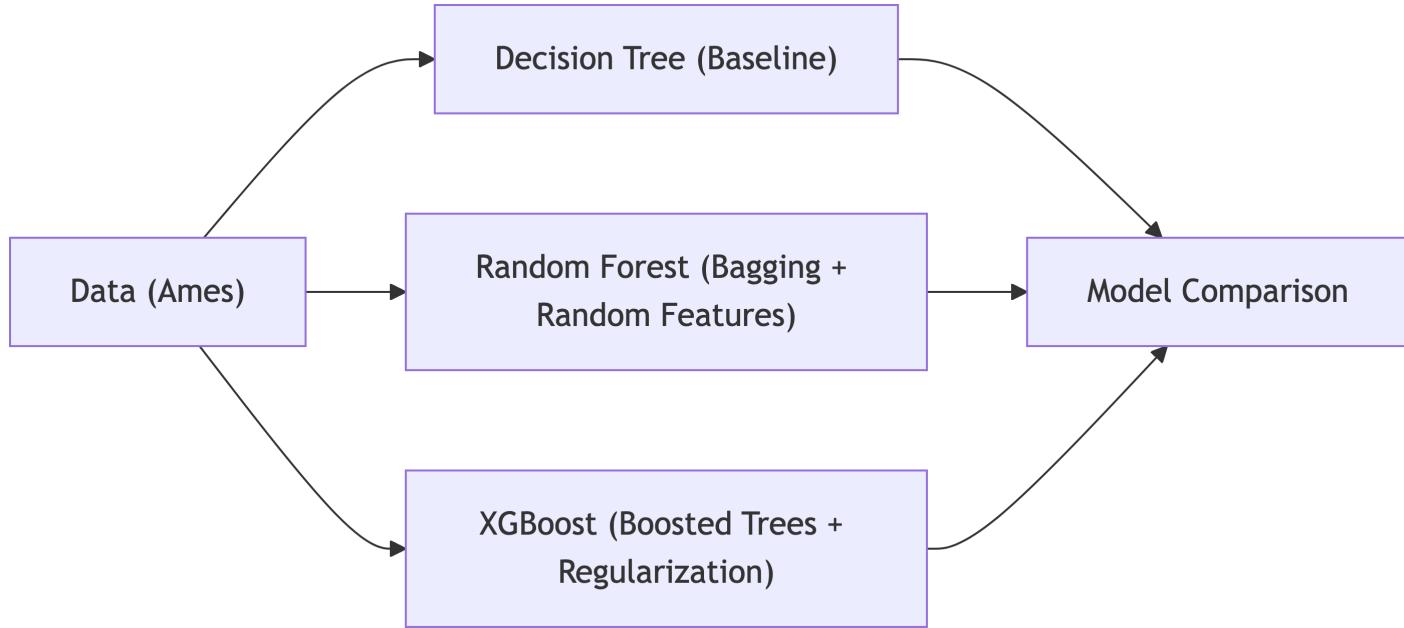
- **Piecewise:** adds simple thresholds, easy to explain.
- **GAMs:** provide smooth, interpretable nonlinearities.
- **Splines:** flexible approximations that behave like local polynomials.
- **MARS (optional):** similar conceptually but less maintained in Python; consider using R's `earth` instead.

6 Tree Based Models - R Version

Comparison: **Decision Tree**, **Random Forest**, and **XGBoost** on the Ames dataset. Includes random-variable baselines, variable importance, partial dependence plots.

6.1 Visual Roadmap

```
graph TD
    A["Data (Ames)"] --> DT["Decision Tree (Baseline)"]
    A --> RF["Random Forest (Bagging + Random Features)"]
    A --> XGB["XGBoost (Boosted Trees + Regularization)"]
    DT --> Compare["Model Comparison"]
    RF --> Compare
    XGB --> Compare
```



6.2 Libraries, Data, Split

```

library(tidyverse)
library(caret)
library(rpart)
library(rpart.plot)
library(randomForest)
library(xgboost)
library(Ckmeans.1d.dp)
library(pdp)
library(AmesHousing)

ames <- make_ordinal_ames()
ames <- ames %>% mutate(id = row_number())

set.seed(4321)
training <- ames %>% sample_frac(0.7)
testing <- anti_join(ames, training, by = 'id')

training <- training %>%
  select(Sale_Price,

```

```
Bedroom_AbvGr, Year_Built, Mo_Sold, Lot_Area, Street, Central_Air,
First_Flr_SF, Second_Flr_SF, Full_Bath, Half_Bath, Fireplaces,
Garage_Area, Gr_Liv_Area, TotRms_AbvGrd)
```

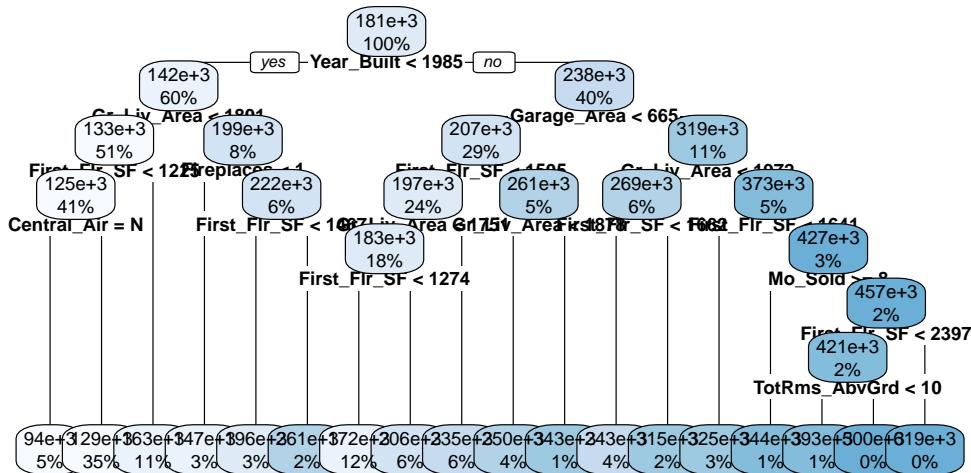
```
training.df <- as.data.frame(training)
```

6.3 Decision Tree (Baseline)

```
set.seed(12345)
tree.ames <- rpart(Sale_Price ~ ., data = training.df, method = "anova",
                     control = rpart.control(cp = 0.005))

rpart.plot(tree.ames, type = 2, fallen.leaves = TRUE, cex = 0.6,
           main = "Decision Tree for Ames Housing")
```

Decision Tree for Ames Housing



```
testing_sel <- testing %>% select(names(training.df))
pred_tree <- predict(tree.ames, newdata = testing_sel)

tree_perf <- testing_sel %>%
```

```

    mutate(pred = pred_tree,
           AE = abs(Sale_Price - pred),
           APE = 100*abs((Sale_Price - pred)/Sale_Price)) %>%
  summarise(Model = "Decision Tree",
            RMSE = sqrt(mean((Sale_Price - pred)^2)),
            MAE = mean(AE),
            MAPE = mean(APE),
            R2 = 1 - sum((Sale_Price - pred)^2)/sum((Sale_Price - mean(Sale_Price))^2))
tree_perf

```

```

# A tibble: 1 x 5
  Model      RMSE     MAE   MAPE     R2
  <chr>     <dbl>   <dbl> <dbl> <dbl>
1 Decision Tree 42938. 28322. 17.5 0.727

```

6.4 Random Forest: Random Variable, PDPs, and Importance

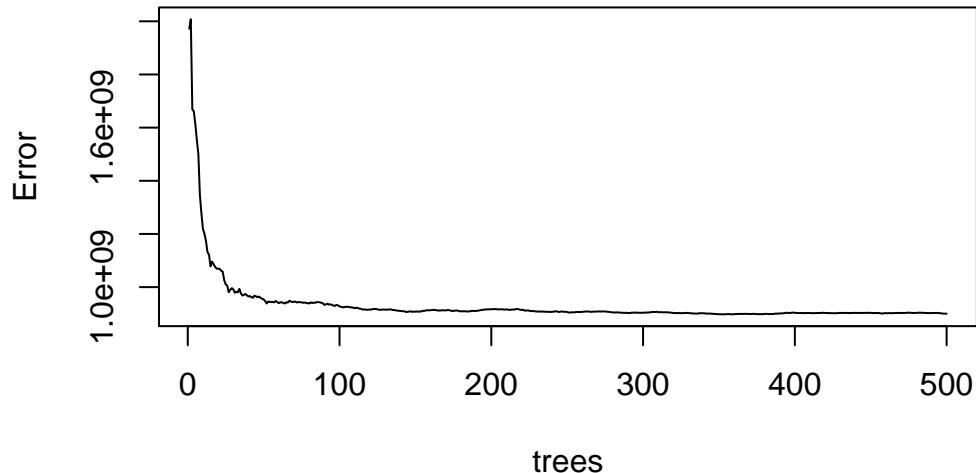
```

# Add random variable for feature selection sanity check
training.df$random <- rnorm(nrow(training.df))

set.seed(12345)
rf.ames <- randomForest(Sale_Price ~ ., data = training.df, ntree = 500, mtry = 4, importance = TRUE)
plot(rf.ames, main = "Random Forest: Error vs Trees")


```

Random Forest: Error vs Trees



```
varImpPlot(rf.ames, sort = TRUE, n.var = 15, main = "Random Forest: Variables vs Random Baseline")
```


Random Forest: Variables vs Random Baseline




```
importance(rf.ames)
```

	%IncMSE	IncNodePurity
Bedroom_AbvGr	16.7830622	2.081292e+11
Year_Built	59.6827740	2.608576e+12
Mo_Sold	0.6717216	1.820404e+11
Lot_Area	16.2410098	5.818823e+11
Street	4.6451725	5.141533e+09
Central_Air	15.9709319	1.023533e+11
First_Flr_SF	37.4760827	1.884177e+12
Second_Flr_SF	24.6324055	6.159270e+11
Full_Bath	12.3716046	6.191162e+11
Half_Bath	13.6183708	1.153802e+11
Fireplaces	26.1893181	5.796209e+11
Garage_Area	28.4006898	1.961373e+12
Gr_Liv_Area	35.1279551	2.445543e+12
TotRms_AbvGrd	13.9865453	4.189078e+11
random	1.2436860	1.937137e+11

```
# Partial dependence for key variables
partialPlot(rf.ames, training.df, Year_Built, main = "Partial Dependence: Year_Built (RF)")
```



```
partialPlot(rf.ames, training.df, Garage_Area, main = "Partial Dependence: Garage_Area (RF)")
```



6.4.1 Random Forest Results

```
testing_sel$random <- rnorm(nrow(testing_sel))
testing_sel$pred_rf <- predict(rf.ames, testing_sel)

rf_perf <- testing_sel %>%
  mutate(AE = abs(Sale_Price - pred_rf),
        APE = 100*abs((Sale_Price - pred_rf)/Sale_Price)) %>%
  summarise(Model = "Random Forest",
            RMSE = sqrt(mean((Sale_Price - pred_rf)^2)),
            MAE = mean(AE),
            MAPE = mean(APE),
            R2 = 1 - sum((Sale_Price - pred_rf)^2)/sum((Sale_Price - mean(Sale_Price))^2))
rf_perf
```



```
# A tibble: 1 x 5
  Model          RMSE     MAE   MAPE     R2
  <chr>        <dbl>   <dbl> <dbl> <dbl>
1 Random Forest 30147. 19223. 11.6  0.866
```

6.5 XGBoost: Random Variable, CV, Tuning, PDPs

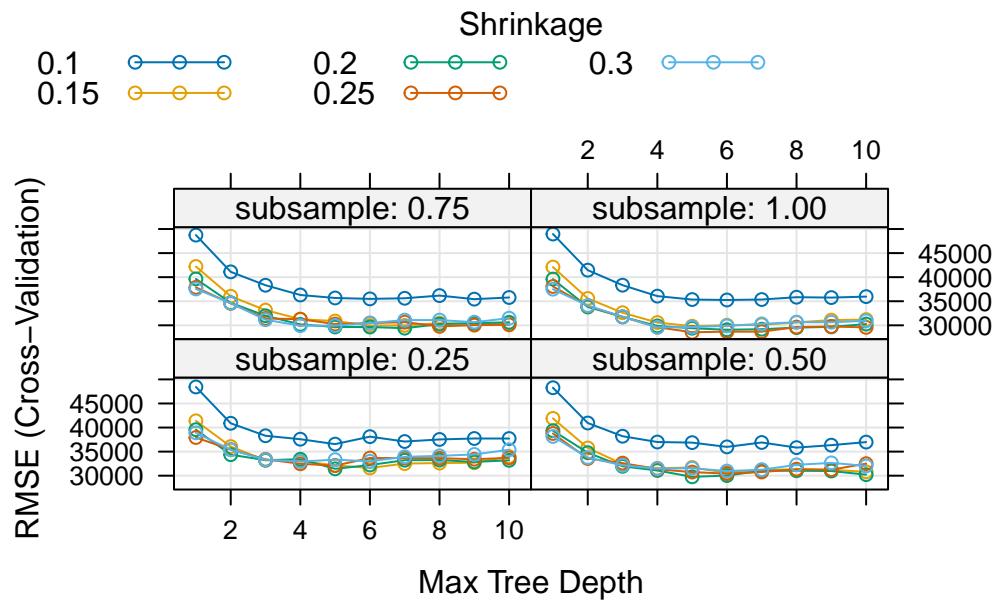
```
training$random <- rnorm(nrow(training))
train_x <- model.matrix(Sale_Price ~ ., data = training)[, -1]
train_y <- training$Sale_Price

set.seed(12345)
xgb.ames <- xgboost(
  data = train_x, label = train_y, subsample = 0.5, nrounds = 100,
  objective = "reg:squarederror", verbose = 0
)

xgbcv.ames <- xgb.cv(
  data = train_x, label = train_y, subsample = 0.5, nrounds = 100,
  nfold = 10, objective = "reg:squarederror", verbose = 0
)

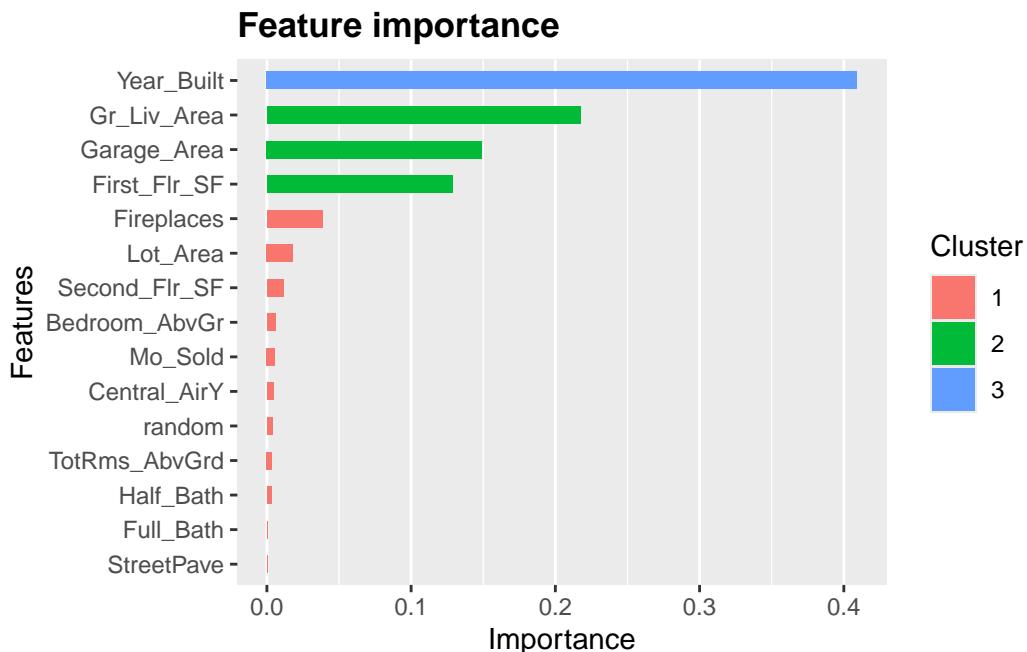
tune_grid <- expand.grid(
  nrounds = 24,
  eta = c(0.1, 0.15, 0.2, 0.25, 0.3),
  max_depth = c(1:10),
  gamma = c(0),
  colsample_bytree = 1,
  min_child_weight = 1,
  subsample = c(0.25, 0.5, 0.75, 1)
)

set.seed(12345)
xgb.ames.caret <- train(
  x = train_x, y = train_y,
  method = "xgbTree",
  tuneGrid = tune_grid,
  trControl = trainControl(method = 'cv', number = 10)
)
plot(xgb.ames.caret)
```



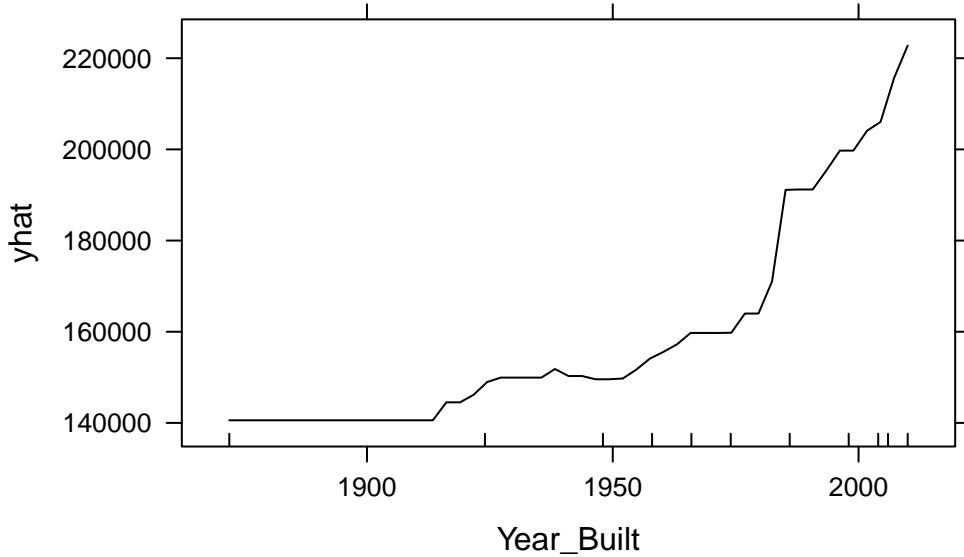
```
# Fit tuned model
xgb.ames <- xgboost(
  data = train_x, label = train_y, subsample = 1,
  nrounds = 24, eta = 0.25, max_depth = 5,
  objective = "reg:squarederror", verbose = 0
)

# Variable importance including random variable
xgb.ggplot.importance(xgb.importance(feature_names = colnames(train_x), model = xgb.ames))
```

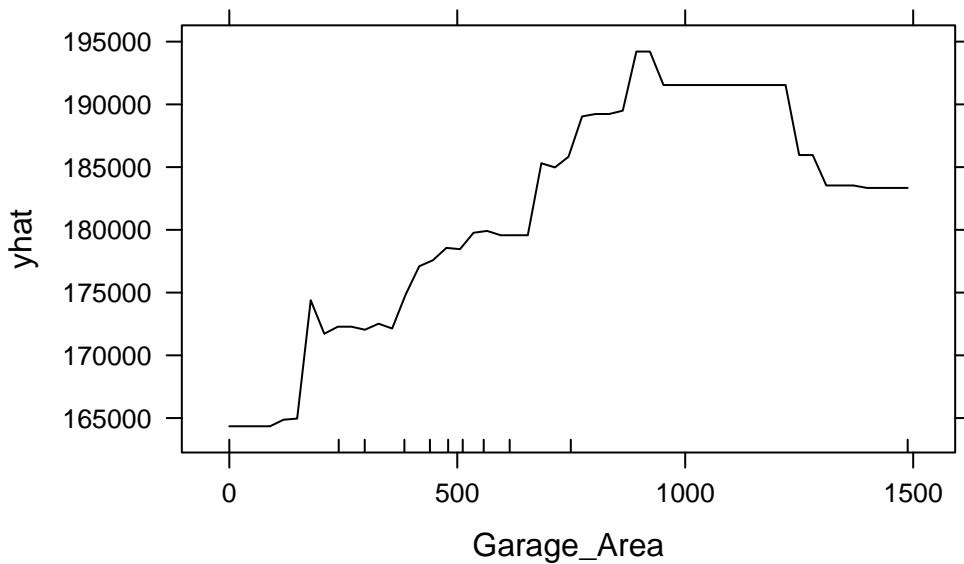


```
# Partial dependence plots
train_df <- as.data.frame(train_x)

partial(xgb.ames, pred.var = "Year_Built",
        plot = TRUE, rug = TRUE, alpha = 0.1, plot.engine = "lattice",
        train = train_df, pdp.color = "red")
```



```
partial(xgb.ames, pred.var = "Garage_Area",
       plot = TRUE, rug = TRUE, alpha = 0.1, plot.engine = "lattice",
       train = train_df)
```



6.5.1 XGBoost Results (Aligned Columns)

```
testing$random <- rnorm(nrow(testing))
testing_x <- model.matrix(Sale_Price ~ ., data = testing)[, -1]

# Align test columns to training columns
missing_cols <- setdiff(colnames(train_x), colnames(testing_x))
for (col in missing_cols) {
  testing_x <- cbind(testing_x, 0)
  colnames(testing_x)[ncol(testing_x)] <- col
}
testing_x <- testing_x[, colnames(train_x)]

testing_sel$pred_xgb <- predict(xgb.ames, testing_x)

xgb_perf <- testing_sel %>%
  mutate(AE = abs(Sale_Price - pred_xgb),
        APE = 100*abs((Sale_Price - pred_xgb)/Sale_Price)) %>%
  summarise(Model = "XGBoost",
            RMSE = sqrt(mean((Sale_Price - pred_xgb)^2)),
            MAE = mean(AE),
            MAPE = mean(APE),
            R2 = 1 - sum((Sale_Price - pred_xgb)^2)/sum((Sale_Price - mean(Sale_Price))^2))
xgb_perf
```

```
# A tibble: 1 x 5
  Model      RMSE     MAE   MAPE     R2
  <chr>    <dbl>  <dbl> <dbl> <dbl>
1 XGBoost 30355. 19820. 11.8 0.864
```

6.6 Final Model Comparison

```
bind_rows(tree_perf, rf_perf, xgb_perf) %>%
  arrange(RMSE)
```

```
# A tibble: 3 x 5
  Model      RMSE     MAE   MAPE     R2
  <chr>    <dbl>  <dbl> <dbl> <dbl>
```

1	Random Forest	30147.	19223.	11.6	0.866
2	XGBoost	30355.	19820.	11.8	0.864
3	Decision Tree	42938.	28322.	17.5	0.727

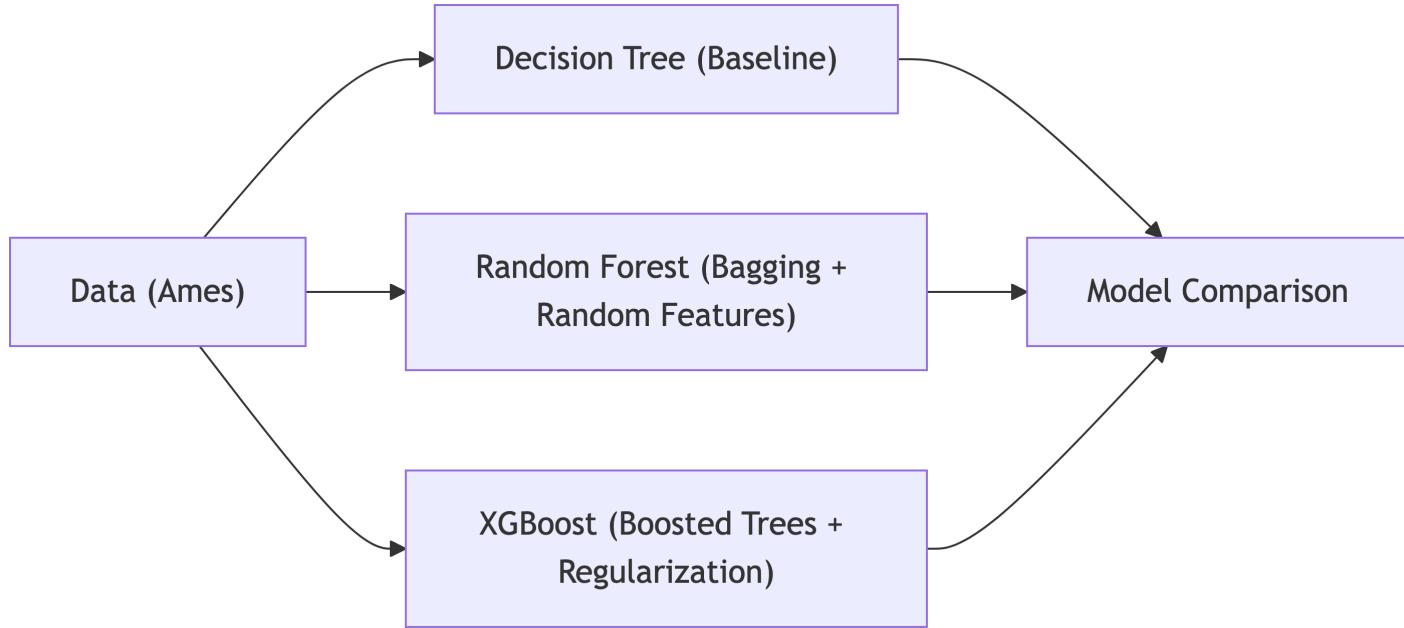
Summary: The Decision Tree provides interpretability, Random Forest and XG-Boost attempt to further reduce bias and manage variance error through randomization and ensembles of weak learners. The issue is that they become less and less interpretable, so we have to introduce other tools to understand what variables are important and how they are relating to the outcome.

7 Tree Based Models - Python Version

Tree Models: **Decision Tree**, **Random Forest**, and **XGBoost** on the Ames Housing dataset. Includes random-variable baselines, variable importance, partial dependence plots, and hyperparameter tuning (random search).

7.1 Visual Roadmap

```
graph TD; A["Data (Ames)"] --> DT["Decision Tree (Baseline)"]; A --> RF["Random Forest (Bagging + Random Features)"]; A --> XGB["XGBoost (Boosted Trees + Regularization)"]; DT --> Compare["Model Comparison"]; RF --> Compare; XGB --> Compare;
```



7.2 Setup and Data Preparation

```

import pandas as pd
import numpy as np
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split, RandomizedSearchCV, KFold
from sklearn.tree import DecisionTreeRegressor, plot_tree
from sklearn.ensemble import RandomForestRegressor
from sklearn.inspection import PartialDependenceDisplay
from xgboost import XGBRegressor, plot_importance
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

# Set random seed for reproducibility
seed = 12345
np.random.seed(seed)

# Load Ames dataset
ames = fetch_openml(name="house_prices", as_frame=True)

```

```

df = ames.frame

# Feature selection matching R version
features = [
    'BedroomAbvGr', 'YearBuilt', 'MoSold', 'LotArea', 'Street', 'CentralAir',
    '1stFlrSF', '2ndFlrSF', 'FullBath', 'HalfBath', 'Fireplaces',
    'GarageArea', 'GrLivArea', 'TotRmsAbvGrd'
]
df = df[features + ['SalePrice']].dropna()

# Encode categoricals
df = pd.get_dummies(df, drop_first=True)

# Train/test split
train, test = train_test_split(df, test_size=0.3, random_state=seed)

# Add random variable for variable-importance baseline
train['random'] = np.random.randn(len(train))
test['random'] = np.random.randn(len(test))

X_train = train.drop(columns=['SalePrice'])
y_train = train['SalePrice']
X_test = test.drop(columns=['SalePrice'])
y_test = test['SalePrice']

```

7.3 Decision Tree (Baseline)

```

tree = DecisionTreeRegressor(random_state=seed, max_depth=5)
tree.fit(X_train, y_train)

plt.figure(figsize=(16,8))
plot_tree(tree, feature_names=X_train.columns, filled=True, max_depth=3)
plt.title("Decision Tree for Ames Housing")
plt.show()

# Metrics
pred_tree = tree.predict(X_test)
rmse_tree = mean_squared_error(y_test, pred_tree)
mae_tree = mean_absolute_error(y_test, pred_tree)


```

```

mape_tree = np.mean(np.abs((y_test - pred_tree) / y_test)) * 100
r2_tree = r2_score(y_test, pred_tree)

pd.DataFrame([["Decision Tree", rmse_tree, mae_tree, mape_tree, r2_tree]],
             columns=["Model", "RMSE", "MAE", "MAPE", "R2"])

```



	Model	RMSE	MAE	MAPE	R2
0	Decision Tree	2.378377e+09	28465.348099	17.701235	0.522975

7.4 Random Forest: Random Variable, Importance, PDPs, and Tuning

```

rf = RandomForestRegressor(random_state=seed)

# Define parameter grid (similar to R tuning)
param_dist_rf = {
    'n_estimators': [100, 200, 300, 400, 500],
    'max_features': ['sqrt', 'log2', None, 4, 6, 8],
    'max_depth': [None, 5, 10, 15, 20],
    'min_samples_split': [2, 5, 10]
}

```

```

}

cv = KFold(n_splits=10, shuffle=True, random_state=seed)

rf_random = RandomizedSearchCV(
    estimator=rf,
    param_distributions=param_dist_rf,
    n_iter=20,
    cv=cv,
    random_state=seed,
    n_jobs=-1,
    scoring='neg_root_mean_squared_error'
)

rf_random.fit(X_train, y_train)
rf_best = rf_random.best_estimator_
print("Best Random Forest Parameters:", rf_random.best_params_)

# Importance with random variable
importances = pd.Series(rf_best.feature_importances_, index=X_train.columns).sort_values(ascending=False)
importances[:15].plot(kind='barh', title='Random Forest: Top Variables (with Random Baseline)')
plt.show()

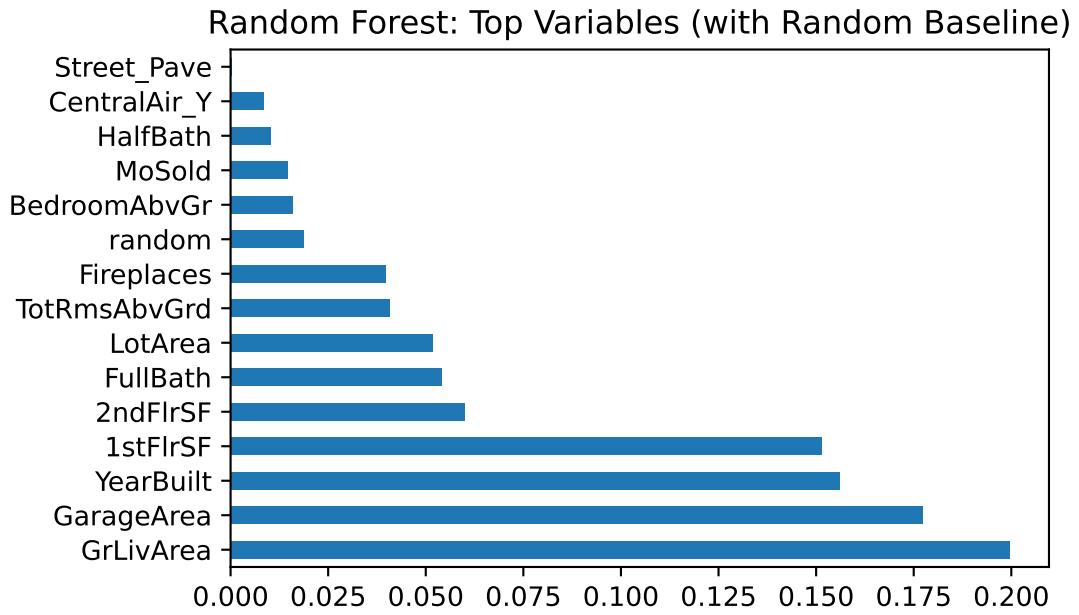
# Partial dependence plots
plt.figure()
PartialDependenceDisplay.from_estimator(rf_best, X_train, ['YearBuilt', 'GarageArea'])
plt.suptitle("Partial Dependence (Random Forest)")
plt.show()

# Performance metrics
pred_rf = rf_best.predict(X_test)
rmse_rf = mean_squared_error(y_test, pred_rf)
mae_rf = mean_absolute_error(y_test, pred_rf)
mape_rf = np.mean(np.abs((y_test - pred_rf) / y_test)) * 100
r2_rf = r2_score(y_test, pred_rf)

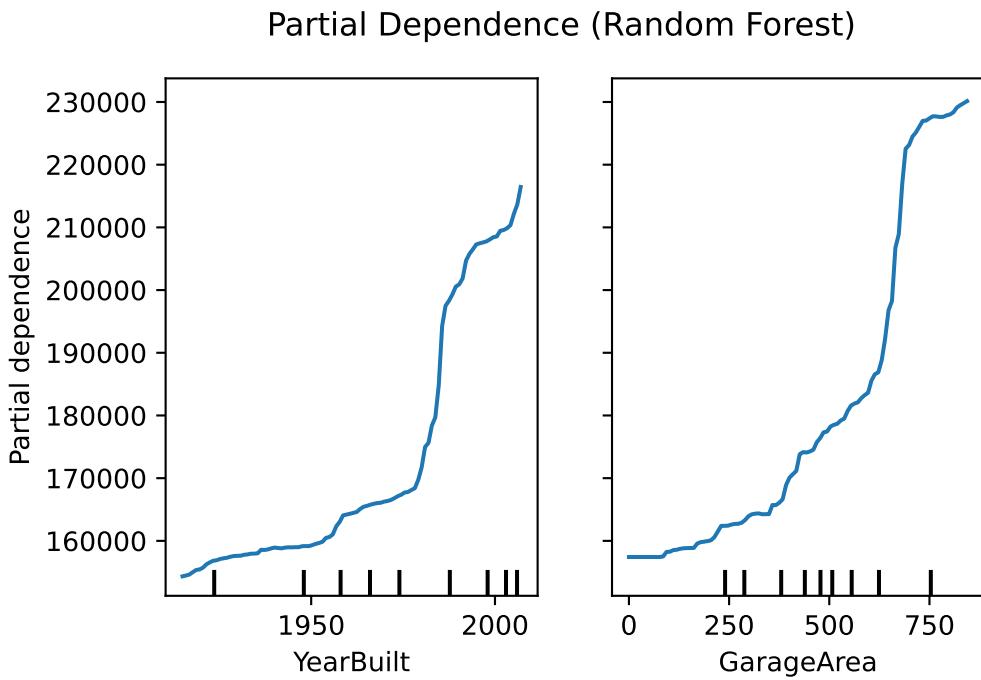
pd.DataFrame([["Random Forest (Tuned)", rmse_rf, mae_rf, mape_rf, r2_rf]],
             columns=["Model", "RMSE", "MAE", "MAPE", "R2"])

```

Best Random Forest Parameters: {'n_estimators': 500, 'min_samples_split': 2, 'max_features': 10, 'max_depth': 10, 'n_jobs': -1}



<Figure size 1650x1050 with 0 Axes>



Model	RMSE	MAE	MAPE	R2
0 Random Forest (Tuned)	1.093313e+09	19889.772151	12.832756	0.780717

7.5 XGBoost: Random Variable, Importance, PDPs, and Tuning

```

xgb = XGBRegressor(
    objective='reg:squarederror',
    random_state=seed,
    n_jobs=-1
)

# Define tuning grid similar to R (eta = learning_rate, subsample, max_depth)
param_dist_xgb = {
    'n_estimators': [24, 50, 100, 200],
    'learning_rate': [0.1, 0.15, 0.2, 0.25, 0.3],
    'max_depth': [1, 3, 5, 7, 10],
    'subsample': [0.25, 0.5, 0.75, 1],
    'colsample_bytree': [0.5, 0.75, 1]
}

xgb_random = RandomizedSearchCV(
    estimator=xgb,
    param_distributions=param_dist_xgb,
    n_iter=20,
    cv=cv,
    random_state=seed,
    n_jobs=-1,
    scoring='neg_root_mean_squared_error'
)

xgb_random.fit(X_train, y_train)
xgb_best = xgb_random.best_estimator_
print("Best XGBoost Parameters:", xgb_random.best_params_)

# Importance plot with random variable
plot_importance(xgb_best, max_num_features=15, importance_type='gain')
plt.title("XGBoost: Variable Importance (with Random Baseline)")
plt.show()

```

```


# Partial dependence plots
plt.figure()
PartialDependenceDisplay.from_estimator(xgb_best, X_train, ['YearBuilt', 'GarageArea'])
plt.suptitle("Partial Dependence (XGBoost)")
plt.show()

# Metrics
pred_xgb = xgb_best.predict(X_test)
rmse_xgb = mean_squared_error(y_test, pred_xgb)
mae_xgb = mean_absolute_error(y_test, pred_xgb)
mape_xgb = np.mean(np.abs((y_test - pred_xgb) / y_test)) * 100
r2_xgb = r2_score(y_test, pred_xgb)

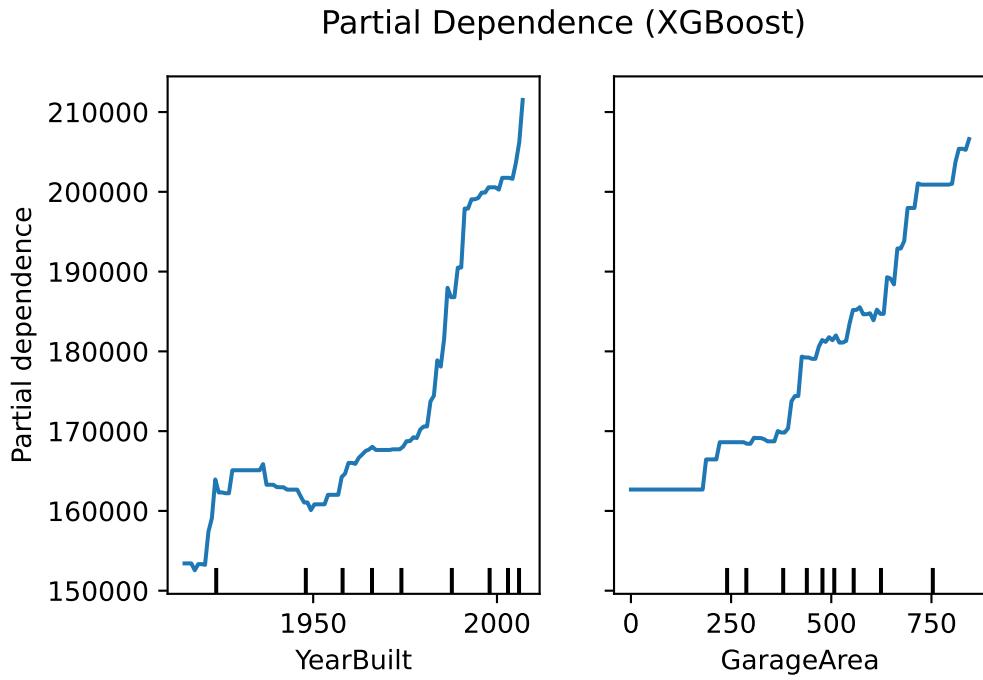
pd.DataFrame([["XGBoost (Tuned)", rmse_xgb, mae_xgb, mape_xgb, r2_xgb]],
             columns=["Model", "RMSE", "MAE", "MAPE", "R2"])

```

Best XGBoost Parameters: {'subsample': 0.75, 'n_estimators': 50, 'max_depth': 5, 'learning_rate': 0.1, 'gamma': 0.4, 'min_child_weight': 6, 'colsample_bytree': 0.75, 'reg_alpha': 0.01, 'reg_lambda': 0.01}



<Figure size 1650x1050 with 0 Axes>



	Model	RMSE	MAE	MAPE	R2
0	XGBoost (Tuned)	1.133839e+09	20819.921875	13.330445	0.772589

7.6 Final Model Comparison

```
results = pd.DataFrame([
    ["Decision Tree", rmse_tree, mae_tree, mape_tree, r2_tree],
    ["Random Forest (Tuned)", rmse_rf, mae_rf, mape_rf, r2_rf],
    ["XGBoost (Tuned)", rmse_xgb, mae_xgb, mape_xgb, r2_xgb]
], columns=["Model", "RMSE", "MAE", "MAPE", "R2"]).sort_values("RMSE")
results
```

	Model	RMSE	MAE	MAPE	R2
1	Random Forest (Tuned)	1.093313e+09	19889.772151	12.832756	0.780717
2	XGBoost (Tuned)	1.133839e+09	20819.921875	13.330445	0.772589
0	Decision Tree	2.378377e+09	28465.348099	17.701235	0.522975

Summary: The Decision Tree provides interpretability, Random Forest and XG-Boost attempt to further reduce bias and manage variance error through randomization and ensembles of weak learners. The issue is that they become less and less interpretable, so we have to introduce other tools to understand what variables are important and how they are relating to the outcome.

8 Understanding Random Forest and XGBoost Hyperparameters

8.1 Random Forest Hyperparameters

8.1.1 Switching Between Regression and Classification

Language	How to Specify Regression	How to Specify Classification
R	<code>randomForest(y ~ ., data = df)</code> where <code>y</code> is numeric	<code>randomForest(y ~ ., data = df)</code> where <code>y</code> is a factor
Python	<code>RandomForestRegressor()</code>	<code>RandomForestClassifier()</code>

In **R**, the outcome type (`numeric` vs `factor`) determines whether the model performs **regression** or **classification** automatically.

In **Python**, you explicitly choose the class (`RandomForestRegressor` or `RandomForestClassifier`).

Concept	R (<code>randomForest</code>)	Python (<code>RandomForestRegressor</code> / <code>RandomForestClassifier</code>)	Meaning / Class	Typical Values	Impact / In- terpretation
Model type	Determined by outcome type (<code>numeric</code> vs <code>factor</code>)	Different estimator class	Defines whether the forest predicts continuous or categorical outcomes	—	Choose based on target type

Concept	Python <code>(RandomForestRegressor / RandomForestClassifier)</code>			Typical Values	Impact / Interpretation
	R <code>(randomForest)</code>	Meaning /			
Number of trees	<code>ntree</code>	<code>n_estimators</code>	Number of trees to grow in the forest	100–1000 (500 common)	More trees reduce variance but increase runtime
Features per split	<code>mtry</code>	<code>max_features</code>	# of variables randomly selected at each split	Regression: \sqrt{p} ; Classification: $p/3$	Smaller \rightarrow higher bias, more diversity; larger \rightarrow lower bias, more correlation
Max tree depth	<i>(no direct arg, via nodesize)</i>	<code>max_depth</code>	Maximum depth of each tree	None, or 5–30	Deeper trees capture complexity but risk overfitting
Min samples per split	<i>(via nodesize indirectly)</i>	<code>min_samples_split</code>	Minimum # of samples to split a node	2–10	Larger \rightarrow smoother, higher bias; smaller \rightarrow lower bias, higher variance
Min samples per leaf	<i>(via nodesize)</i>	<code>min_samples_leaf</code>	Minimum samples per terminal node	1–5	Prevents tiny leaves, reduces overfitting
Bootstrap sampling	<code>replace</code>	<code>bootstrap</code>	Whether to sample with replacement	TRUE/FALSE	Adds randomness; TRUE recommended
Sample fraction per tree	<i>(rarely tuned)</i>	<code>max_samples</code>	Fraction of training data used per tree	0.5–1.0	Lower \rightarrow more randomness, higher bias

Concept	R (randomForest)	Python (RandomForestRegressor / RandomForestClassifier)	Meaning /	Typical Values	Impact / In-terpretation
Learning objective	Implicit: minimize MSE (regression) or Gini/entropy (classifica- tion)	Implicit: minimize MSE or Gini/entropy	Objective changes au- tomatically with model type	–	Regression averages numeric values; classification votes on class labels
Random seed	<code>set.seed()</code>	<code>random_state</code>	Ensures re- producibility	any integer	Use same seed for consistency

8.1.2 Thinking About Random Forest Tuning

- Start with **many trees** (500).
- Tune **mtry / max_features** to balance bias vs variance.
- Use **max_depth**, **min_samples_split**, or **min_samples_leaf** to prevent overfitting.
- RF tuning is relatively forgiving — focus on **efficiency and stability** rather than micro-optimization.
- For **classification**, check **OOB error** or **ROC-AUC** rather than RMSE.

8.2 XGBoost Hyperparameters

Concept	R (<code>xgboost</code>)	Python (<code>XGBRegressor</code> / <code>XGBClassifier</code>)		Meaning / Role	Typical Values	Impact / Interpretation
Boosting rounds	<code>nrounds</code>	<code>n_estimators</code>		Total # of trees (boosting steps)	100–1000	Too high → overfit unless learning rate small
Learning rate	<code>eta</code>	<code>learning_rate</code>		Shrinks each tree's contribution	0.01–0.3	Lower = slower, safer; higher = faster, riskier
Tree depth	<code>max_depth</code>	<code>max_depth</code>		Max depth per tree	3–10	Shallow = less overfit; deep = more complex
Min child weight	<code>min_child_weight</code>	<code>min_child_weight</code>		Min sum of instance weights per leaf	1–10	Larger = more conservative splits
Subsample (rows)	<code>subsample</code>	<code>subsample</code>		Fraction of training rows used per round	0.5–1.0	Lower adds randomness, improves generalization
Col sample (features)	<code>colsample_bytree</code>	<code>colsample_bytree</code>		Fraction of features per tree	0.5–1.0	Reduces feature correlation
Gamma	<code>gamma</code>	<code>gamma</code>		Min loss reduction to split a node	0–5	Higher = more conservative, prevents overfitting
L1 regularization	<code>alpha</code>	<code>reg_alpha</code>		Penalty on abs(weights)	0–10	Encourages sparsity (feature selection)
L2 regularization	<code>lambda</code>	<code>reg_lambda</code>		Penalty on squared(weights)	0–10	Controls weight magnitudes (stability)

Concept	R (<code>xgboost</code>)	Python (<code>XGBRegressor</code> / <code>XGBClassifier</code>)		Meaning / Role	Typical Values	Impact / Interpretation
Learning objective	<code>objective</code>	<code>objective</code>		Specifies prediction type	"reg:squarederror", "binary:logistic"	match target type
Early stopping	via <code>xgb.cv()</code>	<code>early_stopping_rounds</code>		no improvement	10–50 rounds	Prevents unnecessary rounds
Seed	<code>set.seed()</code>	<code>random_state</code>		Ensures reproducibility	any integer	Consistent results across runs

8.2.1 Thinking About XGBoost Tuning

1. Start simple — tune `n_estimators`, `max_depth`, and `learning_rate` first.
2. Add regularization (`gamma`, `alpha`, `lambda`) to control overfitting.
3. Use `subsample` and `colsample_bytree` to add randomness.

8.3 Practical Tips and Heuristics

Situation	Random Forest Strategy	XGBoost Strategy
Overfitting	Increase <code>min_samples_split</code> , reduce <code>max_depth</code>	Lower <code>learning_rate</code> , increase <code>gamma</code> or <code>min_child_weight</code>
Underfitting	Increase <code>n_estimators</code> , allow deeper trees	Lower regularization, increase <code>max_depth</code> or <code>n_estimators</code>
Slow training	Reduce trees or sample fraction	Increase <code>learning_rate</code> , reduce trees

Situation	Random Forest Strategy	XGBoost Strategy
Unstable results	Fix <code>random_state</code> , use more trees	Fix <code>random_state</code> , enable early stopping
High variance	Lower <code>max_features</code> , increase randomness	Increase <code>subsample</code> diversity, lower <code>max_depth</code>

8.4 Rule-of-Thumb Summary

Parameter	Low Value →	High Value →
<code>n_estimators</code>	Underfitting	Stable but slower
<code>max_depth</code>	Simpler model (underfit)	Complex model (overfit)
<code>learning_rate</code>	Slow, steady, robust	Fast, risk of overfit
<code>subsample</code> / <code>colsample_bytree</code>	More bias, less variance	Lower bias, higher variance
<code>mtry</code> / <code>max_features</code>	More randomization	More correlation between trees
<code>min_samples_split</code> / <code>min_child_weight</code>	Flexible, low bias	Conservative, higher bias

8.4.1 Key Takeaways

- **Random Forest**
 - Regression/classification determined by outcome type (R) or model class (Python).
 - Reduces variance through bagging.
 - Tuning focuses on **stability** and **efficiency**.
- **XGBoost**
 - Requires explicit objective selection.
 - Reduces error via boosting and regularization.
- Use interpretability tools (importance, PDPs, SHAP) to understand **why** tuning changes model behavior.

9 Neural Networks in R

9.1 Introduction

This chapter introduces **neural networks in R**, progressing from feed-forward to recurrent (RNN) and convolutional (CNN) architectures.

It uses the Ames housing dataset, a sine wave sequence, and the MNIST digit dataset to demonstrate increasingly complex network structures.

Each example illustrates a different ANN architecture.

9.2 1. Feed-Forward Neural Networks (Ames Data)

Feed-forward networks are the simplest type of neural networks. They map input features (like home size, rooms, or year built) to an output (home price) through **layers of neurons**. Each neuron applies a weighted sum and activation function, allowing the network to learn complex relationships between variables.

```
#Note: This is specific to my environment
library(reticulate)
use_python("/Users/donaldhale/r-arm/bin/python", required = TRUE)

library(keras3)

tf <- reticulate::import("tensorflow", delay_load = FALSE)
cat("TensorFlow version:", tf$`__version__`, "\n")
```

TensorFlow version: 2.16.1

```

library(tidyverse)
library(caret)
library(nnet)
library(NeuralNetTools)
library(AmesHousing)

set.seed(4321)

# Load and sample data
ames <- make_ordinal_ames() %>% mutate(id = row_number())
train_idx <- sample(1:nrow(ames), 0.7 * nrow(ames))
training <- ames[train_idx, ]
testing <- ames[-train_idx, ]

# Select a manageable subset of variables
vars <- c("Sale_Price", "Bedroom_AbvGr", "Year_Built", "Mo_Sold", "Lot_Area",
         "Street", "Central_Air", "First_Flr_SF", "Second_Flr_SF",
         "Full_Bath", "Half_Bath", "Fireplaces", "Garage_Area",
         "Gr_Liv_Area", "TotRms_AbvGrd")
training <- training[, vars]
testing <- testing[, vars]

# Convert categorical predictors into dummy numeric columns
x_train <- model.matrix(Sale_Price ~ ., data = training)[, -1]
y_train <- training$Sale_Price
x_test <- model.matrix(Sale_Price ~ ., data = testing)[, -1]
y_test <- testing$Sale_Price

# Drop low-variance predictors and scale for stable training
nzv <- nearZeroVar(x_train)
if (length(nzv) > 0) {
  x_train <- x_train[, -nzv, drop = FALSE]
  x_test <- x_test[, -nzv, drop = FALSE]
}

x_train <- scale(x_train)
x_test <- scale(x_test, center = attr(x_train, "scaled:center"),
                 scale = attr(x_train, "scaled:scale"))

y_train_s <- scale(y_train)

# Combine into a single training dataframe


```

```

train_df <- data.frame(y = as.numeric(y_train_s), x_train)

set.seed(123)
nn.ames <- nnet(y ~ ., data = train_df, size = 5, linout = TRUE, trace = FALSE)
plotnet(nn.ames)

```



9.2.1 1.2 Cross-Validation Tuning

We tune two important parameters:

- **size** — number of neurons in the hidden layer
- **decay** — regularization strength (prevents overfitting)

`caret::train()` automates this search using k-fold cross-validation.

```

tune_grid <- expand.grid(size = c(3,5,7,9), decay = c(0,0.1,0.5,1))
ctrl <- trainControl(method = "cv", number = 5)

set.seed(123)

```

```

nn.tuned <- train(
  x = x_train, y = as.numeric(y_train_s),
  method = "nnet",
  tuneGrid = tune_grid,
  trControl = ctrl,
  linout = TRUE,
  trace = FALSE
)

nn.tuned$bestTune

```

```

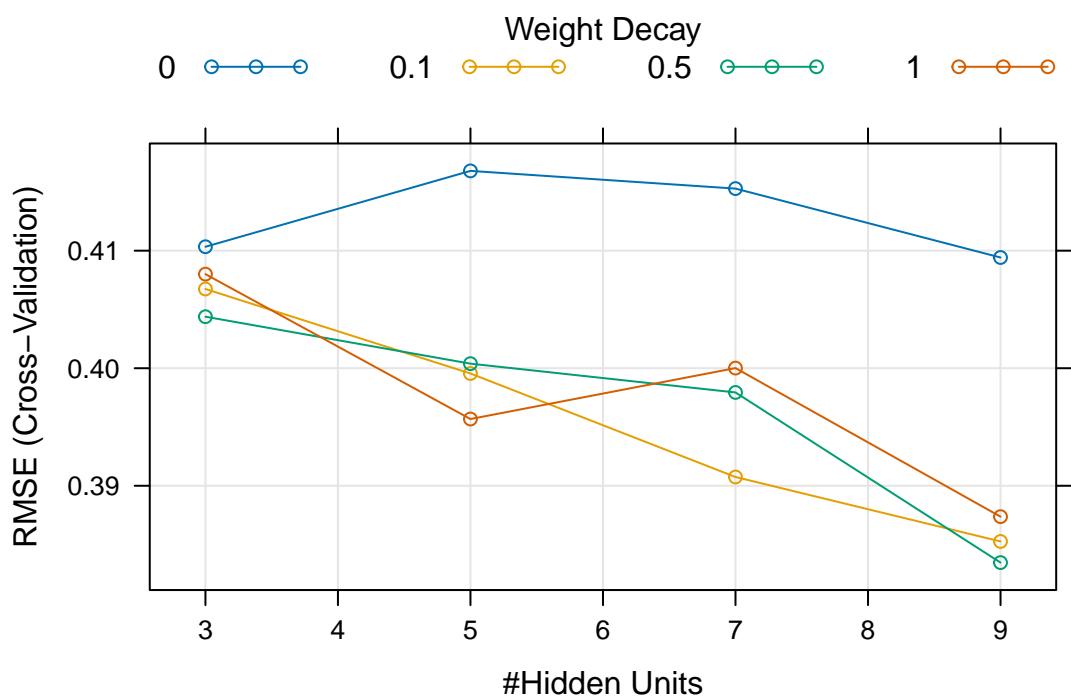
  size decay
15      9    0.5

```

```

plot(nn.tuned)

```



9.2.2 1.3 Evaluate Performance

After identifying the best parameters, we retrain the model and evaluate its predictive accuracy.

```
best_size <- nn.tuned$bestTune$size
best_decay <- nn.tuned$bestTune$decay

set.seed(123)
nn.final <- nnet(
  y ~ .,
  data = train_df,
  size = best_size,
  decay = best_decay,
  linout = TRUE,
  trace = FALSE
)

pred_std <- predict(nn.final, newdata = as.data.frame(x_test))
pred_nn <- pred_std * attr(y_train_s, "scaled:scale") + attr(y_train_s, "scaled:center")

perf <- data.frame(
  MAE = mean(abs(y_test - pred_nn)),
  MAPE = mean(abs((y_test - pred_nn) / y_test)) * 100
)
perf
```

	MAE	MAPE
1	20392.57	12.63864

9.3 2. Feed-Forward Neural Networks (Keras 3)

The **Keras 3** API is a modern deep learning interface. It supports more complex models and GPUs while maintaining a simple syntax.

This example shows the same concept as above using a multi-layer perceptron (MLP).

```

#library(keras3)
set.seed(123)


model_ff <- keras_model_sequential() |>
  layer_dense(units = 16, activation = "relu", input_shape = ncol(x_train)) |>
  layer_dropout(rate = 0.1) |>
  layer_dense(units = 1)

compile(model_ff, optimizer = optimizer_adam(learning_rate = 0.01),
        loss = "mse", metrics = "mae")

history <- fit(model_ff, x_train, as.numeric(y_train_s),
                 epochs = 40, batch_size = 32, validation_split = 0.2, verbose = 0)

plot(history)

```




9.4 3. Recurrent Neural Network (RNN)

RNNs are ideal for **sequential data** such as time series, text, or speech.

We use a sine wave as a toy example because its repeating pattern requires the model to “remember” recent values — something feed-forward networks cannot do.

```
set.seed(42)
t <- seq(0, 50, by = 0.1)
y <- sin(t)
plot(t, y, type='l', main="Sine Wave Example", ylab="y", xlab="Time")
```



```
window_size <- 20
X <- sapply(1:(length(y) - window_size), function(i) y[i:(i + window_size - 1)])
X <- t(X)
y_out <- y[(window_size + 1):length(y)]
X <- array(X, dim = c(nrow(X), window_size, 1))

split <- round(0.8 * nrow(X))
x_train <- X[1:split,,]; y_train <- y_out[1:split]
```

```

x_test <- X[(split + 1):nrow(X),]; y_test <- y_out[(split + 1):length(y_out)]

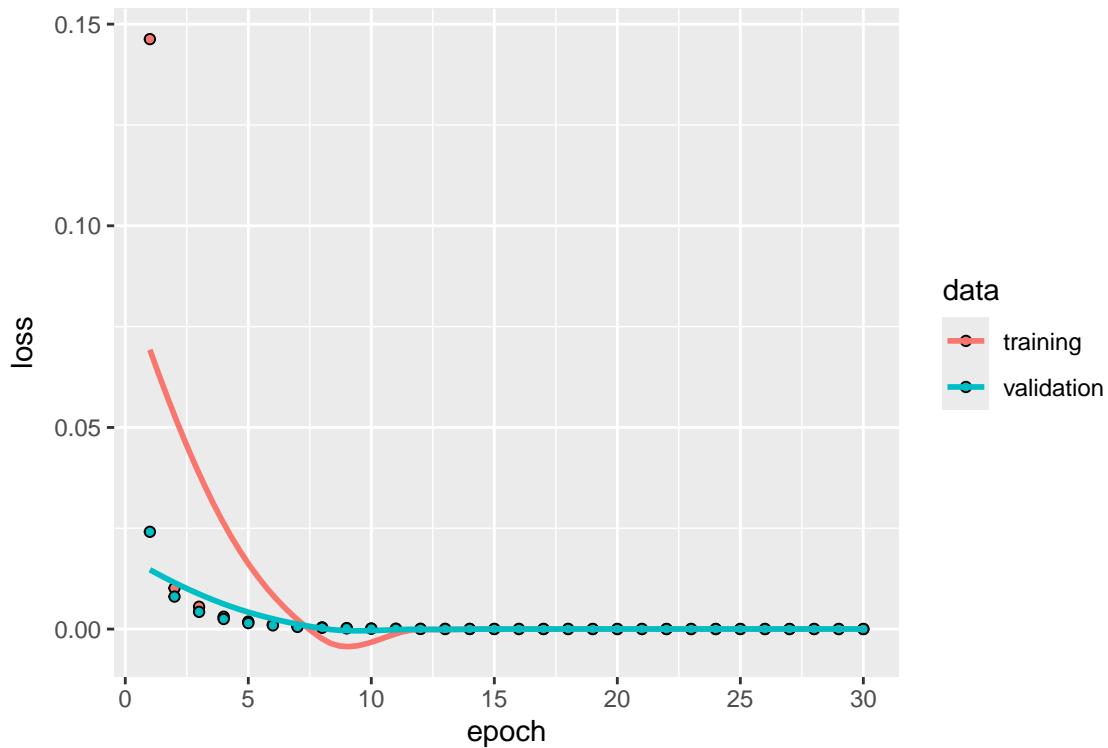
model_rnn <- keras_model_sequential() |>
  layer_simple_rnn(units = 16, activation = "tanh", input_shape = c(window_size, 1)) |>
  layer_dense(units = 1)

compile(model_rnn, loss = "mse", optimizer = "adam")

history_rnn <- fit(model_rnn, x_train, y_train, epochs = 30, batch_size = 16,
                     validation_split = 0.2, verbose = 0)

plot(history_rnn)

```



```

preds <- predict(model_rnn, x_test)


```

3/3 - 0s - 11ms/step

```

plot(y_test, type='l', main="RNN Sine Prediction", col='blue', ylab="y")
lines(preds, col='red')
legend("bottomleft", legend=c("Actual","Predicted"), col=c("blue","red"), lty=1, bty="n")

```



9.5 4. Convolutional Neural Network (CNN)

CNNs are built for **spatial pattern recognition** — ideal for images.

The **MNIST** dataset contains 70,000 grayscale images of handwritten digits (0–9), each 28×28 pixels.

These lines prepare the data for use in a CNN.

```

#library(keras3)

# Load and preview the MNIST data
mnist <- dataset_mnist()

```

```

# We only use a subset (5k train, 1k test) for quick demonstration
x_train <- mnist$train$x[1:5000,,]
y_train <- mnist$train$y[1:5000]
x_test  <- mnist$test$x[1:1000,,]
y_test  <- mnist$test$y[1:1000]

# Normalize pixel values: 0-255 -> 0-1
x_train <- x_train / 255
x_test  <- x_test / 255

# Reshape to [samples, height, width, channels]
# MNIST images are grayscale, so channels = 1
x_train <- array_reshape(x_train, c(nrow(x_train), 28, 28, 1))
x_test  <- array_reshape(x_test,  c(nrow(x_test), 28, 28, 1))

# Convert labels to categorical one-hot vectors (10 classes for digits 0-9)
y_train_cat <- to_categorical(y_train, 10)
y_test_cat  <- to_categorical(y_test, 10)

```

9.5.1 Explanation of the MNIST Preprocessing

1. **Load** the built-in dataset from Keras.
2. **Subsample** for faster training (5,000 train / 1,000 test).
3. **Normalize** pixel intensities to the 0–1 range.
4. **Reshape** to the 4D structure required by CNNs: [samples, height, width, channels].
5. **One-hot encode** labels, converting a digit like 3 into a vector [0,0,0,1,0,0,0,0,0,0].

These steps turn raw pixel data into numeric tensors suitable for learning.

9.5.2 4.1 Train and Evaluate the CNN

```

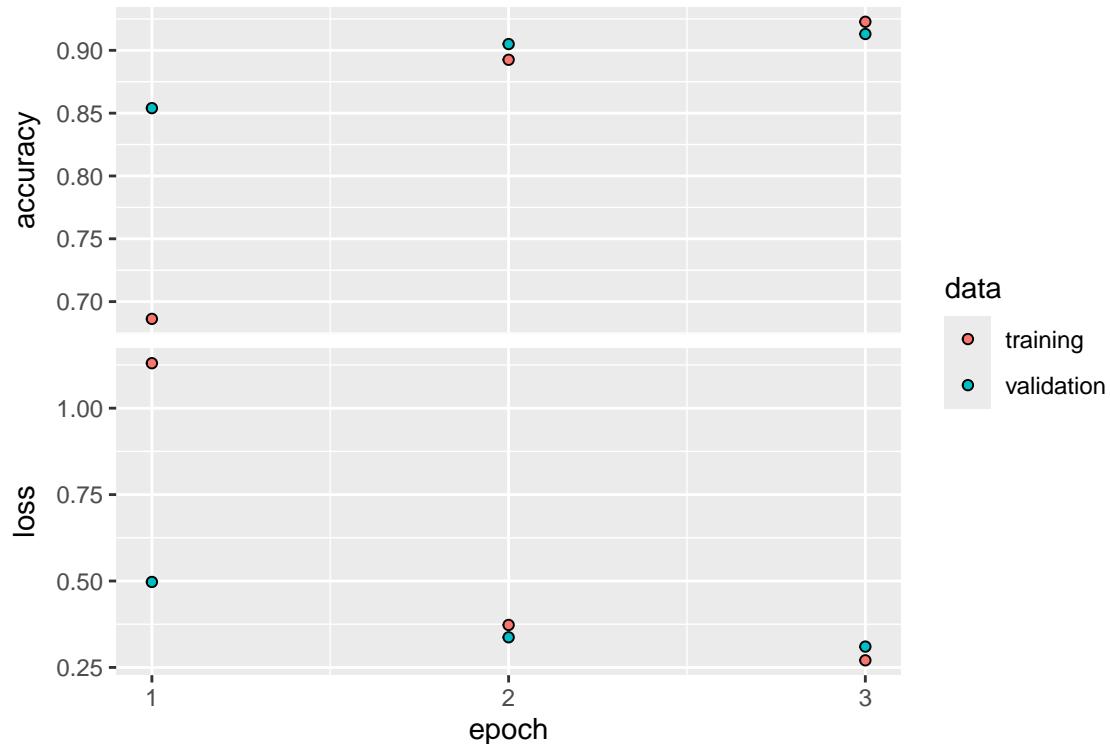
model_cnn <- keras_model_sequential() |>
  layer_conv_2d(filters = 16, kernel_size = c(3,3), activation = "relu", input_shape = c(28,28,3)) |>
  layer_max_pooling_2d(pool_size = c(2,2)) |>
  layer_flatten() |>
  layer_dense(units = 32, activation = "relu") |>
  layer_dense(units = 10, activation = "softmax")

compile(model_cnn, optimizer = "adam",
        loss = "categorical_crossentropy",
        metrics = "accuracy")

history_cnn <- fit(model_cnn, x_train, y_train_cat,
                     epochs = 3, batch_size = 64, validation_split = 0.2, verbose = 0)

plot(history_cnn)

```



```
score <- evaluate(model_cnn, x_test, y_test_cat)
```

32/32 - 0s - 745us/step - accuracy: 0.9050 - loss: 0.3117


```
score
```

```
$accuracy  
[1] 0.905  
  
$loss  
[1] 0.3116771
```

9.5.3 4.2 Visualizing Correct and Incorrect Predictions

We visualize examples of digits that the CNN classified correctly and incorrectly.

```
pred_probs <- predict(model_cnn, x_test)  
  
32/32 - 0s - 1ms/step  
  
pred_classes <- apply(pred_probs, 1, which.max) - 1  
  
correct_idx <- which(pred_classes == y_test)[1:6]  
incorrect_idx <- which(pred_classes != y_test)[1:6]  
  
par(mfrow=c(2,6), mar=c(1,1,2,1))  
for (i in correct_idx) {  
  img <- x_test[i,,,1]  
  image(1:28, 1:28, t(apply(img, 2, rev)), col=gray.colors(255), axes=FALSE,  
        main=paste(" ", y_test[i]))  
}  
for (i in incorrect_idx) {  
  img <- x_test[i,,,1]  
  image(1:28, 1:28, t(apply(img, 2, rev)), col=gray.colors(255), axes=FALSE,  
        main=paste(" ", y_test[i], "→", pred_classes[i]))  
}
```



9.6 Summary

- **Feed-forward networks:** handle static, tabular data (e.g., housing prices).
- **RNNs:** handle sequential data with temporal dependencies (e.g., sine wave, stock prices, text, timer series).
- **CNNs:** handle image and spatial data by learning patterns like edges and textures.

10 Feedforward Neural Network - Python

10.1 Feedforward Neural Network (MLP) Using Ames Housing (Python)

This chapter fits a feed-forward neural network to the Ames Housing subset used in the R version.

RMSE is computed manually for compatibility with older scikit-learn versions.

10.1.1 1. Load and Prepare Data

```
import numpy as np
import pandas as pd

from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_squared_error

np.random.seed(12345)

ames = fetch_openml(name="house_prices", as_frame=True)
df = ames.frame

selected_cols = [
    "SalePrice", "BedroomAbvGr", "YearBuilt", "MoSold", "LotArea",
    "Street", "CentralAir", "1stFlrSF", "2ndFlrSF", "FullBath",
    "HalfBath", "Fireplaces", "GarageArea", "GrLivArea", "TotRmsAbvGrd"
]

df_sub = df[selected_cols].dropna()
```

```
df_sub = pd.get_dummies(df_sub, columns=["Street", "CentralAir"], drop_first=True)

y = df_sub["SalePrice"]
X = df_sub.drop("SalePrice", axis=1)
```

10.1.2 2. Train/Test Split & Scaling

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=12345
)

scaler = StandardScaler()
scaler.fit(X_train)

X_train_s = scaler.transform(X_train)
X_test_s = scaler.transform(X_test)
```

10.1.3 3. Simple MLP Model

```
nn_simple = MLPRegressor(
    solver="lbfgs",
    alpha=1e-5,
    hidden_layer_sizes=(5,),
    random_state=12345,
    max_iter=5000
)

nn_simple.fit(X_train_s, y_train)
pred_simple = nn_simple.predict(X_test_s)

rmse_simple = mean_squared_error(y_test, pred_simple) ** 0.5
rmse_simple
```

38907.904754996634

10.1.4 4. Grid Search Tuning

```
param_grid = {
    "hidden_layer_sizes": [(h,) for h in [3,4,5,6,7]],
    "alpha": [0.00005,0.0005],
    "solver": ["lbfgs"]
}

nn_base = MLPRegressor(max_iter=5000, random_state=12345)

grid = GridSearchCV(
    estimator=nn_base,
    param_grid=param_grid,
    cv=10,
    scoring="neg_mean_squared_error",
    n_jobs=-1
)

grid.fit(X_train_s, y_train)

grid.best_params_

{'alpha': 5e-05, 'hidden_layer_sizes': (5,), 'solver': 'lbfgs'}

best_model = grid.best_estimator_
pred_gs = best_model.predict(X_test_s)

rmse_gs = mean_squared_error(y_test, pred_gs) ** 0.5
rmse_gs
```

39176.951700141195

11 Naive Bayes Classification R version

11.1 1. Overview and Learning Objectives

In this chapter, we will:

1. Use the **Ames Housing** data to create a **binary classification** problem:
 - Predict whether a house is **high price** vs **low price** based on its features..
2. Understand how Naive Bayes handles:
 - **Categorical vs numeric** predictors
 - **Laplace correction** (additive smoothing)
 - **Kernel density estimation** for numeric variables
3. Use **caret** (`train(method = "nb")`) to:
 - **Tune** key hyperparameters (`usekernel`, `fL`)
 - Evaluate performance with **cross-validation**
4. Interpret:
 - Class priors
 - Conditional feature distributions
 - Predictions & confusion matrix
 - ROC curve and AUC

11.2 2. Data Setup and Binary Target Creation

```
library(tidyverse)
library(caret)
library(AmesHousing)
library(pROC)

set.seed(4321)

ames <- make_ordinal_ames() %>% mutate(id = row_number())

training <- ames %>% sample_frac(0.7)
testing <- anti_join(ames, training, by = "id")

predictors <- c(
  "Sale_Price", "Bedroom_AbvGr", "Year_Built", "Mo_Sold", "Lot_Area",
  "Street", "Central_Air", "First_Flr_SF", "Second_Flr_SF", "Full_Bath",
  "Half_Bath", "Fireplaces", "Garage_Area", "Gr_Liv_Area", "TotRms_AbvGrd"
)

training <- training %>% select(all_of(predictors))
testing <- testing %>% select(all_of(predictors))
```

We turn `Sale_Price` into a **classification target**:

- "High" if above the training-set median
- "Low" otherwise

This keeps the task aligned with Naive Bayes as a **classification** model.

```
price_median <- median(training$Sale_Price)

training <- training %>% mutate(Price_High = factor(ifelse(Sale_Price > price_median, "High", "Low")))
testing <- testing %>% mutate(Price_High = factor(ifelse(Sale_Price > price_median, "High", "Low")))

training_nb <- training %>% select(-Sale_Price)
testing_nb <- testing %>% select(-Sale_Price)
```

11.3 3. Naive Bayes Intuition

Naive Bayes is a **generative classifier** based on Bayes' theorem:

$$[P(Y = k | X) = P(Y=k) \prod_{j=1}^p P(x_j | Y=k)]$$

Key assumptions:

11.3.0.1 1. Conditional Independence

Given the class, features are treated as independent:

$$[P(x_1, x_2, \dots, x_p | Y=k) = \prod_j P(x_j | Y=k)]$$

This is rarely exactly true, yet often works surprisingly well.

11.3.0.2 2. Modeling Numeric Features

Naive Bayes supports two approaches:

- **Gaussian (Normal):** assumes each numeric feature is normally distributed within each class.
- **Kernel Density Estimation:** more flexible, handles skewness and multimodal distributions.

11.3.0.3 3. Laplace Correction

Prevents zero probabilities when a category never appears in a given class:

$$\text{P}(x=v | Y=k) = (\text{count} + 1) / (N + K)$$

This keeps the model from collapsing.

11.4 4. Fitting Naive Bayes (Baseline Model)

```

set.seed(12345)

nb.ames <- naiveBayes(
  Price_High ~ .,
  data = training_nb,
  laplace = 0,
  usekernel = TRUE
)

nb.ames

```

Naive Bayes Classifier for Discrete Predictors

Call:

```
naiveBayes.default(x = X, y = Y, laplace = laplace, usekernel = TRUE)
```

A-priori probabilities:

	High	Low
Y	0.4987811	0.5012189

Conditional probabilities:

	Bedroom_AbvGr	
Y	[,1]	[,2]
High	2.948192	0.8218245
Low	2.770428	0.8207514

	Year_Built	
Y	[,1]	[,2]
High	1989.604	23.12763
Low	1953.780	25.88135

	Mo_Sold	
Y	[,1]	[,2]
High	6.353861	2.732694
Low	6.157588	2.691989

	Lot_Area	
Y	[,1]	[,2]
High	11726.913	10345.651
Low	8673.011	4673.732

Street

	Grvl	Pave
Y	High 0.0009775171	0.9990224829
	Low 0.0058365759	0.9941634241

Central_Air

	N	Y
Y	High 0.008797654	0.991202346
	Low 0.122568093	0.877431907

First_Flr_SF

	[,1]	[,2]
Y	High 1333.7087	408.4059
	Low 986.6527	264.8090

Second_Flr_SF

	[,1]	[,2]
Y	High 447.7625	493.0434
	Low 222.7490	317.7090

Full_Bath

	[,1]	[,2]
Y	High 1.913001	0.4415516
	Low 1.233463	0.4412614

Half_Bath

	[,1]	[,2]
Y	High 0.5259042	0.5168999
	Low 0.2324903	0.4406718

Fireplaces

	[,1]	[,2]
Y	High 0.8670577	0.6271108
	Low 0.3180934	0.5359368

Garage_Area

	[,1]	[,2]
Y	High 581.0430	173.7891
	Low 362.3016	190.0580

Gr_Liv_Area

	[,1]	[,2]
Y		

```
High 1783.837 473.5654
Low 1215.255 342.4179
```

```
TotRms_AbvGrd
Y      [,1]      [,2]
High 7.043988 1.516711
Low  5.825875 1.373640
```

11.4.0.1 Explanation of key settings

- **laplace = 0**

No smoothing; probability estimates reflect raw counts.
Zero probabilities may occur for rare categorical levels.

- **usekernel = TRUE**

Numeric features are modeled using **kernel density estimates** instead of assuming normality.

11.4.0.2 What the output means

- **A-priori probabilities:**

Estimated from class frequencies (High vs Low).

- **Conditional tables:**

- For categorical features: probability of each level given class.
- For numerical features: kernel-based density info (not printed in summary).

11.5 5. Predicting on Test Data

```
test_pred_class <- predict(nb.ames, testing_nb)
test_pred_prob  <- predict(nb.ames, testing_nb, type="raw")
head(test_pred_prob)
```

	High	Low
[1,]	0.9999965	3.462116e-06
[2,]	0.9850961	1.490391e-02
[3,]	0.9830462	1.695380e-02
[4,]	0.5936144	4.063856e-01
[5,]	0.9437367	5.626333e-02
[6,]	0.9998253	1.746816e-04

- `test_pred_class`: predicted High/Low label
 - `test_pred_prob`: estimated probability for each class
-

11.6 6. Confusion Matrix Interpretation

```
conf_matrix <- confusionMatrix(
  test_pred_class,
  testing_nb$Price_High,
  positive="High"
)

conf_matrix
```

Confusion Matrix and Statistics

		Reference	
		Prediction	High Low
		High	373 49
		Low	55 402
Accuracy : 0.8817			
95% CI : (0.8585, 0.9023)			
No Information Rate : 0.5131			
P-Value [Acc > NIR] : <2e-16			
Kappa : 0.7631			

Mcnemar's Test P-Value : 0.6239

```

Sensitivity : 0.8715
Specificity : 0.8914
Pos Pred Value : 0.8839
Neg Pred Value : 0.8796
Prevalence : 0.4869
Detection Rate : 0.4243
Detection Prevalence : 0.4801
Balanced Accuracy : 0.8814

'Positive' Class : High

```

11.6.0.1 What the confusion matrix tells us

- **Accuracy:** overall correctness
- **Sensitivity (Recall for High):** proportion of expensive homes correctly identified
- **Specificity (Recall for Low):** proportion of cheaper homes correctly identified
- **Balanced Accuracy:** average of sensitivity & specificity
- **Kappa:** agreement beyond chance

This is the primary evaluation for classification problems.

11.7 7. ROC Curve & AUC

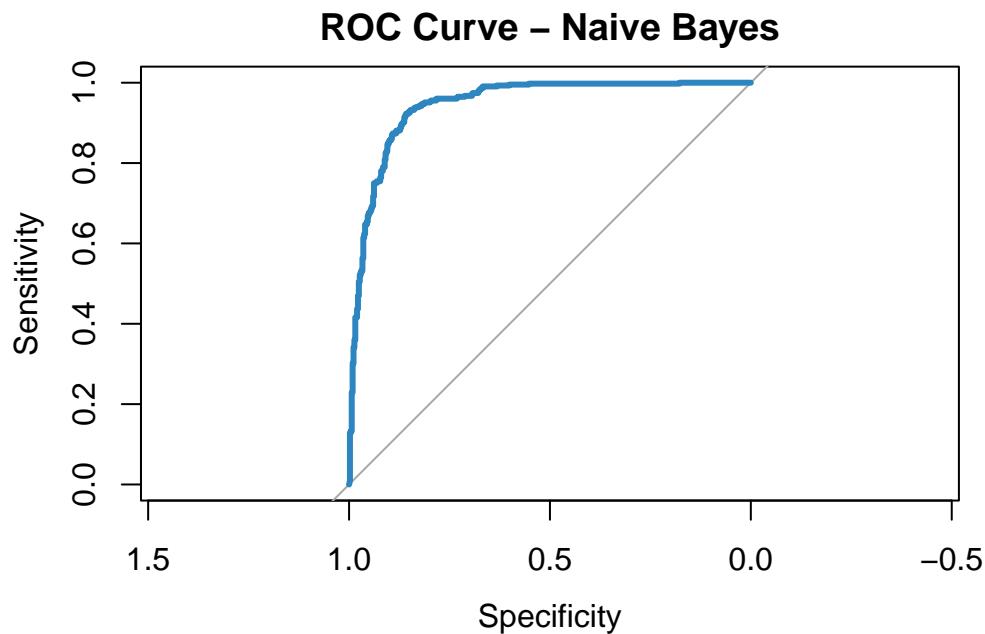
ROC curves evaluate model ranking ability — how well it separates High from Low across thresholds.

```

roc_obj <- roc(
  response = testing_nb$Price_High,
  predictor = test_pred_prob[, "High"],
  levels = c("Low", "High")
)

plot(roc_obj, col="#2E86C1", lwd=3, main="ROC Curve - Naive Bayes")

```



```
auc_value <- auc(roc_obj)
auc_value
```

Area under the curve: 0.9457

11.7.0.1 Interpretation

- **ROC Curve:** visual representation of sensitivity vs (1 - specificity)
- **AUC:** single summary metric
 - 0.5 = no better than random
 - 0.7 = decent
 - 0.8 = strong
 - 0.9 = outstanding

11.8 8. Tuning Naive Bayes Using caret::train()

We now tune:

- usekernel {TRUE, FALSE}
- fL {0, 0.5, 1} (*fractional Laplace*)

```
set.seed(12345)

tune_grid <- expand.grid(
  usekernel = c(TRUE, FALSE),
  fL        = c(0, 0.5, 1),
  adjust    = 1
)

ctrl <- trainControl(
  method = "cv",
  number = 10,
  classProbs = TRUE,
  summaryFunction = twoClassSummary
)

training_nb$Price_High <- relevel(training_nb$Price_High, ref="High")
testing_nb$Price_High  <- relevel(testing_nb$Price_High, ref="High")

nb.ames.caret <- train(
  Price_High ~.,
  data = training_nb,
  method = "nb",
  tuneGrid = tune_grid,
  trControl = ctrl,
  metric ="ROC"
)

nb.ames.caret
```

Naive Bayes

```
2051 samples
14 predictor
2 classes: 'High', 'Low'
```

```
No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 1845, 1846, 1846, 1846, 1846, 1845, ...
Resampling results across tuning parameters:
```

usekernel	fL	ROC	Sens	Spec
FALSE	0.0	0.9325587	0.9217484	0.8453794
FALSE	0.5	0.9325587	0.9217484	0.8453794
FALSE	1.0	0.9325587	0.9217484	0.8453794
TRUE	0.0	0.9549057	0.9197316	0.8589473
TRUE	0.5	0.9549057	0.9197316	0.8589473
TRUE	1.0	0.9549057	0.9197316	0.8589473

```
Tuning parameter 'adjust' was held constant at a value of 1
ROC was used to select the optimal model using the largest value.
The final values used for the model were fL = 0, usekernel = TRUE and adjust
= 1.
```

##\$ How tuning works

- **Cross-validation** evaluates each combination of hyperparameters.
- Best model chosen by the metric (ROC).
- caret automatically handles class probabilities and ROC scoring.

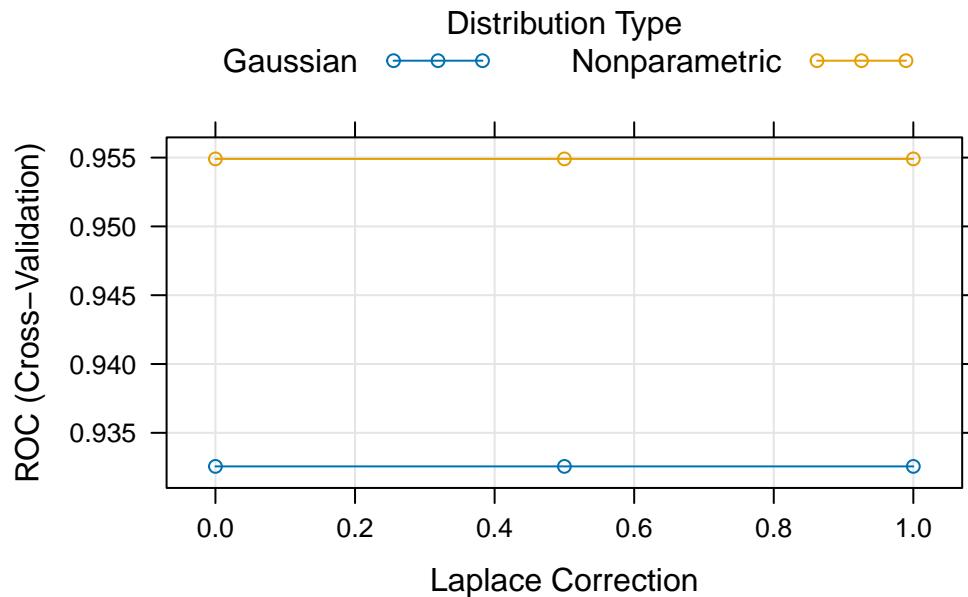
```
nb.ames.caret$results
```

	usekernel	fL	adjust	ROC	Sens	Spec	ROCSD	SensSD
1	FALSE	0.0	1	0.9325587	0.9217484	0.8453794	0.02290800	0.03508486
4	TRUE	0.0	1	0.9549057	0.9197316	0.8589473	0.01199907	0.04510391
2	FALSE	0.5	1	0.9325587	0.9217484	0.8453794	0.02290800	0.03508486
5	TRUE	0.5	1	0.9549057	0.9197316	0.8589473	0.01199907	0.04510391
3	FALSE	1.0	1	0.9325587	0.9217484	0.8453794	0.02290800	0.03508486
6	TRUE	1.0	1	0.9549057	0.9197316	0.8589473	0.01199907	0.04510391
				SpecSD				
1				0.02659457				
4				0.03971772				
2				0.02659457				
5				0.03971772				
3				0.02659457				
6				0.03971772				

```
nb.ames.caret$bestTune
```

```
fL usekernel adjust  
4 0      TRUE      1
```

```
plot(nb.ames.caret)
```



11.9 9. Evaluating the Tuned Model on the Test Set

```
test_pred_class_tuned <- predict(nb.ames.caret, testing_nb)  
test_pred_prob_tuned <- predict(nb.ames.caret, testing_nb, type="prob")  
  
confusionMatrix(  
  test_pred_class_tuned,  
  testing_nb$Price_High,  
  positive="High"  
)
```

Confusion Matrix and Statistics

		Reference
Prediction	High	Low
High	398	77
Low	30	374

Accuracy : 0.8783
95% CI : (0.8548, 0.8992)
No Information Rate : 0.5131
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.7571

McNemar's Test P-Value : 8.708e-06

Sensitivity : 0.9299
Specificity : 0.8293
Pos Pred Value : 0.8379
Neg Pred Value : 0.9257
Prevalence : 0.4869
Detection Rate : 0.4528
Detection Prevalence : 0.5404
Balanced Accuracy : 0.8796

'Positive' Class : High

11.10 10. Summary

- Naive Bayes is a **probabilistic, generative classifier**.
- It models **class priors + conditional likelihoods** for each predictor.
- Numeric predictors can use:
 - **Gaussian distributions**, or
 - **Kernel density estimates** (`usekernel = TRUE`).

- **Laplace smoothing** protects against zero probabilities.
- **ROC** and **AUC** evaluate ranking performance.
- **Confusion matrix** evaluates classification performance.
- **caret** enables systematic tuning over smoothing and density options.

12 Naive Bayes Classification - Python Version

12.1 1. Overview and Learning Objectives

In this Python chapter, we will:

1. Load the **Ames Housing** dataset from OpenML.
 2. Create a **High vs Low** price classification target.
 3. Fit **Gaussian Naive Bayes**.
 4. Produce predictions and class probabilities.
 5. Compute a **correct ROC curve and AUC**.
 6. Mirror the structure of the R version.
-

12.2 2. Setup and Data Preparation

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix, classification_report, roc_curve, auc
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_openml

# Load Ames Housing dataset
ames = fetch_openml(name="house_prices", as_frame=True)
df = ames.frame.copy()

df.rename(columns={"SalePrice": "Sale_Price"}, inplace=True)

df.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPu
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPu
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPu
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPu
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPu

12.3 2.1 Select Variables of Interest

```
predictors = [
    "Sale_Price",
    "BedroomAbvGr",
    "YearBuilt",
    "MoSold",
    "LotArea",
    "Street",
    "CentralAir",
    "1stFlrSF",
    "2ndFlrSF",
    "FullBath",
    "HalfBath",
    "Fireplaces",
    "GarageArea",
    "GrLivArea",
    "TotRmsAbvGrd"
]

df = df[predictors].dropna()
```

12.4 2.2 Create a Binary Target Variable

```
median_price = df["Sale_Price"].median()

df["Price_High"] = np.where(df["Sale_Price"] > median_price, "High", "Low")
df["Price_High"] = df["Price_High"].astype("category")

df["Price_High"].value_counts(normalize=True)
```

```
Price_High
Low      0.50137
High     0.49863
Name: proportion, dtype: float64
```

12.5 2.3 Train/Test Split

```
X = df.drop(columns=["Sale_Price", "Price_High"])
y = df["Price_High"]

# Encode categorical predictors
X = pd.get_dummies(X, drop_first=True)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.30, random_state=4321, stratify=y
)

X_train.shape, X_test.shape
```

```
((1022, 14), (438, 14))
```

12.6 3. Gaussian Naive Bayes Model

```
gnb = GaussianNB()
gnb.fit(X_train, y_train)
```

```
GaussianNB()
```

12.7 4. Predictions

```
y_pred = gnb.predict(X_test)

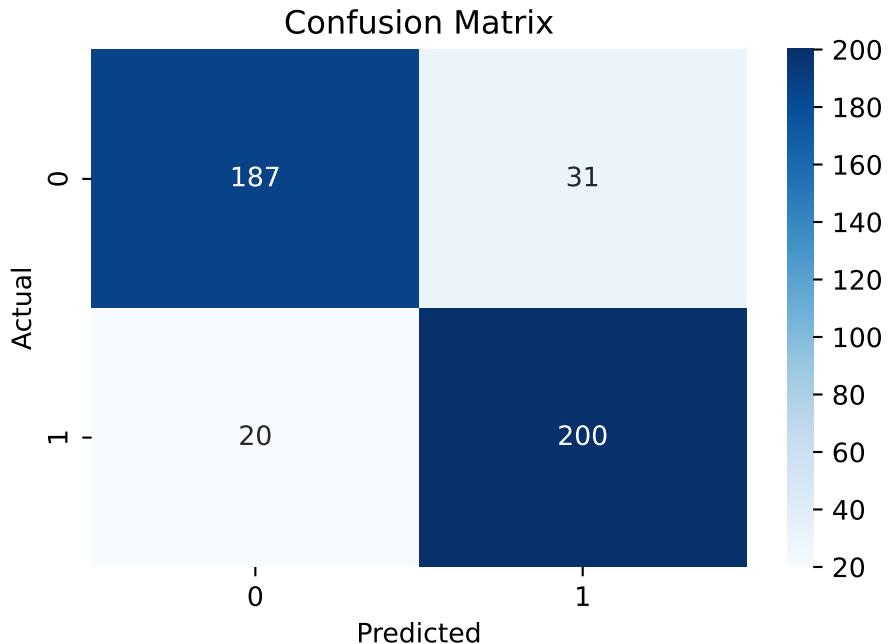
# Extract the probability of the TRUE positive class ("High")
pos_index = np.where(gnb.classes_ == "High")[0][0]
y_prob_high = gnb.predict_proba(X_test)[:, pos_index]
```

12.8 5. Confusion Matrix & Classification Report

```
cm = confusion_matrix(y_test, y_pred)
cm

array([[187,  31],
       [ 20, 200]])

sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
High	0.90	0.86	0.88	218
Low	0.87	0.91	0.89	220
accuracy			0.88	438
macro avg	0.88	0.88	0.88	438
weighted avg	0.88	0.88	0.88	438

12.9 6. ROC Curve (Corrected) and AUC

```
# Convert y_test to binary (1 = High, 0 = Low)
y_test_binary = (y_test == "High").astype(int)

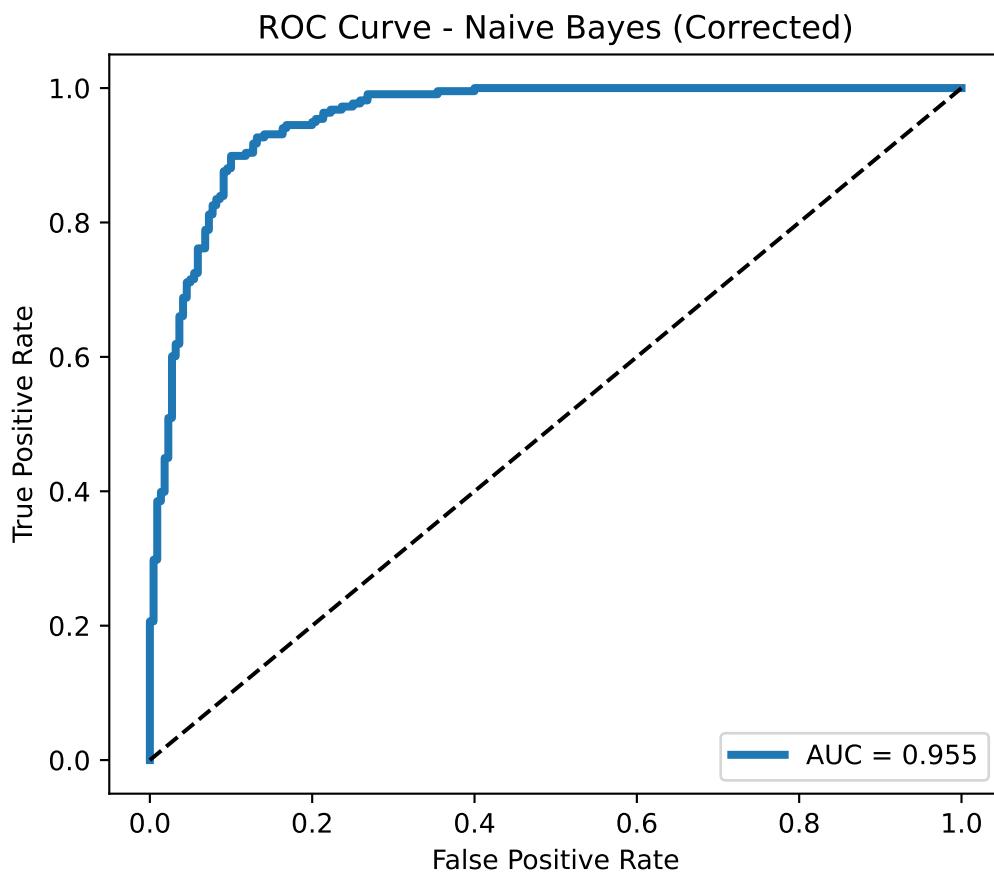
# Compute ROC using correct positive class probability
```

```
fpr, tpr, thresholds = roc_curve(y_test_binary, y_prob_high)
roc_auc = auc(fpr, tpr)
```

```
roc_auc
```

```
np.float64(0.9547122602168474)
```

```
plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f"AUC = {roc_auc:.3f}", linewidth=3)
plt.plot([0,1], [0,1], "k--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Naive Bayes (Corrected)")
plt.legend()
plt.show()
```



12.10 7. Interpretation Notes

12.10.1 Naive Bayes Assumptions

- Predictors are independent given the class.
- Numeric features follow Gaussian distributions.
- One-hot encoded categorical variables behave like Bernoulli variables.

12.10.2 Strengths

- Fast and simple.
- Works well as a baseline.
- Highly interpretable.

12.10.3 Weaknesses

- Independence assumption rarely holds.
 - Gaussian assumption is an approximation.
-

12.11 8. Summary

This chapter:

- Loaded the Ames Housing dataset.
- Created a binary classification target.
- Fit Gaussian Naive Bayes.
- Correctly computed ROC & AUC.
- Provided fully aligned interpretation with the R version.

13 model-agnostic-interpretability

14 support-vector-machines