

GAM— Piecewise, LOESS, and GAM Splines - Python Version

Note: The original MARS section used the `sklearn-contrib-py-earth` package, which is **no longer actively maintained** and may fail to install on modern Python versions.

The MARS code is left **commented out** for reference — you can un-comment it if you successfully install `py-earth` (try `conda install -c conda-forge sklearn-contrib-py-earth`).

Equivalent flexibility can be achieved through **GAM** or **spline-based models**, which are shown below.

1. Setup and Data

```
1 import sys
2 print("Active Python interpreter:", sys.executable)
3
4 import numpy as np, pandas as pd
5 import matplotlib.pyplot as plt
6
7 from sklearn.model_selection import train_test_split
8 from sklearn.compose import ColumnTransformer
9 from sklearn.preprocessing import OneHotEncoder, StandardScaler
10 from sklearn.pipeline import Pipeline
11 from sklearn.metrics import mean_squared_error, r2_score
12 from sklearn.linear_model import LinearRegression, Ridge
13 import inspect
14
15 # --- Fix for SciPy >= 1.13 removing .A from sparse matrices ---
16 import scipy.sparse as sp
17
18 if not hasattr(sp.spmatrix, "A"):
```

```

19     def _toarray(self):
20         return self.toarray()
21     sp.spmatrix.A = property(_toarray)
22 # -----
23
24 np.random.seed(4321)

```

Active Python interpreter: /opt/anaconda3/bin/python

```

1 from sklearn.datasets import fetch_openml
2
3 ames = fetch_openml(name="house_prices", as_frame=True).frame
4
5 keep = {
6     "SalePrice": "SalePrice",
7     "BedroomAbvGr": "Bedroom_AbvGr",
8     "YearBuilt": "Year_Built",
9     "MoSold": "Mo_Sold",
10    "LotArea": "Lot_Area",
11    "Street": "Street",
12    "CentralAir": "Central_Air",
13    "1stFlrSF": "First_Flr_SF",
14    "2ndFlrSF": "Second_Flr_SF",
15    "FullBath": "Full_Bath",
16    "HalfBath": "Half_Bath",
17    "Fireplaces": "Fireplaces",
18    "GarageArea": "Garage_Area",
19    "GrLivArea": "Gr_Liv_Area",
20    "TotRmsAbvGrd": "TotRms_AbvGrd"
21 }
22
23 df = ames[list(keep.keys())].rename(columns=keep).dropna().copy()
24
25 X = df.drop(columns=["SalePrice"])
26 y = df["SalePrice"].values
27
28 num_cols = X.select_dtypes(include=[np.number]).columns.tolist()
29 cat_cols = [c for c in X.columns if c not in num_cols]
30
31 # Handle OneHotEncoder version changes and ensure dense output
32 if "sparse_output" in inspect.signature(OneHotEncoder).parameters:

```

```

33     encoder = OneHotEncoder(drop="first", handle_unknown="ignore", sparse_output=False)
34 else:
35     encoder = OneHotEncoder(drop="first", handle_unknown="ignore", sparse=False)
36
37 pre = ColumnTransformer([
38     ("num", StandardScaler(), num_cols),
39     ("cat", encoder, cat_cols)
40 ], sparse_threshold=0)
41
42 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=4321)
43 df.head()

```


	SalePrice	Bedroom_AbvGr	Year_Built	Mo_Sold	Lot_Area	Street	Central_Air	First_Flr_SF
0	208500	3	2003	2	8450	Pave	Y	856
1	181500	3	1976	5	9600	Pave	Y	1262
2	223500	3	2001	9	11250	Pave	Y	920
3	140000	3	1915	2	9550	Pave	Y	961
4	250000	4	2000	12	14260	Pave	Y	1145

2. Visual Check for Nonlinearity

```

1 import statsmodels.api as sm
2 fig = plt.figure()
3 plt.scatter(df["Gr_Liv_Area"], df["SalePrice"], s=8, alpha=0.3)
4 low = sm.nonparametric.lowess(df["SalePrice"], df["Gr_Liv_Area"], frac=0.3, return_sorted=True)
5 plt.plot(low[:,0], low[:,1], linewidth=2)
6 plt.title("SalePrice vs Gr_Liv_Area (LOESS smoother)")
7 plt.xlabel("Gr_Liv_Area"); plt.ylabel("SalePrice")
8 plt.tight_layout()

```



3. Piecewise Regression Example

```


1  class PiecewiseLinear:
2      def __init__(self, knot=2000.0):
3          self.knot = knot
4          self.lin = LinearRegression()
5
6      def _transform(self, x):
7          x1 = np.asarray(x).reshape(-1,1)
8          hx = np.maximum(0, x1 - self.knot)
9          return np.hstack([x1, hx])
10
11     def fit(self, x, y):
12         Z = self._transform(x)
13         self.lin.fit(Z, y)
14         return self
15
16     def predict(self, x):
17         Z = self._transform(x)

```

```

18     return self.lin.predict(Z)
19
20 pw = PiecewiseLinear(knot=2000).fit(df["Gr_Liv_Area"].values, df["SalePrice"].values)
21
22 grid = np.linspace(df["Gr_Liv_Area"].min(), df["Gr_Liv_Area"].max(), 200)
23 pred = pw.predict(grid)
24
25 plt.figure()
26 plt.scatter(df["Gr_Liv_Area"], df["SalePrice"], s=8, alpha=0.2)
27 plt.plot(grid, pred, linewidth=2)
28 plt.axvline(2000, linestyle="--")
29 plt.title("Piecewise Linear Fit at Knot = 2000")
30 plt.xlabel("Gr_Liv_Area"); plt.ylabel("SalePrice")
31 plt.tight_layout()

```



4. MARS (Commented Out — For Reference Only)

```

1 # The MARS implementation (py-earth) is not actively maintained and may fail to install on m
2 # If you wish to try it, install from conda-forge:
3 #     conda install -c conda-forge sklearn-contrib-py-earth
4 #
5 # Example (commented out):
6 #
7 # from pyearth import Earth
8 #
9 # mars = Pipeline([("prep", pre), ("model", Earth(max_degree=2))]).fit(X_train, y_train)
10 # print(mars.named_steps["model"].summary())
11 #
12 # xs = np.linspace(X["Garage_Area"].min(), X["Garage_Area"].max(), 200)
13 # yhat = mars.predict(pd.DataFrame({"Garage_Area": xs}))
14 #
15 # plt.figure()
16 # plt.scatter(X_train["Garage_Area"], y_train, s=8, alpha=0.2)
17 # plt.plot(xs, yhat, linewidth=2)
18 # plt.title("MARS with Garage_Area (Commented Out Example)")
19 # plt.xlabel("Garage_Area"); plt.ylabel("SalePrice")
20 # plt.tight_layout()


```

5. LOESS Visualization

```

1 fig = plt.figure()
2 plt.scatter(df["Gr_Liv_Area"], df["SalePrice"], s=8, alpha=0.2)
3 low = sm.nonparametric.lowess(df["SalePrice"], df["Gr_Liv_Area"], frac=0.4, return_sorted=True)
4 plt.plot(low[:,0], low[:,1], linewidth=2)
5 plt.title("LOESS Smoother (Visual Aid)")
6 plt.xlabel("Gr_Liv_Area"); plt.ylabel("SalePrice")
7 plt.tight_layout()

```



6. GAM and Spline Approaches (Modern Alternatives)

```

1 from pygam import LinearGAM, s
2 gam1 = LinearGAM(s(0)).fit(X_train[["Garage_Area"]].values, y_train)
3 print(gam1.summary())
4 gam2 = LinearGAM(s(0) + s(1)).fit(X_train[["Garage_Area", "Gr_Liv_Area"]].values, y_train)
5 print(gam2.summary())
6 num_train = X_train[num_cols].values
7 gam3 = LinearGAM().fit(num_train, y_train)
8 print(gam3.summary())

```

```

LinearGAM
=====
Distribution: NormalDist Effective DoF:
Link Function: IdentityLink Log Likelihood:
23263.1874
Number of Samples: 1022 AIC:
AICc:
GCV:

```

30

Feature Function	Lambda	Rank	EDoF	P > x
s(0)	[0.6]	20	11.7	1.11e-
16 ***				
intercept		1	0.0	1.11e-
16 ***				

Significance codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

WARNING: Fitting splines and a linear function to a feature introduces a model identifiability problem which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models with known smoothing parameters, but when smoothing parameters have been estimated, the p-values are typically lower than they should be, meaning that the tests reject the null too often.

None

LinearGAM

Distribution:	NormalDist	Effective DoF:
Link Function:	IdentityLink	Log Likelihood:
22770.8323		
Number of Samples:	1022	AIC:
		AICc:
		GCV:
		Scale:
		Pseudo R-Squared:

Feature Function Lambda Rank EDoF P > x

Feature Function	Lambda	Rank	EDoF	P > x
s(0)	[0.6]	20	12.6	1.11e-
16 ***				
s(1)	[0.6]	20	8.3	1.11e-
16 ***				
intercept		1	0.0	1.11e-
16 ***				

Significance codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

WARNING: Fitting splines and a linear function to a feature introduces a model identifiability problem which can cause p-values to appear significant when they are not.

which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models with known smoothing parameters, but when smoothing parameters have been estimated, the p-values

are typically lower than they should be, meaning that the tests reject the null too easily.

None

LinearGAM

Distribution:	NormalDist	Effective DoF:	1022	AIC:	1022
Link Function:	IdentityLink	Log Likelihood:	21995.395	AICc:	21995.395
Number of Samples:				GCV:	1022
Feature Function	Lambda	Rank	EDoF	P > x	
s(0)	[0.6]	20	8.7	6.58e-	
07 ***					
s(1)	[0.6]	20	12.8	1.11e-	
16 ***					
s(2)	[0.6]	20	11.0	5.15e-	
01					
s(3)	[0.6]	20	8.4	1.11e-	
16 ***					
s(4)	[0.6]	20	8.6	1.11e-	
16 ***					
s(5)	[0.6]	20	10.1	2.17e-	
03 **					
s(6)	[0.6]	20	3.2	9.76e-	
02 .					
s(7)	[0.6]	20	2.2	4.19e-	
01					
s(8)	[0.6]	20	2.9	3.16e-	
10 ***					
s(9)	[0.6]	20	9.0	1.11e-	
16 ***					
s(10)	[0.6]	20	5.7	1.11e-	
16 ***					
s(11)	[0.6]	20	6.5	3.78e-	

```
01
intercept                               1          0.0      1.11e-
16    ***
=====
Significance codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

WARNING: Fitting splines and a linear function to a feature introduces a model identifiability
         which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models
         with known smoothing parameters, but when smoothing parameters have been estimated, the p-
         values
         are typically lower than they should be, meaning that the tests reject the null too
         often.

None
```

```
/var/folders/v0/1fynz4dx4sd1h7nxwmvzrxf00000gn/T/ipykernel_95010/2456984288.py:3: UserWarning:
```

KNOWN BUG: p-values computed in this summary are likely much smaller than they should be.

Please do not make inferences based on these values!

Collaborate on a solution, and stay up to date at:
github.com/dswah/pyGAM/issues/163

```
/var/folders/v0/1fynz4dx4sd1h7nxwmvzrxf00000gn/T/ipykernel_95010/2456984288.py:5: UserWarning:
```

KNOWN BUG: p-values computed in this summary are likely much smaller than they should be.

Please do not make inferences based on these values!

Collaborate on a solution, and stay up to date at:
github.com/dswah/pyGAM/issues/163

```
/var/folders/v0/1fynz4dx4sd1h7nxwmvzrxf00000gn/T/ipykernel_95010/2456984288.py:8: UserWarning:
```

KNOWN BUG: p-values computed in this summary are likely much smaller than they should be.

Please do not make inferences based on these values!

Collaborate on a solution, and stay up to date at:

7. Regression Spline Example (Fallback)

```
1 from patsy import dmatrix
2
3 design_tr = dmatrix("bs(Garage_Area, df=6) + bs(Gr_Liv_Area, df=6)", data=X_train, return_type='dataframe')
4 design_te = dmatrix("bs(Garage_Area, df=6) + bs(Gr_Liv_Area, df=6)", data=X_test, return_type='dataframe')
5
6 ridge = Ridge(alpha=1.0).fit(design_tr, y_train)
7 pred_spline = ridge.predict(design_te)
```

8. Model Comparison (Piecewise vs GAM vs Spline)

```
1 def rmse(y, yhat):
2     try:
3         return mean_squared_error(y, yhat, squared=False)
4     except TypeError:
5         return np.sqrt(mean_squared_error(y, yhat))
6
7 def mse(y, yhat): return mean_squared_error(y, yhat)
8 def rsq(y, yhat): return r2_score(y, yhat)
9
10 p_piece = pw.predict(X_test["Gr_Liv_Area"].values)
11
12 try:
13     p_gam = gam3.predict(X_test[num_cols].values)
14 except Exception:
15     p_gam = np.full_like(y_test, np.nan, dtype=float)
16
17 p_spline = pred_spline
18
19 cmp = pd.DataFrame({
```

```

20 "Model": ["Piecewise", "GAM (pyGAM)" if not np.isnan(p_gam).all() else "GAM (not available)", "RMSE": [rmse(y_test, p_piece), rmse(y_test, p_gam) if not np.isnan(p_gam).all() else np.nan], "MSE": [mse(y_test, p_piece), mse(y_test, p_gam) if not np.isnan(p_gam).all() else np.nan], "R2": [rsq(y_test, p_piece), rsq(y_test, p_gam) if not np.isnan(p_gam).all() else np.nan]}, {"Model": "Regression Spline", "RMSE": 60289.000524, "MSE": 3.634764e+09, "R2": 0.490419}], cmp

```


	Model	RMSE	MSE	R2
1	GAM (pyGAM)	40805.867162	1.665119e+09	0.766556
0	Piecewise	58108.537553	3.376602e+09	0.526612
2	Regression Spline	60289.000524	3.634764e+09	0.490419

9. Predicted vs Actual Comparison

```

1 plt.figure()
2 plt.scatter(y_test, p_piece, s=8, alpha=0.4, label="Piecewise")
3 if not np.isnan(p_gam).all():
4     plt.scatter(y_test, p_gam, s=8, alpha=0.4, label="GAM")
5 plt.scatter(y_test, p_spline, s=8, alpha=0.4, label="Spline")
6 plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], linestyle="--")
7 plt.legend()
8 plt.title("Predicted vs Actual (Test Set)")
9 plt.xlabel("Actual SalePrice"); plt.ylabel("Predicted SalePrice")
10 plt.tight_layout()

```



10. Interpretation Summary

- **Piecewise:** adds simple thresholds, easy to explain.
- **GAMs:** provide smooth, interpretable nonlinearities.
- **Splines:** flexible approximations that behave like local polynomials.
- **MARS (optional):** similar conceptually but less maintained in Python; consider using R's `earth` instead.