

GAM – Piecewise, MARS, LOESS, and Splines

- R Version

Goal: Predict `Sale_Price` from a focused set of Ames features using increasingly flexible methods in the GAM framework.

1. Setup and Data

```
1 set.seed(4321)
2
3 library(tidyverse)
4 library(dplyr)
5 library(tidyr)
6 library(AmesHousing)    # make_ordinal_ames()
7 library(earth)          # MARS
8 library(segmented)      # piecewise (segmented) regression
9 library(splines)         # regression splines (bs, ns) - kept for reference
10 library(mgcv)           # smoothing splines via GAM
11 library(caret)          # data splitting / utilities
12 library(Metrics)        # rmse, mse helpers (plus we'll compute R2 manually)
13 theme_set(theme_minimal())
```

```
1 # Limit to requested columns
2 cols <- c("Sale_Price", "Bedroom_AbvGr", "Year_Built", "Mo_Sold", "Lot_Area", "Street",
3          "Central_Air", "First_Flr_SF", "Second_Flr_SF", "Full_Bath", "Half_Bath",
4          "Fireplaces", "Garage_Area", "Gr_Liv_Area", "TotRms_AbvGrd")
5
6 ames <- make_ordinal_ames() |>
7   dplyr::select(dplyr::all_of(cols)) |>
8   tidyr::drop_na()
```

```

10 # Train/test split
11 set.seed(4321)
12 idx <- sample.int(nrow(ames), size = floor(0.7*nrow(ames)))
13 train <- ames[idx,]
14 test <- ames[-idx,]
15
16 glimpse(train)

```

Rows: 2,051
 Columns: 15
\$ Sale_Price <int> 274000, 75200, 329900, 145400, 108000, 184000, 176000, 1~
\$ Bedroom_AbvGr <int> 3, 2, 4, 3, 2, 2, 3, 3, 3, 3, 3, 3, 4, 3, 2, 3, 3, 2,~
\$ Year_Built <int> 2001, 1922, 2005, 1926, 1949, 1999, 1962, 1915, 1999, 19~
\$ Mo_Sold <int> 1, 9, 8, 5, 5, 6, 5, 6, 9, 6, 8, 5, 7, 7, 4, 8, 5, 11, 5~
\$ Lot_Area <int> 9720, 3672, 11643, 7000, 8777, 5858, 19296, 8094, 3768, ~
\$ Street <fct> Pave, Pave, Pave, Pave, Pave, Pave, Pave, Pave, Pave, Pa~
\$ Central_Air <fct> Y, Y,~
\$ First_Flr_SF <int> 1366, 816, 1544, 861, 1126, 1337, 1382, 1048, 713, 792, ~
\$ Second_Flr_SF <int> 581, 0, 814, 424, 0, 0, 720, 739, 725, 0, 1151, 695, ~
\$ Full_Bath <int> 2, 1, 2, 1, 2, 2, 1, 2, 1, 1, 2, 2, 2, 1, 2, 2, 1, 1,~
\$ Half_Bath <int> 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0,~
\$ Fireplaces <int> 1, 0, 1, 0, 0, 1, 1, 0, 0, 2, 0, 1, 2, 1, 0, 1, 0, 0, 1,~
\$ Garage_Area <dbl> 725, 100, 784, 506, 520, 511, 884, 576, 506, 400, 0, 434~
\$ Gr_Liv_Area <int> 1947, 816, 2358, 1285, 1126, 1337, 1382, 1768, 1452, 151~
\$ TotRms_AbvGrd <int> 7, 5, 10, 6, 5, 5, 6, 8, 6, 7, 5, 8, 7, 5, 6, 7, 6, 5~

```
1 summary(train$Sale_Price)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
12789	130000	161000	180897	215000	755000	

2. Quick Visual: Nonlinearity is Common

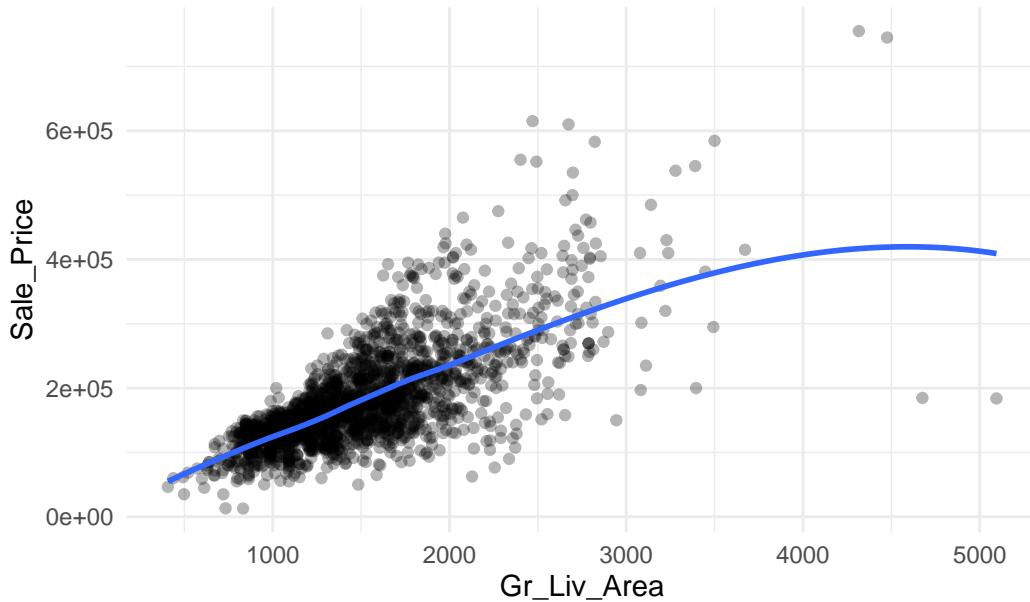
```

1 ggplot(train, aes(Gr_Liv_Area, Sale_Price)) +
2   geom_point(alpha=.3) +
3   geom_smooth(method="loess", se=FALSE) +
4   labs(title="Sale_Price vs Gr_Liv_Area (LOESS smoother)")

```

```
`geom_smooth()` using formula = 'y ~ x'
```

Sale_Price vs Gr_Liv_Area (LOESS smoother)



Takeaway: The relationship bends—nonlinear methods can help.

3. Piecewise Regression

```
1 m_lin <- lm(Sale_Price ~ Gr_Liv_Area, data=train)
2 m_seg_init <- lm(Sale_Price ~ Gr_Liv_Area, data=train)
3 m_seg <- segmented(m_seg_init, seg.Z = ~ Gr_Liv_Area, psi = list(Gr_Liv_Area = 2000))
4 summary(m_seg)
```

Regression Model with Segmented Relationship(s)

Call:
segmented.lm(obj = m_seg_init, seg.Z = ~Gr_Liv_Area, psi = list(Gr_Liv_Area = 2000))

Estimated Break-Point(s):

```

      Est.  St.Err
psi1.Gr_Liv_Area 2466 332.881

Coefficients of the linear terms:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 7901.186   4698.641   1.682   0.0928 .
Gr_Liv_Area   115.705     3.171  36.488 <2e-16 ***
U1.Gr_Liv_Area -23.479    12.436  -1.888      NA
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 55110 on 2047 degrees of freedom
Multiple R-Squared: 0.5124, Adjusted R-squared: 0.5116

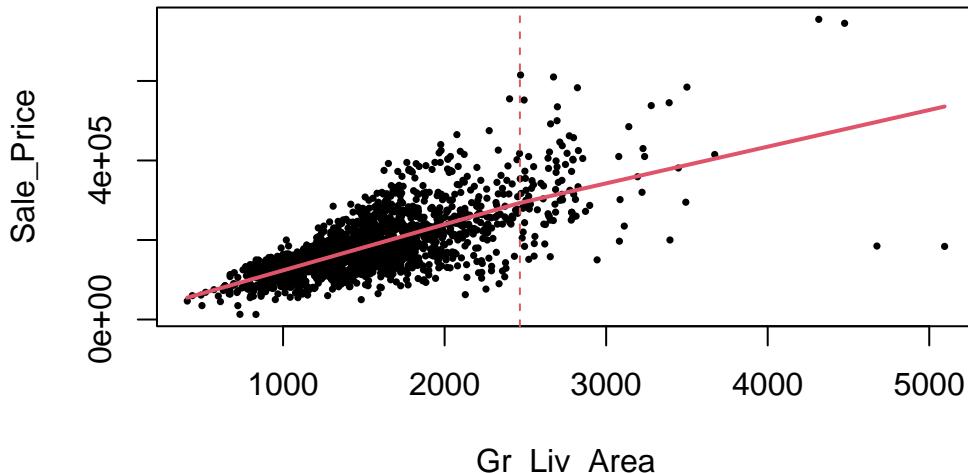
Boot restarting based on 6 samples. Last fit:
Convergence attained in 1 iterations (rel. change 0)

```

1 plot(train$Gr_Liv_Area, train$Sale_Price, pch=16, cex=.5,
2       xlab="Gr_Liv_Area", ylab="Sale_Price", main="Piecewise fit at estimated knot")
3 plot(m_seg, add=TRUE, col=2, lwd=2)
4 abline(v = m_seg$psi[, "Est."], lty=2, col=2)

```

Piecewise fit at estimated knot



Interpretation (piecewise): Two linear slopes before/after a knot. Useful when you expect

a threshold effect.

4. MARS (earth) – From Simple to Interactions

MARS = *Multivariate Adaptive Regression Splines*. It builds **hinge** basis functions like $h(x - c) = \max(0, x - c)$ and can **interact** them.

- A term such as $h(\text{Gr_Liv_Area} - 1500)$ means *once* Gr_Liv_Area exceeds 1500, the fitted line's slope can change.
- An interaction like $h(\text{Gr_Liv_Area} - 1500) * \text{Central_AirY}$ means the slope change applies when $\text{Central_Air} == "Y"$.

4A. Univariate MARS (Garage_Area)

```
1 mars1 <- earth(Sale_Price ~ Garage_Area, data=train)
2 summary(mars1)
```

Call: `earth(formula=Sale_Price~Garage_Area, data=train)`

	coefficients
(Intercept)	124159.039
$h(286-Garage_Area)$	-60.257
$h(Garage_Area-286)$	297.277
$h(Garage_Area-521)$	-483.642
$h(Garage_Area-576)$	733.859
$h(Garage_Area-758)$	-356.460
$h(Garage_Area-1043)$	-490.873

Selected 7 of 7 terms, and 1 of 1 predictors

Termination condition: RSq changed by less than 0.001 at 7 terms

Importance: Garage_Area

Number of terms at each degree of interaction: 1 6 (additive model)

GCV 3427475346 RSS 6.94092e+12 GRSq 0.4492014 RSq 0.4556309

```
1 grid <- tibble(Garage_Area = seq(min(train$Garage_Area), max(train$Garage_Area), length.out=200))
2 grid$pred <- predict(mars1, newdata=grid)
3
4 ggplot(train, aes(Garage_Area, Sale_Price)) +
```

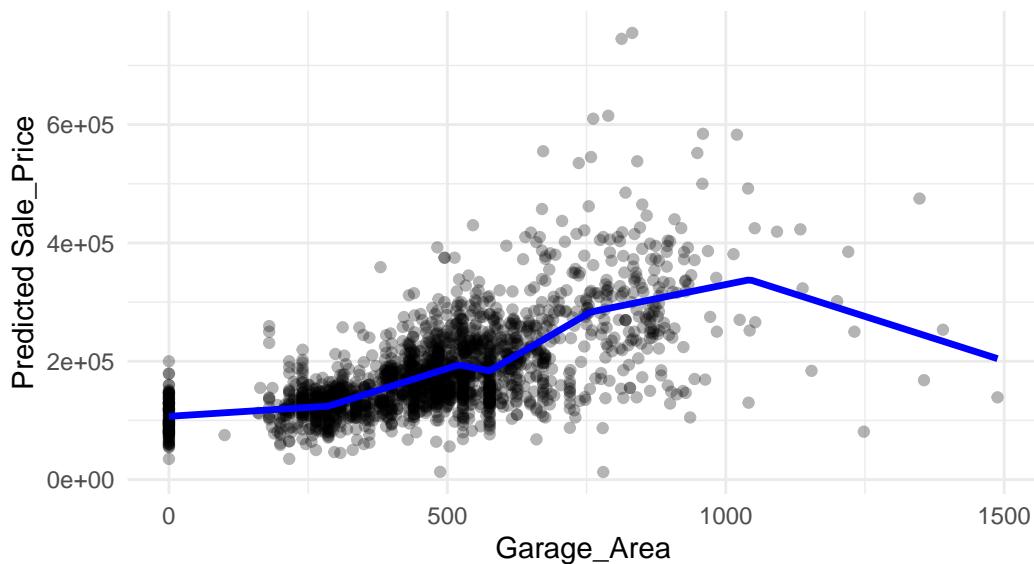
```

5   geom_point(alpha=.3) +
6   geom_line(data=grid, aes(Garage_Area, pred), color="blue", linewidth=1.2) +
7   labs(title="Step 1: MARS with One Predictor (Garage_Area)",
8       subtitle="Piecewise linear fit with automatically chosen knots",
9       y="Predicted Sale_Price")

```

Step 1: MARS with One Predictor (Garage_Area)

Piecewise linear fit with automatically chosen knots



4B. Additive MARS (degree = 1; no interactions)

```

1 mars2 <- earth(Sale_Price ~ Bedroom_AbvGr + Year_Built + Mo_Sold + Lot_Area +
2                               Street + Central_Air + First_Flr_SF + Second_Flr_SF + Full_Bath +
3                               Half_Bath + Fireplaces + Garage_Area + Gr_Liv_Area + TotRms_AbvGrd,
4                               data=train, degree=1, nfold=5)
5 summary(mars2)

```

```
Call: earth(formula=Sale_Price~Bedroom_AbvGr+Year_Built+Mo_Sold+Lot_Ar...),
      data=train, degree=1, nfold=5)
```

	coefficients
(Intercept)	319493.46
Central_AirY	20289.49

h(4-Bedroom_AbvGr)	9214.66
h(Bedroom_AbvGr-4)	-23009.05
h(Year_Built-1977)	1275.57
h(2004-Year_Built)	-336.64
h(Year_Built-2004)	5315.57
h(13869-Lot_Area)	-2.09
h(Lot_Area-13869)	0.22
h(First_Flr_SF-1600)	104.91
h(2402-First_Flr_SF)	-71.56
h(First_Flr_SF-2402)	-176.61
h(1523-Second_Flr_SF)	-53.13
h(Second_Flr_SF-1523)	426.63
h(Half_Bath-1)	-45378.31
h(2-Fireplaces)	-14408.56
h(Fireplaces-2)	-26072.58
h(Garage_Area-539)	101.97
h(Garage_Area-1043)	-294.30
h(Gr_Liv_Area-2049)	65.21
h(Gr_Liv_Area-3194)	-159.79

Selected 21 of 24 terms, and 10 of 14 predictors

Termination condition: Reached nk 29

Importance: First_Flr_SF, Second_Flr_SF, Year_Built, Garage_Area, ...

Number of terms at each degree of interaction: 1 20 (additive model)

GCV 1033819964 RSS 2.036439e+12 GRSq 0.8338641 RSq 0.8402842 CVRSq 0.7978671

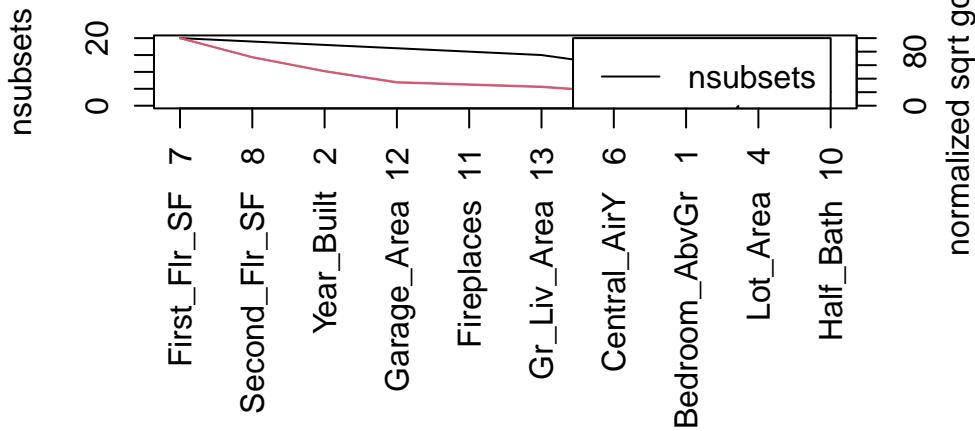
Note: the cross-validation sd's below are standard deviations across folds

Cross validation: nterms 21.00 sd 0.71 nvars 9.20 sd 0.84

CVRSq	sd	MaxErr	sd
0.798	0.043	-428941	282429

```
1 ev2 <- evimp(mars2); plot(ev2, main="Step 2: Variable Importance (degree=1, additive)")
```

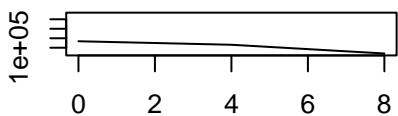
Step 2: Variable Importance (degree=1, additive)



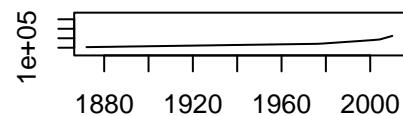
```
1 par(mfrow=c(2,2)); plotmo(mars2, type="response", nresponse=1, do.par=FALSE); par(mfrow=c(1,1))
```

```
plotmo grid: Bedroom_AbvGr Year_Built Mo_Sold Lot_Area Street Central_Air
            3           1974      6     9480   Pave       Y
First_Flr_SF Second_Flr_SF Full_Bath Half_Bath Fireplaces Garage_Area
        1092             0         2         0         1        480
Gr_Liv_Area TotRms_AbvGrd
        1442             6
```

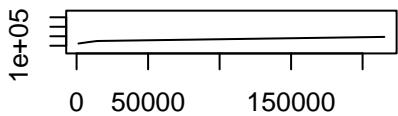
1 Bdrm_AbvG



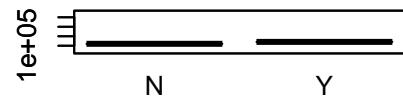
2 Year_Bult



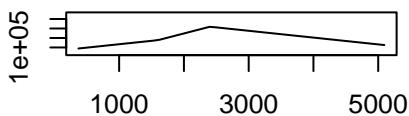
3 Lot_Area



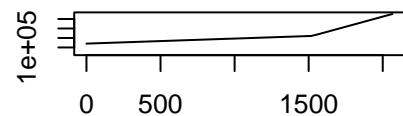
4 Centrl_Ar



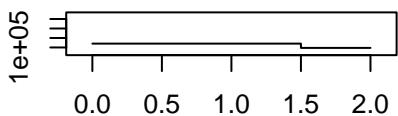
5 Frst_F_SF



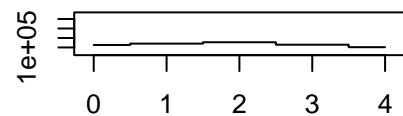
6 Scnd_F_SF

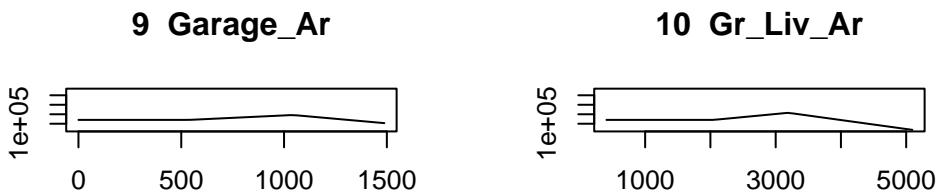


7 Half_Bath



8 Fireplacs





4C. MARS with 2-way Interactions (degree = 2)

```

1 mars3 <- earth(Sale_Price ~ Bedroom_AbvGr + Year_Built + Mo_Sold + Lot_Area +
2                               Street + Central_Air + First_Flr_SF + Second_Flr_SF + Full_Bath +
3                               Half_Bath + Fireplaces + Garage_Area + Gr_Liv_Area + TotRms_AbvGrd,
4                               data=train, degree=2, nfold=5)
5 summary(mars3)

```

Call: earth(formula=Sale_Price~Bedroom_AbvGr+Year_Built+Mo_Sold+Lot_Ar...),
 data=train, degree=2, nfold=5)

	coefficients
(Intercept)	411949.27
h(Year_Built-2004)	169869.37
h(14803-Lot_Area)	-1.75
h(Lot_Area-14803)	0.48
h(1570-First_Flr_SF)	-150.79
h(First_Flr_SF-1570)	224.90
h(1523-Second_Flr_SF)	-86.91
h(Second_Flr_SF-1523)	156.02
h(1-Half_Bath)	-7743.19

h(Half_Bath-1)	-54819.24
h(2-Fireplaces)	-10516.92
h(1043-Garage_Area)	-70.97
h(Garage_Area-1043)	-141.86
h(Gr_Liv_Area-3194)	106.27
h(3-Bedroom_AbvGr) * h(1523-Second_Flr_SF)	7.01
h(Bedroom_AbvGr-3) * h(1523-Second_Flr_SF)	-9.61
h(2004-Year_Built) * h(First_Flr_SF-876)	-1.86
h(2004-Year_Built) * h(3194-Gr_Liv_Area)	-0.17
h(Year_Built-2004) * h(Gr_Liv_Area-2320)	-159.80
h(Year_Built-2004) * h(2320-Gr_Liv_Area)	156.64
h(Year_Built-2004) * h(3194-Gr_Liv_Area)	-173.05
h(First_Flr_SF-1778) * h(2-Fireplaces)	-78.19
h(1043-Garage_Area) * h(2122-Gr_Liv_Area)	0.05

Selected 23 of 29 terms, and 9 of 14 predictors

Termination condition: Reached nk 29

Importance: First_Flr_SF, Second_Flr_SF, Year_Built, Gr_Liv_Area, ...

Number of terms at each degree of interaction: 1 13 9

GCV 795873462 RSS 1.544416e+12 GRSq 0.8721024 RSq 0.8788731 CVRSq 0.7824362

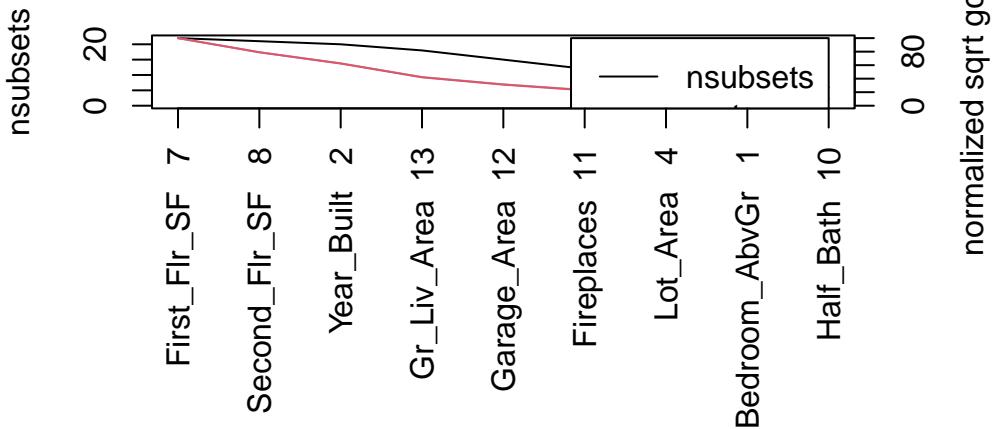
Note: the cross-validation sd's below are standard deviations across folds

Cross validation: nterms 22.60 sd 1.82 nvars 9.40 sd 1.14

CVRSq	sd	MaxErr	sd
0.782	0.114	-945502	422348

```
1 ev3 <- evimp(mars3); plot(ev3, main="Step 3: Variable Importance (degree=2, interactions)")
```

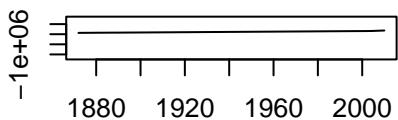
Step 3: Variable Importance (degree=2, interactions)



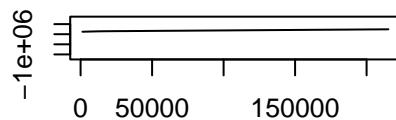
```
1 par(mfrow=c(2,2)); plotmo(mars3, type="response", nresponse=1, do.par=FALSE); par(mfrow=c(1,1))
```

plotmo grid:	Bedroom_AbvGr	Year_Built	Mo_Sold	Lot_Area	Street	Central_Air
	3	1974	6	9480	Pave	Y
First_Flr_SF	Second_Flr_SF	Full_Bath	Half_Bath	Fireplaces	Garage_Area	
1092		0	2	0	1	480
Gr_Liv_Area	TotRms_AbvGrd					
1442	6					

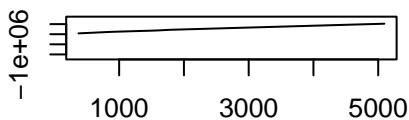
1 Year_Bult



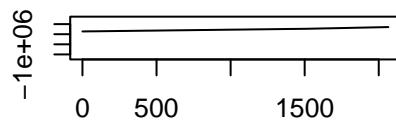
2 Lot_Area



3 Frst_F_SF



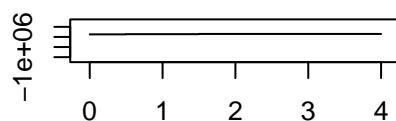
4 Scnd_F_SF



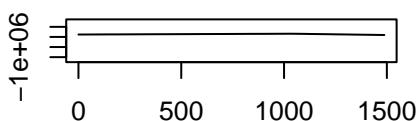
5 Half_Bath



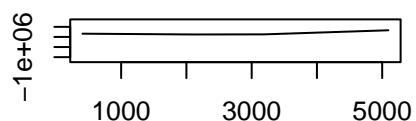
6 Fireplaces



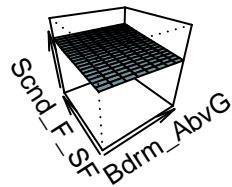
7 Garage_Ar



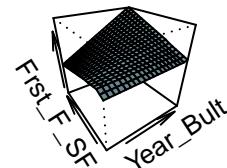
8 Gr_Liv_Ar



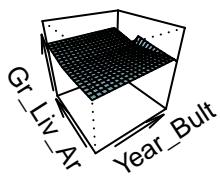
1 Bdrm_AbvG: Scnd_F_SF



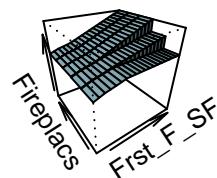
2 Year_Bult: Frst_F_SF



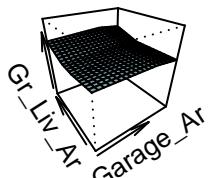
3 Year_Bult: Gr_Liv_Ar



4 Frst_F_SF: Fireplacs



5 Garage_Ar: Gr_Liv_Ar

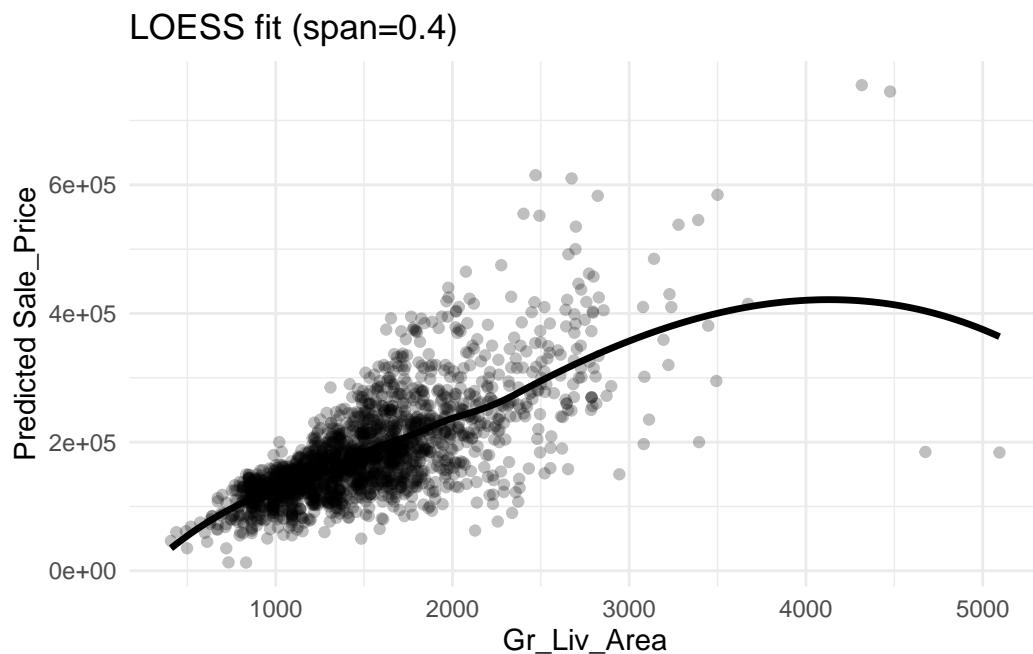


Interpreting MARS terms:

Look for $h()$ pieces and products. Coefficients on $h(x-c)$ indicate how slope changes after the knot c . Interaction products mean “the slope change depends on another variable”.

5. LOESS (Visual local regression)

```
1 loess_fit <- loess(Sale_Price ~ Gr_Liv_Area, data=train, span=0.4)
2 grid_lo <- tibble(Gr_Liv_Area = seq(min(train$Gr_Liv_Area), max(train$Gr_Liv_Area), length.out=100))
3 grid_lo$pred <- predict(loess_fit, newdata=grid_lo)
4
5 ggplot() +
6   geom_point(data=train, aes(Gr_Liv_Area, Sale_Price), alpha=.25) +
7   geom_line(data=grid_lo, aes(Gr_Liv_Area, pred), linewidth=1.2) +
8   labs(title="LOESS fit (span=0.4)", y="Predicted Sale_Price")
```



We keep LOESS as a visualization-oriented smoother (not used in the multivariate comparison below).

6. GAM Splines with mgcv – Simple → Complex

GAM smooths are written `s(x)` and estimated with penalized splines. The `edf` (effective degrees of freedom) in the summary tells you the smooth's complexity (larger edf = wigglier function).

6A. Univariate GAM: Sale_Price ~ s(Garage_Area)

```
1 gam1 <- gam(Sale_Price ~ s(Garage_Area, k=9), data=train, method="REML")
2 summary(gam1) # inspect edf and significance
```

```
Family: gaussian
Link function: identity

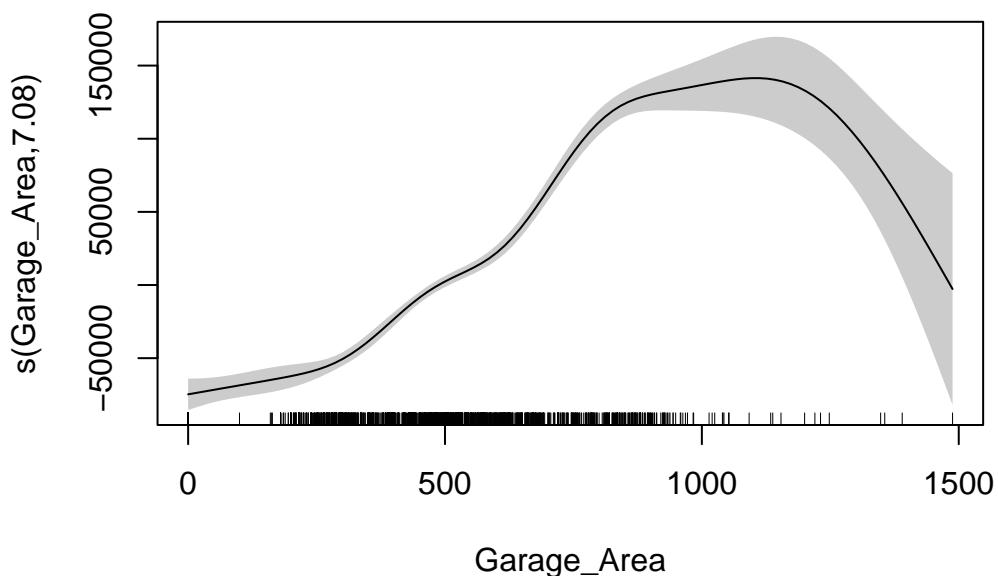
Formula:
Sale_Price ~ s(Garage_Area, k = 9)

Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) 180897        1292      140    <2e-16 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:
            edf Ref.df     F p-value
s(Garage_Area) 7.078 7.719 217.1 <2e-16 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.45 Deviance explained = 45.2%
-REML = 25418 Scale est. = 3.4232e+09 n = 2051
```

```
1 plot(gam1, pages=1, shade=TRUE, scheme=1, scale=0)
```



6B. Two-smooth Additive GAM: $s(\text{Garage_Area}) + s(\text{Gr_Liv_Area})$

```

1 gam2 <- gam(Sale_Price ~ s(Garage_Area, k=9) + s(Gr_Liv_Area, k=9), data=train, method="REML")
2 summary(gam2)

```

Family: gaussian
 Link function: identity

Formula:
 $\text{Sale_Price} \sim s(\text{Garage_Area}, k = 9) + s(\text{Gr_Liv_Area}, k = 9)$

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	180897	1026	176.2	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
$s(\text{Garage_Area})$	6.446	7.293	103.3	<2e-16 ***

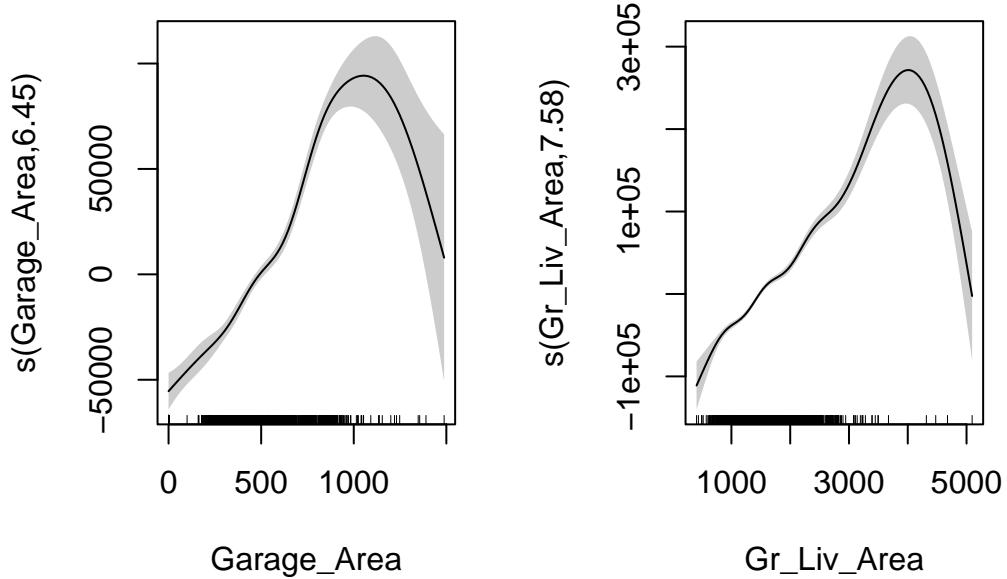
```

s(Gr_Liv_Area) 7.585 7.941 152.6 <2e-16 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.653 Deviance explained = 65.5%
-REML = 24949 Scale est. = 2.1608e+09 n = 2051

1 plot(gam2, pages=1, shade=TRUE, scheme=1, scale=0)

```



6C. Full Multivariate GAM (selected smooths + factors)

```

1 gam3 <- gam(Sale_Price ~
2   s(Garage_Area, k=9) +
3   s(Gr_Liv_Area, k=9) +
4   s(Year_Built, k=7) +
5   s(Lot_Area, k=7) +
6   Bedroom_AbvGr + Mo_Sold +
7   Street + Central_Air +
8   First_Flr_SF + Second_Flr_SF +
9   Full_Bath + Half_Bath +

```

```

10         Fireplaces + TotRms_AbvGrd,
11         data=train, method="REML")
12 summary(gam3)

```

Family: gaussian
Link function: identity

Formula:
Sale_Price ~ s(Garage_Area, k = 9) + s(Gr_Liv_Area, k = 9) +
s(Year_Built, k = 7) + s(Lot_Area, k = 7) + Bedroom_AbvGr +
Mo_Sold + Street + Central_Air + First_Flr_SF + Second_Flr_SF +
Full_Bath + Half_Bath + Fireplaces + TotRms_AbvGrd

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-21237.92	30106.84	-0.705	0.4806
Bedroom_AbvGr	-10350.20	1382.38	-7.487	1.05e-13 ***
Mo_Sold	-409.12	270.24	-1.514	0.1302
StreetPave	30589.21	13486.61	2.268	0.0234 *
Central_AirY	17189.68	3366.03	5.107	3.59e-07 ***
First_Flr_SF	140.27	17.44	8.044	1.47e-15 ***
Second_Flr_SF	101.93	17.36	5.871	5.06e-09 ***
Full_Bath	-4548.94	2211.69	-2.057	0.0398 *
Half_Bath	2355.98	2175.40	1.083	0.2789
Fireplaces	13066.61	1371.59	9.527	< 2e-16 ***
TotRms_AbvGrd	-1640.87	948.67	-1.730	0.0838 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

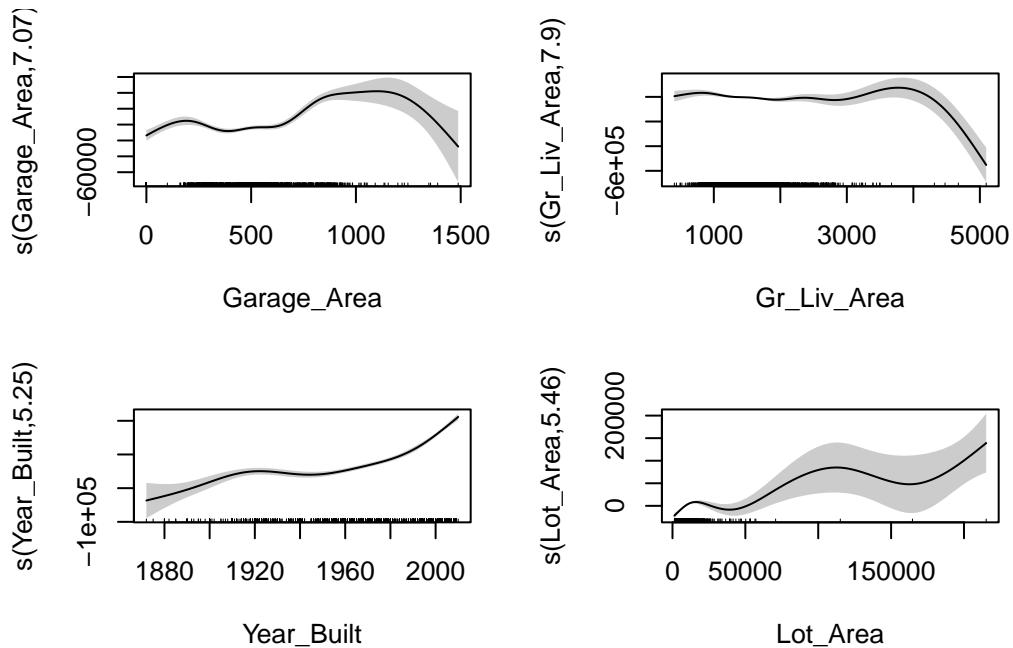
Approximate significance of smooth terms:

edf	Ref.df	F	p-value
s(Garage_Area)	7.067	7.712	24.96 <2e-16 ***
s(Gr_Liv_Area)	7.899	7.996	46.44 <2e-16 ***
s(Year_Built)	5.255	5.764	128.05 <2e-16 ***
s(Lot_Area)	5.465	5.883	16.44 <2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.827 Deviance explained = 83%
-REML = 24167 Scale est. = 1.0786e+09 n = 2051

```
1 plot(gam3, pages=1, shade=TRUE, scheme=1, scale=0)
```



Interpreting GAM splines:

A smooth $s(x)$ shows how the expected outcome varies with x *holding other terms constant*. The edf indicates complexity; wide confidence bands suggest greater uncertainty in those regions.

7. Model Comparison (R^2 , RMSE, MSE)

```
1 # Helper for R^2 on test
2 rsq <- function(y, yhat) {
3   1 - sum((y - yhat)^2) / sum((y - mean(y))^2)
4 }
5
6 # Predictions
7 p_piece <- predict(m_seg, newdata=test)
8 p_mars  <- predict(mars3, newdata=test)    # final MARS with interactions
9 p_gam   <- predict(gam3, newdata=test)      # full multivariate GAM
```

```

10
11 # Metrics
12 tbl <- tibble(
13   Model = c("Piecewise (segmented)", "MARS (degree=2)", "GAM splines (mgcv full)"),
14   RMSE  = c(rmse(test$Sale_Price, p_piece),
15             rmse(test$Sale_Price, p_mars),
16             rmse(test$Sale_Price, p_gam)),
17   MSE   = c(mse(test$Sale_Price, p_piece),
18             mse(test$Sale_Price, p_mars),
19             mse(test$Sale_Price, p_gam)),
20   R2    = c(rsq(test$Sale_Price, p_piece),
21             rsq(test$Sale_Price, p_mars),
22             rsq(test$Sale_Price, p_gam))
23 ) |> arrange(desc(R2))
24
25 tbl

```

```

# A tibble: 3 x 4
  Model           RMSE       MSE      R2
  <chr>        <dbl>     <dbl>  <dbl>
1 MARS (degree=2) 33028. 1090840972. 0.839
2 GAM splines (mgcv full) 36796. 1353918715. 0.800
3 Piecewise (segmented) 59034. 3484963067. 0.484

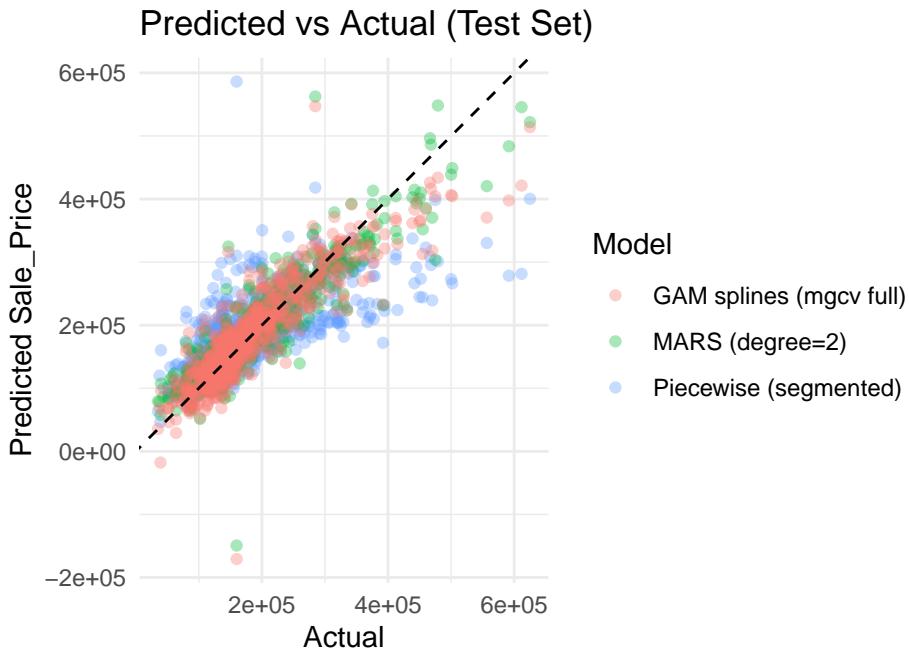
```

8. Final Plot: Predicted vs Actual (Test Set)

```

1 plot_df <- bind_rows(
2   tibble(Model="Piecewise (segmented)", Actual=test$Sale_Price, Pred=p_piece),
3   tibble(Model="MARS (degree=2)", Actual=test$Sale_Price, Pred=p_mars),
4   tibble(Model="GAM splines (mgcv full)", Actual=test$Sale_Price, Pred=p_gam)
5 )
6
7 ggplot(plot_df, aes(Actual, Pred, color=Model)) +
8   geom_point(alpha=.35) +
9   geom_abline(intercept=0, slope=1, linetype="dashed") +
10  labs(title="Predicted vs Actual (Test Set)", y="Predicted Sale_Price") +
11  coord_equal()

```



9. How to Explain

- **Piecewise:** “There’s a threshold where the effect changes.” Simple story, limited flexibility.
- **MARS:** “The model finds breakpoints and sometimes *when* they matter (interactions).” Explain $h(x-c)$ as *extra slope after c*.
- **GAM (splines):** “We model smooth curves for key variables; edf shows wigginess. More edf more flexibility.”
- **Tradeoff:** More flexibility usually improves error (RMSE/MSE) and R^2 , but reduce interpretability. Validate with cross-validation and keep the story aligned with business intuition.